



Use SANtricity solutions

E-Series Systems

NetApp
December 19, 2022

Table of Contents

- Use SANtricity solutions 1
 - Web services proxy 1
 - Remote volume mirroring 34
 - Storage plugin for vCenter 42
 - Legacy solutions 163

Use SANtricity solutions

Web services proxy

SANtricity Web Services Proxy overview

The SANtricity Web Services Proxy is a RESTful API server installed separately on a host system to manage hundreds of new and legacy NetApp E-Series storage systems. The proxy includes SANtricity Unified Manager, which is a web-based interface that provides similar functions.

Installation overview

Installing and configuring the Web Services Proxy involves the following steps:

1. [Review installation and upgrade requirements.](#)
2. [Download and install Web Services Proxy file.](#)
3. [Log in to API and Unified Manager.](#)
4. [Configure Web Services Proxy.](#)

Find more information

- Unified Manager — The proxy installation includes SANtricity Unified Manager, a web-based interface that provides configuration access to newer E-Series and EF-Series storage systems. For more information, see the Unified Manager online help, which is available from its user interface or from the [SANtricity software doc site](#).
- GitHub repository — GitHub contains a repository for the collection and organization of sample scripts illustrating the use of the NetApp SANtricity Web Services API. To access the repository, see [NetApp Webservices samples](#).
- Representational state transfer (REST) — Web Services is a RESTful API that provides access to virtually all the SANtricity management capabilities, so you should be familiar with REST concepts. For more information, see [Architectural Styles and the Design of Network-based Software Architectures](#).
- JavaScript Object Notation (JSON) — Because data within Web Services is encoded through JSON, you should be familiar with JSON programming concepts. For more information, see [Introducing JSON](#).

Learn about Web Services

Web Services and Unified Manager overview

Before you install and configure the Web Services proxy, read the overview of Web Services and SANtricity Unified Manager.

Web Services

Web Services is an Application Programming Interface (API) that allows you to configure, manage, and monitor NetApp E-Series and EF-Series storage systems. By issuing API requests, you can complete workflows such as configuration, provisioning, and performance monitoring for E-Series storage systems.

When using the Web Services API to manage storage systems, you should be familiar with the following:

- JavaScript Object Notation (JSON) – Because data within Web Services is encoded through JSON, you should be familiar with JSON programming concepts. For more information, see [Introducing JSON](#).
- Representational state transfer (REST) – Web Services is a RESTful API that provides access to virtually all the SANtricity management capabilities, so you should be familiar with REST concepts. For more information, see [Architectural Styles and the Design of Network-based Software Architectures](#).
- Programming language concepts – Java and Python are the most common programming languages used with the Web Services API, but any programming language that can make HTTP requests is sufficient for API interaction.

Web Services is available in two implementations:

- **Embedded** — A RESTful API server is embedded on each controller of an E2800/EF280 storage system running NetApp SANtricity 11.30 or later versions, an E5700/EF570 running SANtricity 11.40 or later versions, and an EF300 or EF600 running SANtricity 11.60 or later versions. No installation is required.
- **Proxy** — The SANtricity Web Services Proxy is a RESTful API server installed separately on a Windows or Linux server. This host-based application can manage hundreds of new and legacy NetApp E-Series storage systems. In general, you should use the proxy for networks with more than 10 storage systems. The proxy can handle numerous requests more efficiently than the embedded API.

The core of the API is available in both implementations.



The following table provides a comparison of the proxy and the embedded version.

Consideration	Proxy	Embedded
Installation	Requires a host system (Linux or Windows). The proxy is available for download at the NetApp Support Site or on DockerHub .	No installation or enablement required.

Consideration	Proxy	Embedded
Security	Minimal security settings by default. Security settings are low so that developers can get started with the API quickly and easily. If desired, you can configure the proxy with the same security profile as the embedded version.	High security settings by default. Security settings are high because the API runs directly on the controllers. For example, it does not allow HTTP access, and it disables all SSL and older TLS encryption protocols for HTTPS.
Central management	Manages all storage systems from one server.	Manages only the controller on which it is embedded.

Unified Manager

The proxy installation package includes Unified Manager, a web-based interface that provides configuration access to newer E-Series and EF-Series storage systems, such as the E2800, E5700, EF300, and EF600.

From Unified Manager, you can perform the following batch operations:

- View the status of multiple storage systems from a central view
- Discover multiple storage systems in your network
- Import settings from one storage system to multiple systems
- Upgrade firmware for multiple storage systems

Compatibility and restrictions

The following compatibility and restrictions apply to using the Web Services Proxy.

Consideration	Compatibility or restriction
HTTP support	The Web Services Proxy allows use of HTTP or HTTPS. (The embedded version of Web Services requires HTTPS for security reasons.)
Storage systems and firmware	The Web Services Proxy can manage all E-Series storage systems, including a mixture of older systems and the latest E2800, EF280, E5700, EF570, EF300, and EF600 series systems.

Consideration	Compatibility or restriction
IP Support	<p>The Web Services Proxy supports either the IPv4 protocol or IPv6 protocol.</p> <div>  <p>The IPv6 protocol might fail when the Web Services Proxy tries to automatically discover the management address from the controller configuration. Possible causes for the failure include problems during IP address forwarding or IPv6 being enabled on the storage systems but not on the server.</p> </div>
NVSRAM file name constraints	<p>The Web Services Proxy uses NVSRAM file names to identify version information accurately. Therefore, you cannot change NVSRAM filenames when they are used with the Web Services Proxy. The Web Services Proxy might not recognize a renamed NVSRAM file as a valid firmware file.</p>
Symbol Web	<p>Symbol Web is a URL in the REST API. It provides access to almost all symbol calls. The symbol function is part of the following URL:</p> <pre>http://host:port/devmgr/storage-system/storage array ID/symbol/symbol function</pre> <div>  <p>Symbol-disabled storage systems are supported through the Web Services Proxy.</p> </div>

API basics

In the Web Services API, HTTP communications involve a request-response cycle.

URL elements in requests

Regardless of the programming language or tool used, each call to the Web Services API has a similar structure, with a URL, HTTP verb, and an Accept header.



All requests include a URL, as in the following example, and contain the elements described in the table.

`https://webservices.name.com:8443/devmgr/v2/storage-systems`

Area	Description
HTTP transport <code>https://</code>	<p>The Web Services Proxy enables the use of HTTP or HTTPS.</p> <p>The embedded Web Services requires HTTPS for security reasons.</p>
Base URL and port <code>webservices.name.com:8443</code>	<p>Each request must be correctly routed to an active instance of Web Services. The FQDN (fully qualified domain name) or the IP address of the instance is required, along with the listening port. By default, Web Services communicates over port 8080 (for HTTP) and port 8443 (for HTTPS).</p> <p>For the Web Services Proxy, both ports can be changed during the proxy installation or in the <code>wsconfig.xml</code> file. Port contention is common on data center hosts running various management applications.</p> <p>For the embedded Web Services, the port on the controller cannot be changed; it defaults to port 8443 for secure connections.</p>
API path <code>devmgr/v2/storage-systems</code>	<p>A request is made to a specific REST resource or endpoint within the Web Services API. Most endpoints are in the form of:</p> <p><code>devmgr/v2/<resource>/[id]</code></p> <p>The API path consists of three parts:</p> <ul style="list-style-type: none">• <code>devmgr</code> (Device Manager) is the namespace of the Web Services API.• <code>v2</code> denotes the version of the API that you are accessing. You can also use <code>utils</code> to access login endpoints.• <code>storage-systems</code> is a category within the documentation.

Supported HTTP verbs

Supported HTTP verbs include GET, POST, and DELETE:

- GET requests are used for read-only requests.
- POST requests are used to create and update objects, and also for read requests that might have security implications.

- DELETE requests are typically used to remove an object from management, remove an object entirely, or to reset the state of the object.



Currently, the Web Services API does not support PUT or PATCH. Instead, you can use POST to provide the typical functionality for these verbs.

Accept headers

When returning a request body, Web Services returns the data in JSON format (unless otherwise specified). Certain clients default to requesting “text/html” or something similar. In these cases, the API responds with an HTTP code 406, denoting that it cannot provide data in this format. As a best practice, you should define the Accept header as “application/json” for any cases in which you expect JSON as the response type. In other cases where a response body is not returned (for example, DELETE), providing the Accept header does not cause any unintended effects.

Responses

When a request is made to the API, a response returns two critical pieces of information:

- HTTP status code — Indicates whether the request was successful.
- Optional response body — Usually provides a JSON body representing the state of the resource or a body providing more details on the nature of a failure.

You must check the status code and the content-type header to determine what the resulting response body looks like. For HTTP status codes 200-203 and 422, Web Services returns a JSON body with the response. For other HTTP status codes, Web Services generally does not return an additional JSON body, either because the specification does not allow it (204) or because the status is self-explanatory. The table lists common HTTP status codes and definitions. It also indicates whether information associated with each HTTP code is returned in a JSON body.

HTTP status code	Description	JSON body
200 OK	Denotes a successful response.	Yes
201 Created	Indicates that an object was created. This code is used in a few rare cases instead of a 200 status.	Yes
202 Accepted	Indicates that the request is accepted for processing as an asynchronous request, but you must make a subsequent request to get the actual result.	Yes
203 Non-Authoritative Information	Similar to a 200 response, but Web Services cannot guarantee that the data is up-to-date (for example, only cached data is available at this time).	Yes

HTTP status code	Description	JSON body
204 No Content	Indicates a successful operation, but there is no response body.	No
400 Bad Request	Indicates that the JSON body provided in the request is not valid.	No
401 Unauthorized	Indicates that an authentication failure has occurred. Either no credentials were provided, or the username or password was invalid.	No
403 Forbidden	An authorization failure, which indicates that the authenticated user does not have permission to access the requested endpoint.	No
404 Not Found	Indicates that the requested resource could not be located. This code is valid for nonexistent APIs or nonexistent resources requested by the identifier.	No
422 Unprocessable Entity	Indicates the request is generally well-formed, but either the input parameters are invalid, or the state of the storage system does not allow Web Services to satisfy the request.	Yes
424 Failed Dependency	Used in the Web Services Proxy to indicate that the requested storage system is currently inaccessible. Therefore, Web Services cannot satisfy the request.	No
429 Too Many Requests	Indicates that a request limit was exceeded and should be retried at a later time.	No

Sample scripts

GitHub contains a repository for the collection and organization of sample scripts illustrating the use of the NetApp SANtricity Web Services API. To access the repository, see [NetApp Webservices samples](#).

Terms and concepts

The following terms apply to the Web Services Proxy.

Term	Definition
API	An Application Programming Interface (API) is a set of protocols and methods that enables developers to communicate with devices. The Web Services API is used to communicate with E-Series storage systems.
ASUP	The AutoSupport (ASUP) feature collects data in a customer support bundle and automatically sends the message file to technical support for remote troubleshooting and problem analysis.
Endpoint	Endpoints are functions that are available through the API. An endpoint includes an HTTP verb, plus the URI path. In Web Services, endpoints can execute such tasks as discovering storage systems and creating volumes.
HTTP Verb	An HTTP verb is a corresponding action for an endpoint, such as retrieving and creating data. In Web Services, HTTP verbs include POST, GET, and DELETE.
JSON	JavaScript Object Notation (JSON) is a structured data format much like XML, which uses a minimal, readable format. Data within Web Services is encoded through JSON.
REST / RESTful	<p>Representational state transfer (REST) is a loose specification that defines an architectural style for an API. Because most REST APIs do not fully adhere to the specification, they are described as “RESTful” or “REST-like.” Generally, a “RESTful” API is agnostic to programming languages and has the following characteristics:</p> <ul style="list-style-type: none"> • HTTP-based, which follows the general semantics of the protocol • Producer and consumer of structured data (JSON, XML, etc.) • Object-oriented (as opposed to operation-oriented) <p>Web Services is a RESTful API that provides access to virtually all the SANtricity management capabilities.</p>
storage system	A storage system is an E-Series array, which includes shelves, controllers, drives, software, and firmware.

Term	Definition
SYMBOL API	SYMBOL is a legacy API for managing E-Series storage systems. The underlying implementation of the Web Services API uses SYMBOL.
Web Services	Web Services is an API that NetApp designed for developers to manage E-Series storage systems. There are two implementations of Web Services: embedded on the controller and a separate proxy that can be installed on Linux or Windows.

Install and configure

Review installation and upgrade requirements

Before installing the Web Services Proxy, review the installation requirements and upgrade considerations.

Installation requirements

You can install and configure the Web Services Proxy on a Windows or Linux host system.

Proxy installation includes the following requirements.

Requirement	Description
Hostname limitations	Be sure that the hostname of the server where you plan to install the Web Services Proxy contains only ASCII letters, numerical digits, and hyphens (-). This requirement is due to a limitation of Java Keytool, which is used in generating a self-signed certificate for the server. If the hostname of your server contains any other characters, such as an underscore (_), the Webserver will fail to start after installation.
Operating systems	<p>You can install the proxy on the following operating systems:</p> <ul style="list-style-type: none"> • Linux • Windows <p>For a complete list of operating systems and firmware compatibility, see the NetApp Interoperability Matrix Tool.</p>

Requirement	Description
Linux: Additional Considerations	Linux Standard Base libraries (init-functions) are required for the Webserver to function properly. You must install the lsb/insserv packages for your operating system. For more information, refer to the "Additional packages required" section of the Readme file.
Multiple instances	You can install only one instance of Web Services Proxy on a server; however, you can install the proxy on multiple servers within your network.
Capacity planning	<p>Web Services Proxy requires adequate space for logging. Make sure that your system meets the following available disk space requirements:</p> <ul style="list-style-type: none"> • Required installation space — 275 MB • Minimum logging space — 200 MB • System memory — 2 GB; heap space is 1 Gb by default <p>You can use a disk-space monitoring tool to verify available disk drive space for persistent storage and logging.</p>
License	The Web Services Proxy is a free, standalone product that does not require a license key. However, applicable copyrights and terms of service apply. If you are installing the proxy in either Graphical or Console mode, you must accept the End User License Agreement (EULA).

Upgrade considerations

If you are upgrading from a previous version, be aware that some items are preserved or removed.

- For the Web Services Proxy, previous configuration settings are preserved. These settings include user passwords, all discovered storage systems, server certificates, trusted certificates, and server runtime configuration.
- For Unified Manager, all SANtricity OS files previously loaded in the repository are removed during the upgrade.

Install or upgrade Web Services Proxy file

Installation involves downloading the file and then installing the proxy package on a Linux or Windows server. You can also upgrade the proxy using these instructions.

Download Web Services Proxy files

You can download the installation file and the readme file from the Software download page of the NetApp Support site.

The download package includes the Web Services Proxy and the Unified Manager interface.

Steps

1. Go to [NetApp Support - Downloads](#).
2. Select **E-Series SANtricity Web Services Proxy**.
3. Follow the instructions to download the file. Make sure you select the correct download package for your server (for example, EXE for Windows; BIN or RPM for Linux).
4. Download the installation file to the server where you want to install the proxy and Unified Manager.

Install on Windows or Linux server

You can install the Web Services Proxy and Unified Manager using one of three modes (Graphical, Console, or Silent), or by using an RPM file (Linux only).

Before you begin

- [Review installation requirements](#).
- Make sure you have downloaded the correct installation file (EXE for Windows; BIN for Linux) to the server where you want to install the proxy and Unified Manager.

Graphical mode install

You can run the installation in Graphical mode for either Windows or Linux. In Graphical mode, the prompts appear in a Windows-style interface.

Steps

1. Access the folder where you downloaded the installation file.
2. Launch the installation for either Windows or Linux, as follows:

- Windows — Double-click the installation file:

```
santricity_webservices-windows_x64-nn.nn.nn.nnnn.exe
```

- Linux — Run the following command: `santricity_webservices-linux_x64-nn.nn.nn.nnnn.bin`

In the above filenames, `nn.nn.nn.nnnn` represents the version number.

The installation process starts and the NetApp SANtricity Web Services Proxy + Unified Manager splash screen appears.

3. Follow the on-screen prompts.

During the installation, you are prompted to enable several features and enter some configuration parameters. If necessary, you can change any of these selections later in the configuration files.



During an upgrade, you are not prompted for configuration parameters.

4. When the Webserver Started message appears, click **OK** to complete the installation.

The Install Complete dialog box appears.

5. Click the check boxes if you want to launch Unified Manager or the interactive API documentation, and then click **Done**.

Console mode install

You can run the installation in Console mode for either Windows or Linux. In Console mode, the prompts appear in the terminal window.

Steps

1. Run the following command: `<install filename> -i console`

In the above command, `<install filename>` represents the name of the proxy installation file you downloaded (for example: `santricity_webservices-windows_x64-nn.nn.nnnn.exe`).



To cancel the installation at any time during the installation process, type `QUIT` at the command prompt.

The installation process starts and the Launching Installer — Introduction message appears.

2. Follow the on-screen prompts.

During the installation, you are prompted to enable several features and enter some configuration parameters. If necessary, you can change any of these selections later in the configuration files.



During an upgrade, you are not prompted for configuration parameters.

3. When the installation is complete, press **Enter** to exit the installer.

Silent mode install

You can run the installation in Silent mode for either Windows or Linux. In Silent mode, no return messages or scripts appear in the terminal window.

Steps

1. Run the following command: `<install filename> -i silent`

In the above command, `<install filename>` represents the name of the proxy installation file you downloaded (for example: `santricity_webservices-windows_x64-nn.nn.nnnn.exe`).

2. Press **Enter**.

The installation process can take several minutes to complete. After successful installation, a command prompt appears in the terminal window.

RPM command install (Linux only)

For Linux systems that are compatible with the RPM package management system, you can install the Web Services Proxy using an optional RPM file.

Steps

1. Download the RPM file to the server where you want to install the proxy and Unified Manager.
2. Open a terminal window.
3. Enter the following command:

```
rpm -u santricity_webservices-nn.nn.nn.nnnn-n.x86_64.rpm
```



In the above command, nn.nn.nn.nnnn represents the version number.

The installation process can take several minutes to complete. After successful installation, a command prompt appears in the terminal window.

Log in to API and Unified Manager

Web Services includes API documentation, which enables you to directly interact with the REST API. It also includes Unified Manager, a browser-based interface for managing multiple E-Series storage systems.

Log in to Web Services API

After you install the Web Services Proxy, you can access the interactive API documentation in a browser.

The API documentation runs with each instance of Web Services, and is also available in a static PDF format from the NetApp Support site. To access the interactive version, you open a browser and enter the URL pointing to where Web Services resides (either a controller for the embedded version or a server for the proxy).



The Web Services API implements the OpenAPI specification (originally called the Swagger specification).

For initial login, you use the "admin" credentials. "Admin" is considered a super administrator with access to all functions and roles.

Steps

1. Open a browser.
2. Enter the URL for the embedded or proxy implementation:
 - Embedded: `https://<controller>:<port>/devmgr/docs/`

In this URL, <controller> is the IP address or FQDN of the controller, and <port> is the management port number of the controller (defaults to 8443).

- Proxy: `http[s]://<server>:<port>/devmgr/docs/`

In this URL, <server> is the IP address or FQDN of the server where the proxy is installed, and <port> is the listening port number (defaults to 8080 for HTTP or 8443 for HTTPS).




If the listening port is already in use, the proxy detects the conflict and prompts you to choose a different listening port.

The API documentation opens in the browser.

3. When the interactive API documentation opens, go to the drop-down menu in the upper right of the page and select **utils**.
4. Click the **Login** category to see the available endpoints.
5. Click the **POST: /login** endpoint, and then click **Try it out**.
6. For first-time login, enter admin for the username and password.
7. Click **Execute**.
8. To access the endpoints for storage management, go to the drop-down menu in the upper right and select **v2**.

The high-level categories for endpoints are displayed. You can navigate the API documentation as described in the table.

Area	Description
Drop-down menu	<p>At the upper right of the page, a drop-down menu provides options for switching between version 2 of the API documentation (V2), the SYMBol interface (SYMBol V2), and API utilities (utils) for logging in.</p> <div>  <p>Because version 1 of the API documentation was a prerelease and not generally available, V1 is not included in the drop-down menu.</p> </div>
Categories	The API documentation is organized by high-level categories (for example: Administration, Configuration). Click on a category to see the related endpoints.
Endpoints	Select an endpoint to see its URL paths, required parameters, response bodies, and status codes that the URLs are likely to return.
Try It Out	<p>Interact with the endpoint directly by clicking Try It Out. This button is provided in each of the expanded views for endpoints.</p> <p>When you click the button, fields appear for entering parameters (if applicable). You can then enter values and click Execute.</p> <p>The interactive documentation uses JavaScript to make the request directly to the API; it is not a test request.</p>

Log in to Unified Manager

After you install the Web Services Proxy, you can access Unified Manager to manage multiple storage systems in a web-based interface.

To access Unified Manager, you open a browser and enter the URL pointing to where the proxy is installed. The following browsers and versions are supported.

Browser	Minimum version
Google Chrome	79
Microsoft Internet Explorer	11
Microsoft Edge	79
Mozilla Firefox	70
Safari	12

Steps

1. Open a browser and enter the following URL:

```
http[s]://<server>:<port>/um
```

In this URL, <server> represents the IP address or FQDN of the server where the Web Services Proxy is installed, and <port> represents the listening port number (defaults to 8080 for HTTP or 8443 for HTTPS).

The Unified Manager login page opens.

2. For first-time login, enter `admin` for the user name, and then set and confirm a password for the admin user.

The password can include up to 30 characters. For further information about users and passwords, see the Access Management section of the Unified Manager online help.

Configure Web Services Proxy

You can modify the Web Services Proxy settings to meet the unique operating and performance requirements for your environment.

Stop or restart the Webserver

The Webserver service is started during installation and runs in the background. During some configuration tasks, you might need to stop or restart the Webserver service.

Steps

1. Do one of the following:
 - For Windows, go to the **Start** menu, select **Administrative Tools** > **Services**, locate **NetApp SANtricity Web Services** and then select either **Stop** or **Restart**.
 - For Linux, choose the method of stopping and restarting the Webserver for your operating system version. During the installation, a popup dialog indicated what daemon started. For example:

```
web_services_proxy webserver installed and started. You can interact with it
```

```
using systemctl start|stop|restart|status web_services_proxy.service
```

The most common method for interacting with the service is by using `systemctl` commands.

Resolve port conflicts

If the Web Services Proxy is running while another application is available at the defined address or port, you can resolve the port conflict in the `wsconfig.xml` file.

Steps

1. Open the `wsconfig.xml` file, located at:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy`
2. Add the following line to the `wsconfig.xml` file, in which *n* is the port number:

```
<sslport clientauth="request">*n*</sslport>
<port>n</port>
```

The following table shows the attributes that control HTTP ports and HTTPS ports.

Name	Description	Parent Node	Attributes	Required
config	The root node for the config	Null	Version - The version of the config schema is currently 1.0.	Yes
sslport	The TCP port to listen for SSL requests. Defaults to 8443.	config	Clientauth	No
port	The TCP port to listen for HTTP request, defaults to 8080.	config	-	No

3. Save and close the file.
4. Restart the Webserver service so the change takes effect.

Configure load-balancing and/or high-availability

To use the Web Services Proxy in a highly-available (HA) configuration, you can configure load balancing. In an HA configuration, typically either a single node receives all requests while the others are on stand-by, or requests are load-balanced across all nodes.

The Web Services Proxy can exist in a highly-available (HA) environment, with most APIs operating correctly regardless of the recipient of the request. Metadata tags and folders are two exceptions, because tags and folders are stored in a local database and are not shared between Web Services Proxy instances.

However, there are some known timing issues that occur in a small percentage of requests. Specifically, one instance of the proxy can have newer data faster than a second instance for a small window. The Web Services Proxy includes a special configuration that removes this timing issue. This option is not enabled by default, because it increases the amount of time it takes to service requests (for data consistency). To enable this option, you must add a property to an .INI file (for Windows) or an .SH file (for Linux).

Steps

1. Do one of the following:
 - Windows: Open the appserver64.ini file, and then add the `Dload-balance.enabled=true` property.

For example: `vmarg.7=-Dload-balance.enabled=true`

- Linux: Open the webserver.sh file, and then add the `Dload-balance.enabled=true` property.

For example: `DEBUG_START_OPTIONS="-Dload-balance.enabled=true"`

2. Save your changes.
3. Restart the Webserver service so the change takes effect.

Disable SYMBol HTTPS

You can disable SYMBol commands (default setting) and send commands over a remote procedure call (RPC). This setting can be changed in the wsconfig.xml file.

By default, the Web Services Proxy sends SYMBol commands over HTTPS for all E2800 series and E5700 series storage systems running SANtricity OS versions 08.40 or later. SYMBol commands sent over HTTPS are authenticated to the storage system. If needed, you can disable HTTPS SYMBol support and send commands over RPC. Whenever SYMBol over RPC is configured, all passive commands to the storage system are enabled without authentication.



When SYMBol over RPC is used, the Web Services Proxy cannot connect to systems with the SYMBol management port disabled.

Steps

1. Open the wsconfig.xml file, located at:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy`
2. In the `devicemgt.symbolclientstrategy` entry, replace the `httpsPreferred` value with `rpcOnly`.

For example:

```
<env key="devicemgt.symbolclientstrategy">rpcOnly</env>
```

3. Save the file.

Configure cross-origin resource sharing

You can configure cross-origin resource sharing (CORS), which is a mechanism that uses additional HTTP headers to provide a web application running at one origin to have permission to access selected resources from a server at a different origin.

CORS is handled by the cors.cfg file located in the working directory. The CORS configuration is open by default, so cross domain access is not restricted.

If no configuration file is present, CORS is open. But if the cors.cfg file is present, then it is used. If the cors.cfg file is empty, you cannot make a CORS request.

Steps

1. Open the cors.cfg file, which is located in the working directory.
2. Add the desired lines to the file.

Each line in the CORS configuration file is a regular expression pattern to match. The origin header must match a line in the cors.cfg file. If any line pattern matches the origin header, the request is allowed. The complete origin is compared, not just the host element.

3. Save the file.

Requests are matched on the host and according to protocol, such as the following:

- Match localhost with any protocol — `*localhost*`
- Match localhost for HTTPS only — `https://localhost*`

Uninstall Web Services Proxy

To remove the Web Services Proxy and Unified Manager, you can use any mode (Graphical, Console, Silent, or RPM file), regardless of what method you used to install the proxy.

Graphical mode uninstall

You can run the uninstall in Graphical mode for either Windows or Linux. In Graphical mode, the prompts appear in a Windows-style interface.

Steps

1. Launch the uninstall for either Windows or Linux, as follows:
 - Windows — Go to the directory that contains the `uninstall_web_services_proxy` uninstall file. The default directory is at the following location: `C:/Program Files/NetApp/SANtricity Web Services Proxy/`. Double-click `uninstall_web_services_proxy.exe`.



Alternatively, you can go to **Control Panel > Programs > Uninstall a program**, and then select "NetApp SANtricity Web Services Proxy."

- Linux — Go to the directory that contains the Web Services Proxy uninstall file. The default directory is at the following location:
`/opt/netapp/santricity_web_services_proxy/uninstall_web_services_proxy`
2. Run the following command:

```
uninstall_web_services_proxy -i gui
```

The SANtricity Web Services Proxy splash screen appears.

3. From the Uninstall dialog box, click **Uninstall**.

The Uninstaller progress bar appears and shows the progress.

4. When the Uninstall Complete message appears, click **Done**.

Console mode uninstall

You can run the uninstall in Console mode for either Windows or Linux. In Console mode, the prompts appear in the terminal window.

Steps

1. Go to the `uninstall_web_services_proxy` directory.
2. Run the following command:

```
uninstall_web_services_proxy -i console
```

The uninstall process starts.

3. When the uninstall is complete, press **Enter** to exit the installer.

Silent mode uninstall

You can run the uninstall in Silent mode for either Windows or Linux. In Silent mode, no return messages or scripts appear in the terminal window.

Steps

1. Go to the `uninstall_web_services_proxy` directory.
2. Run the following command:

```
uninstall_web_services_proxy -i silent
```

The uninstall process runs, but no return messages or scripts appear in the terminal window. After Web Services Proxy is successfully uninstalled, a command prompt appears in the terminal window.

RPM command uninstall (Linux only)

You can use an RPM command to uninstall the Web Services Proxy from a Linux system.

Steps

1. Open a terminal window.
2. Enter the following command line:

```
rpm -e santricity_webservices
```



The uninstall process might leave files that were not part of the original installation. Manually delete these files to remove Web Services Proxy completely.

Manage user access in Web Services Proxy

You can manage user access to the Web Services API and Unified Manager for security purposes.

Overview of access management

Access management includes role-based logins, password encryption, basic authentication, and LDAP integration.

Role-based access

Role-based access control (RBAC) associates predefined users with roles. Each role grants permissions to a specific level of functionality.

The following table describes each role.

Role	Description
security.admin	SSL and certificate management.
storage.admin	Full read/write access to storage system configuration.
storage.monitor	Read-only access to view storage system data.
support.admin	Access to all hardware resources on storage systems and support operations such as AutoSupport (ASUP) retrieval.

Default user accounts are defined in the `users.properties` file. You can change user accounts by directly modifying the `users.properties` file or by using the Access Management functions in Unified Manager.

The following table lists the user logins available for the Web Services Proxy.

Predefined user login	Description
admin	A super administrator who has access to all functions and includes all roles. For Unified Manager, you must set the password on first-time login.
storage	The administrator responsible for all storage provisioning. This user includes the following roles: <code>storage.admin</code> , <code>support.admin</code> , and <code>storage.monitor</code> . This account is disabled until a password is set.
security	The user responsible for security configuration. This user includes the following roles: <code>security.admin</code> and <code>storage.monitor</code> . This account is disabled until a password is set.
support	The user responsible for hardware resources, failure data, and firmware upgrades. This user includes the following roles: <code>support.admin</code> and <code>storage.monitor</code> . This account is disabled until a password is set.

Predefined user login	Description
monitor	A user with read-only access to the system. This user includes only the storage.monitor role. This account is disabled until a password is set.
rw (legacy for older arrays)	The rw (read/write) user includes the following roles: storage.admin, support.admin, and storage.monitor. This account is disabled until a password is set.
ro (legacy for older arrays)	The ro (read only) user includes only the storage.monitor role. This account is disabled until a password is set.

Password encryption

For each password, you can apply an additional encryption process using the existing SHA256 password encoding. This additional encryption process applies a random set of bytes to each password (salt) for each SHA256 hash encryption. Salted SHA256 encryption is applied to all newly created passwords.



Prior to the Web Services Proxy 3.0 release, passwords were encrypted through SHA256 hashing only. Any existing SHA256 hash-only encrypted passwords retain this encoding and are still valid under the users.properties file. However, SHA256 hash-only encrypted passwords are not as secure as those passwords with salted SHA256 encryption.

Basic authentication

By default, basic authentication is enabled, which means the server returns a basic authentication challenge. This setting can be changed in the wsconfig.xml file.

LDAP

Lightweight Directory Access Protocol (LDAP), an application protocol for accessing and maintaining distributed directory information services, is enabled for the Web Services Proxy. LDAP integration allows for user authentication and mapping of roles to groups.

For information on configuring LDAP functionality, refer to configuration options in the Unified Manager interface or in the LDAP section of the interactive API documentation.

Configure user access

You can manage user access by applying additional encryption to passwords, setting basic authentication, and defining role-based access.

Apply additional encryption to passwords

For the highest level of security, you can apply additional encryption to passwords using the existing SHA256 password encoding.

This additional encryption process applies a random set of bytes to each password (salt) for each SHA256 hash encryption. Salted SHA256 encryption is applied to all newly created passwords.

Steps

1. Open the `users.properties` file, located at:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy\data\config`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy/data/config`
2. Re-enter the encrypted password as plain text.
3. Run the `securepasswd` command line utility to re-encrypt the password or simply restart the Web Services Proxy. This utility is installed in the root install directory for the Web Services Proxy.



Alternatively, you can salt and hash local user passwords whenever password edits are performed through the Unified Manager.

Configure basic authentication

By default basic authentication is enabled, which means the server returns a basic authentication challenge. If desired, you can change that setting in the `wsconfig.xml` file.

1. Open the `wsconfig.xml` file, located at:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy`
2. Modify the following line in the file by specifying `false` (not enabled) or `true` (enabled).

For example: `<env key="enable-basic-auth">true</env>`

3. Save the file.
4. Restart the Webserver service so the change takes effect.

Configure role-based access

To limit user access to specific functions, you can modify which roles are specified for each user account.

The Web Services Proxy includes role-based access control (RBAC), in which roles are associated with predefined users. Each role grants permissions to a specific level of functionality. You can change the roles assigned to user accounts by directly modifying the `users.properties` file.



You can also change user accounts by using Access Management in Unified Manager. For more information, see the online help available with Unified Manager.

Steps

1. Open the `users.properties` file, located in:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy\data\config`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy/data/config`
2. Locate the line for the user account you want to modify (storage, security, monitor, support, rw, or ro).



Do not modify the admin user. This is a super user with access to all functions.

3. Add or remove the specified roles, as desired.

Roles include:

- security.admin — SSL and certificate management.
- storage.admin — Full read/write access to storage system configuration.
- storage.monitor — Read-only access to view storage system data.
- support.admin — Access to all hardware resources on storage systems and support operations such as AutoSupport (ASUP) retrieval.



The storage.monitor role is required for all users, including the administrator.

4. Save the file.

Manage security and certificates in Web Services Proxy

For security in the Web Services Proxy, you can specify an SSL port designation and you can manage certificates. Certificates identify website owners for secure connections between clients and servers.

Enable SSL

The Web Services Proxy uses Secure Sockets Layer (SSL) for security, which is enabled during installation. You can change the SSL port designation in the wsconfig.xml file.

Steps

1. Open the wsconfig.xml file, located at:
 - (Windows) — C:\Program Files\NetApp\SANtricity Web Services Proxy
 - (Linux) — /opt/netapp/santricity_web_services_proxy
2. Add or change the SSL port number, similar to the following example:

```
<sslport clientauth="request">8443</sslport>
```

Result

When the server is started with SSL configured, the server looks for the keystore and truststore files.

- If the server does not find a keystore, the server uses the IP address of the first detected non-loopback IPv4 address to generate a keystore and then add a self-signed certificate to the keystore.
- If the server does not find a truststore, or the truststore is not specified, the server uses the keystore as the truststore.

Bypass certificate validation

To support secure connections, the Web Services Proxy validates the storage systems' certificates against its own trusted certificates. If needed, you can specify that the proxy bypass that validation before connecting to the storage systems.

Before you begin

- All storage system connections must be secure.

Steps

1. Open the wsconfig.xml file, located at:
 - (Windows) — C:\Program Files\NetApp\SANtricity Web Services Proxy
 - (Linux) — /opt/netapp/santricity_web_services_proxy
2. Enter true in the trust.all.arrays entry, as shown in the example:

```
<env key="trust.all.arrays">true</env>
```

3. Save the file.

Generate and import a host management certificate

Certificates identify website owners for secure connections between clients and servers. To generate and import Certificate Authority (CA) certificates for the host system where the Web Services Proxy is installed, you can use API endpoints.

To manage certificates for the host system, you perform the following tasks using the API:

- Create a certificate signing request (CSR) for the host system.
- Send the CSR file to a CA, and then wait for them to send you the certificate files.
- Import the signed certificates to the host system.



You can also manage certificates in the Unified Manager interface. For more information, see the online help available in Unified Manager.

Steps

1. Log in to the [interactive API documentation](#).
2. Go to the drop-down menu in the upper right and then select **v2**.
3. Expand the **Administration** link and scroll down to the **/certificates** endpoints.
4. Generate the CSR file:
 - a. Select **POST:/certificates**, and then select **Try it out**.

The web server regenerates a self-signed certificate. You can then enter information in the fields to define the common name, organization, organization unit, alternate ID, and other information used to generate the CSR.

- b. Add the required information in the **Example values** pane to generate a valid CA certificate, and then execute the commands.



Do not call **POST:/certificates** or **POST:/certificates/reset** again, or you must regenerate the CSR. When you call **POST:/certificates** or **POST:/certificates/reset**, you are generating a new self-signed certificate with a new private key. If you send a CSR that was generated before the last reset of the private key on the server, the new security certificate does not work. You must generate a new CSR and request a new CA certificate.

- c. Execute the **GET:/certificates/server** endpoint to confirm that the current certificate status is the self-signed certificate with the information added from the **POST:/certificates** command.

The server certificate (denoted by the alias `jetty`) is still self-signed at this point.

- d. Expand the **POST:/certificates/export** endpoint, select **Try it out**, enter a file name for the CSR file, and then click **Execute**.
5. Copy and paste the `fileUrl` into a new browser tab to download the CSR file, and then send the CSR file to a valid CA to request a new web server certificate chain.
6. When the CA issues a new certificate chain, use a certificate manager tool to break out the root, intermediate, and web server certificates, and then import them to the Web Services Proxy server:
 - a. Expand the **POST:/sslconfig/server** endpoint and select **Try it out**.
 - b. Enter a name for the CA root certificate in the **alias** field.
 - c. Select **false** in the **replaceMainServerCertificate** field.
 - d. Browse to and select the new CA root certificate.
 - e. Click **Execute**.
 - f. Confirm that the certificate upload was successful.
 - g. Repeat the CA certificate upload procedure for the CA intermediate certificate.
 - h. Repeat the certificate upload procedure for the new web server security certificate file, except in this step, select **true** on the **replaceMainServerCertificate** drop-down.
 - i. Confirm that the web server security certificate import was successful.
 - j. To confirm that the new root, intermediate, and web server certificates are available in the keystore, run **GET:/certificates/server**.
7. Select and expand the **POST:/certificates/reload** endpoint, and then select **Try it out**. When prompted, whether you want to restart both controllers or not, select **false**. ("True" applies only in the case of dual array controllers.) Click **Execute**.

The **/certificates/reload** endpoint usually returns a successful http 202 response. However, the reload of the web server truststore and keystore certificates does create a race condition between the API process and the web server certificate reload process. In rare cases, the web server certificate reload can beat the API processing. In this case, the reload appears to fail even though it completed successfully. If this occurs, continue to the next step anyway. If the reload actually failed, the next step also fails.

8. Close the current browser session to the Web Services Proxy, open a new browser session, and confirm that a new secure browser connection to the Web Services Proxy can be established.

By using an incognito or in-private browsing session, you can open a connection to the server without using any saved data from previous browsing sessions.

Manage storage systems using Web Services Proxy

To manage storage systems in the network, you must first discover them and then add them to the management list.

Discover storage systems

You can set automatic discovery or manually discover storage systems.

Automatically discover storage systems

You can specify that storage systems are automatically discovered in the network by modifying the settings in the `wsconfig.xml` file. By default, IPv6 automatic discovery is disabled and IPv4 is enabled.

You only need to provide one management IP or DNS address to add a storage system. The server automatically discovers all management paths when the paths are either not configured or the paths are configured and rotatable.



If you attempt to use an IPv6 protocol to automatically discover storage systems from the controller configuration after an initial connection has been made, the process might fail. Possible causes for the failure include problems during IP address forwarding or IPv6 being enabled on the storage systems, but not being enabled on the server.

Before you begin

Before enabling IPv6 discovery settings, verify that your infrastructure supports IPv6 connectivity to the storage systems to mitigate any connection issues.

Steps

1. Open the `wsconfig.xml` file, located at:
 - (Windows) — `C:\Program Files\NetApp\SANtricity Web Services Proxy`
 - (Linux) — `/opt/netapp/santricity_web_services_proxy`
2. In the autodiscover strings, change settings from `true` to `false`, as desired. See the following example.

```
<env key="autodiscover.ipv6.enable">true</env>
```



When the paths are configured, but not configured so that the server can route to the addresses, intermittent connection errors happen. If you cannot set the IP addresses to be routable from the host, turn off auto discovery (change the settings to `false`).

3. Save the file.

Discover and add storage systems using API endpoints

You can use API endpoints to discover and add storage systems to the managed list. This procedure creates a management connection between the storage system and the API.



This task describes how to discover and add storage systems using the REST API, so you can manage these systems in the interactive API documentation. However, you might want to manage storage systems in the Unified Manager instead, which provides an easy-to-use interface. For more information, see the online help available with Unified Manager.

Before you begin

For storage systems with SANtricity versions 11.30 and later, the legacy management interface for SYMBol must be enabled in the SANtricity System Manager interface. Otherwise, the Discovery endpoints fail. You can find this setting by opening System Manager, and then going to **Settings > System > Additional Settings > Change Management Interface**.

Steps

1. Log in to the [interactive API documentation](#).
2. Discover storage systems, as follows:
 - a. From the API documentation, make sure **V2** is selected in the drop-down, and then expand the **Storage-Systems** category.
 - b. Click the **POST: /discovery** endpoint, and then click **Try it out**.
 - c. Enter the parameters as described in the table.

startIP endIP	Replace string with the starting and ending IP address range for one or more storage systems in the network.
useAgents	Set this value to either: <ul style="list-style-type: none"> • true = Use in-band agents for the network scan. • false = Do not use in-band agents for the network scan.
connectionTimeout	Enter the seconds allowed for the scan before the connection times out.
maxPortsToUse	Enter a maximum number of ports used for the network scan.

- d. Click **Execute**.



API actions execute without user prompts.

The discovery process runs in the background.

- e. Make sure the code returns a 202.
 - f. Under **Response Body**, locate the value returned for the requestId. You need the Request ID to view the results in the next step.
3. View discovery results, as follows:
 - a. Click the **GET: /discovery** endpoint, and then click **Try it out**.
 - b. Enter the Request ID from the previous step. If you leave the **Request ID** blank, the endpoint defaults to the last request ID executed.
 - c. Click **Execute**.
 - d. Make sure the code returns 200.
 - e. In the response body, locate your Request ID and the strings for storageSystems. The strings look similar to the following example:

```

"storageSystems": [
  {
    "serialNumber": "123456789",
    "wwn": "000A011000AF00000000000001A0C000E",
    "label": "EF570_Array",
    "firmware": "08.41.10.01",
    "nvsram": "N5700-841834-001",
    "ipAddresses": [
      "10.xxx.xx.213",
      "10.xxx.xx.214"
    ],
  },

```

f. Write down the values for wwn, label, and ipAddresses. You need them for the next step.

4. Add storage systems, as follows:

- a. Click the **POST: /storage-system** endpoint, and then click **Try it out**.
- b. Enter the parameters as described in the table.

id	Enter a unique name for this storage system. You can enter the label (displayed in the response for GET: /discovery), but the name can be any string you choose. If you do not provide a value for this field, Web Services automatically assigns a unique identifier.
controllerAddresses	Enter the IP addresses displayed in the response for GET: /discovery. For dual controllers, separate the IP addresses with a comma. For example: "IP address 1", "IP address 2"
validate	Enter <code>true</code> , so you can receive confirmation that Web Services can connect to the storage system.
password	Enter the administrative password for the storage system.
wwn	Enter the WWN of the storage system (displayed in the response for GET: /discovery).

- c. Remove all strings after "enableTrace": `true`, so that the entire string set is similar to the following example:

```
{
  "id": "EF570_Array",
  "controllerAddresses": [
    "Controller-A-Mgmt-IP", "Controller-B-Mgmt_IP"
  ],
  "validate": true,
  "password": "array-admin-password",
  "wwn": "000A011000AF00000000000001A0C000E",
  "enableTrace": true
}
```

d. Click **Execute**.

e. Make sure the code response is 201, which indicates that the endpoint executed successfully.

The **Post: /storage-systems** endpoint is queued. You can view the results using the **GET: /storage-systems** endpoint in the next step.

5. Confirm the list addition, as follows:

a. Click the **GET: /storage-system** endpoint.

No parameters are required.

b. Click **Execute**.

c. Make sure that the code response is 200, which indicates that the endpoint executed successfully.

d. In the response body, look for the storage system details. The returned values indicate that it was successfully added to the list of managed arrays, similar to the following example:

```
[
  {
    "id": "EF570_Array",
    "name": "EF570_Array",
    "wwn": "000A011000AF00000000000001A0C000E",
    "passwordStatus": "valid",
    "passwordSet": true,
    "status": "optimal",
    "ip1": "10.xxx.xx.213",
    "ip2": "10.xxx.xx.214",
    "managementPaths": [
      "10.xxx.xx.213",
      "10.xxx.xx.214"
    ]
  }
]
```

Scale up the number of managed storage systems

By default, the API can manage up to 100 storage systems. If you need to manage more, you must bump the memory requirements for the server.

The server is set to use 512 MB of memory. For every 100 extra storage systems in your network, add 250 MB to that number. Do not add more memory than what you physically have. Allow enough extra for your operating system and other applications.



The default cache size is 8,192 events. The approximate data usage for the MEL events cache is 1MB for each 8,192 events. Therefore, by retaining the defaults, cache usage should be approximately 1MB for a storage system.



In addition to memory, the proxy uses network ports for each storage system. Linux and Windows consider network ports as file handles. As a security measure, most operating systems limit the number of open file handles that a process or a user can have open at one time. Especially in Linux environments, where open TCP connections are considered to be file handles, the Web Services Proxy can easily exceed this limit. Because the fix is system dependent, you should refer to your operating system's documentation for how to raise this value.

Steps

1. Do one of the following:
 - On Windows, go to the `appserver64.init` file. Locate the line, `vmarg.3=-Xmx512M`
 - On Linux, go to the `webserver.sh` file. Locate the line, `JAVA_OPTIONS="-Xmx512M"`
2. To increase the memory, replace 512 with the desired memory in MB.
3. Save the file.

Manage automatic polling for Web Services Proxy statistics

You can configure automatic polling for all disk and volume statistics on discovered storage systems.

Overview of statistics

Statistics provide information about the data collection rates and performance of the storage systems.

The Web Services Proxy provides access to the following types of statistics:

- Raw statistics — Total counters for data points at the time of data collection. Raw statistics can be used for total read operations or total write operations.
- Analyzed statistics — Calculated information for an interval. Examples of analyzed statistics are read input/output operations (IOPs) per second or write throughput.

Raw statistics are linear, typically requiring at least two collected data points to derive usable data from them. The analyzed statistics are a derivation of the raw statistics, which provide important metrics. Many values that can be derived from the raw statistics are shown in a usable, point-in-time format in the analyzed statistics for your convenience.

You can retrieve raw statistics regardless of whether the automatic polling is enabled or not. You can add the

`usecache=true` query string to the end of the URL to retrieve cached statistics from the last poll. Using cached results greatly increases the performance of statistics retrieval. However, multiple calls at a rate equal to or less than the configured polling interval cache retrieves the same data.

Statistics functionality

The Web Services Proxy provides API endpoints that enable the retrieval of raw and analyzed controller and interface statistics from supported hardware models and software versions.

Raw Statistics APIs

- `/storage-systems/{system-id}/controller-statistics`
- `/storage-systems/{system-id}/drive-statistics/{optional list of disk ids}`
- `/storage-systems/{system-id}/interface-statistics/{optional list of interface ids}`
- `/storage-systems/{system-id}/volume-statistics/{optional list of volume ids}`

Analyzed Statistics APIs

- `/storage-systems/{id}/analysed-controller-statistics/`
- `/storage-systems/{id}/analysed-drive-statistics/{optional list of disk ids}`
- `/storage-systems/{id}/analysed-interface-statistics/{optional list of interface ids}`
- `/storage-systems/{id}/analysed-volume-statistics/{optional list of volume ids}`

These URLs retrieve analyzed statistics from the last poll and are only available when polling is enabled. These URLs include the following input-output data:

- Operations per second
- Throughput in megabytes per second
- Response times in milliseconds

The calculations are based on the differences between statistical polling iterations, which are the most common measures of storage performance. These statistics are preferable to unanalyzed statistics.



When the system starts, there is no previous statistics collection to use to calculate the various metrics, so analyzed statistics require at least one polling cycle after startup to return data. In addition, if the cumulative counters are reset, the next polling cycle will have unpredictable numbers for the data.

Configure polling intervals

To configure polling intervals, you modify the `wsconfig.xml` file to specify a polling interval in seconds.



Because the statistics are cached in memory, you might see an increase of about 1.5 MB of memory-use for each storage system.

Before you begin

- The storage systems must be discovered by the proxy.

Steps

1. Open the wsconfig.xml file, located at:
 - (Windows) — C:\Program Files\NetApp\SANtricity Web Services Proxy
 - (Linux) — /opt/netapp/santricity_web_services_proxy
2. Add the following line inside the `<env-entries>` tag, in which `n` is the number of seconds for the interval between polling requests:

```
<env key="stats.poll.interval">n</env>
```

For example, if 60 is entered, polling starts at 60-second intervals. That is, the system requests polling to start 60 seconds after the prior polling period was completed (regardless of the duration of the prior polling period). All statistics are time-stamped with the exact time they were retrieved. The system uses the time stamp or time difference on which to base the 60-second calculation.

3. Save the file.

Manage AutoSupport using Web Services Proxy

You can configure AutoSupport (ASUP), which collects data and then automatically sends that data to technical support for remote troubleshooting and problem analysis.

Overview of AutoSupport (ASUP)

The AutoSupport (ASUP) feature automatically transmits messages to NetApp based on manual and schedule-based criteria.

Each AutoSupport message is a collection of log files, configuration data, state data, and performance metrics. By default, AutoSupport transmits the files listed in the following table to the NetApp Support team once each week.

File Name	Description
x-headers-data.txt	A .txt file containing the X-header information.
manifest.xml	An .xml file detailing the contents of the message.
arraydata.xml	An .xml file containing the list of client persisted data.
appserver-config.txt	A .txt file containing application server configuration data.
wsconfig.txt	A .txt file containing the web service configuration data.

File Name	Description
host-info.txt	A .txt file containing information about the host environment.
server-logs.7z	A .7z file containing every available webserver log file.
client-info.txt	A .txt file with arbitrary key/value pairs for application-specific counters such as method and webpage hits.
webservices-profile.json	<p>These files contain Webservices profile data and Jersey monitoring statistical data. By default, Jersey monitoring statistics are enabled. You can enable and disable them in the wsconfig.xml file, as follows:</p> <ul style="list-style-type: none"> • Enable: <code><env key="enable.jersey.statistics">true</env></code> • Disable: <code><env key="enable.jersey.statistics">false</env></code>

Configure AutoSupport

AutoSupport is enabled by default at installation; however, you can change that setting or modify the delivery types.

Enable or disable AutoSupport

The AutoSupport feature is enabled or disabled during the initial installation of the Web Services Proxy, but you can change that setting in the ASUPConfig file.

You can enable or disable AutoSupport through the ASUPConfig.xml file, as described in the steps below. Alternatively, you can enable or disable this feature through the API using **Configuration** and **POST/asup**, and then entering "true" or "false."

1. Open the ASUPConfig.xml file in the working directory.
2. Locate the lines for `<asupdata enable="Boolean_value" timestamp="timestamp">`
3. Enter `true` (enable) or `false` (disable). For example:

```
<asupdata enabled="false" timestamp="0">
```



The timestamp entry is superfluous.

4. Save the file.

Configure AutoSupport delivery method

You can configure the AutoSupport feature to use HTTPS, HTTP, or SMTP delivery methods. HTTPS is the default delivery method.

1. Access the ASUPConfig.xml file in the working directory.
2. In the string, `<delivery type="n">`, enter 1, 2, or 3 as described in the table:

Value	Description
1	HTTPS (default) <code><delivery type="1"></code>
2	HTTP <code><delivery type="2"></code>
3	SMTP — To properly configure the AutoSupport delivery type to SMTP, you must include the SMTP mail server address, along with the sender and recipient user emails, similar to the following example: <pre><delivery type="3"> <smtp> <mailserver>smtp.example.com</mailserver> <sender>user@example.com</sender> <replyto>user@example.com</replyto> </smtp> </delivery></pre>

Remote volume mirroring

Remote Storage Volumes overview

Use the SANtricity® Remote Storage Volumes feature to import data from a remote storage device directly to a local E-Series volume. This feature helps streamline the process for equipment upgrades and provides data migration capabilities to move data from non-E-Series devices to E-Series systems.

Configuration overview

The Remote Storage Volumes feature is available with SANtricity System Manager on selected submodel IDs. To use this feature, you must configure a remote storage system and an E-Series storage system to

communicate with each other.

Use the following workflow:

1. [Review requirements and restrictions.](#)
2. [Configure hardware.](#)
3. [Import remote storage.](#)

Find more information

- Online help, available in the System Manager user interface or in the [SANtricity software doc site](#).
- For additional technical information on the Remote Storage Volumes feature, see the [Remote Storage Volumes Technical Report](#).

Requirements and restrictions for remote storage

Before configuring the Remote Storage Volumes feature, review the following requirements and restrictions.

Hardware requirements

Supported protocols

For the initial release of the Remote Storage Volumes feature, support is only available for iSCSI and IPv4 protocols.

Refer to the [NetApp Interoperability Matrix Tool](#) for up-to-date support and configuration information between the host and E-Series (destination) array used for the Remote Storage Volumes feature.

Storage system requirements

The E-Series storage system must include:

- Two controllers (duplex mode)
- iSCSI connections for both E-Series controllers to communicate with the remote storage system through one or more iSCSI connections
- SANtricity OS 11.71 or greater
- Remote Storage feature enabled in the Submodel ID (SMID)

The remote system can be either an E-Series storage system or a system from another vendor. It must include iSCSI-capable interfaces.

Volume requirements

Volumes used for imports must meet the requirements for size, status, and other criteria.

Remote storage volume

The source volume of an import is called a "remote storage volume." This volume must meet the following criteria:

- Cannot be part of another import
- Must have an online status

After the import begins, the controller firmware creates a remote storage volume in the background. Due to that background process, the remote storage volume is not manageable in System Manager and can only be used for the import operation.

After it is created, the remote storage volume is treated like any other standard volume on the E-Series system with the following exceptions:

- Can be used as proxies to the remote storage device.
- Cannot be used as candidates for other volume copies or snapshots.
- Cannot have the Data Assurance setting changed while the import is in progress.
- Cannot be mapped to any hosts, because they are reserved strictly for the import operation.

Each remote storage volume is associated with only one remote storage object; however, one remote storage object can be associated with multiple remote storage volumes. The remote storage volume is uniquely identified using a combination of the following:

- Remote storage object identifier
- Remote storage device LUN number

Target volume candidates

The target volume is the destination volume on the local E-Series system.

The destination volume must meet the following criteria:

- Must be a RAID/DDP volume.
- Must have a capacity that is equal to or larger than the remote storage volume.
- Must have a block size that is the same as the remote storage volume.
- Must have a valid state (optimal).
- Cannot have any of the following relationships: volume copy, snapshot copies, asynchronous or synchronous mirroring.
- Cannot be undergoing any reconfiguration operations: Dynamic Volume Expansion, Dynamic Capacity Expansion, Dynamic Segment Size, Dynamic RAID Migration, Dynamic Capacity Reduction, or Defragmentation.
- Cannot be mapped to a host before the import starts (however, it can be mapped after import starts).
- Cannot have Flash Read Cached (FRC) enabled.

System Manager automatically checks these requirements as part of the Import Remote Storage wizard. Only volumes that meet all the requirements are displayed for destination volume selection.

Restrictions

The Remote Storage feature has the following restrictions:

- Mirroring must be disabled.
- Destination volume on the E-Series system must not have snapshots.

- Destination volume on the E-Series system must not be mapped to any hosts before the import is started.
- Destination volume on the E-Series system must have resource-provisioning disabled.
- Direct mappings of the remote storage volume to a host or multiple hosts are not supported.
- Web Services Proxy is not supported.
- iSCSI CHAP secrets are not supported.
- SMcli is not supported.
- VMware Datastore is not supported.
- Only one storage system in the relationship/import pair can be upgraded at a time when there is an import pair present.

Preparation for production imports

You should perform a test or "dry run" import before production imports to verify proper storage and fabric configuration.

Many variables can impact the import operation and completion time. To ensure a production import is successful and to get a duration estimate, you can use these test imports to ensure all connections are working as expected and the import operation is completing in an appropriate amount of time. You can then make adjustments to achieve the desired results before the production import is initiated.

Configure hardware for Remote Storage Volumes

The E-Series storage system must be configured to communicate with the remote storage system through the supported iSCSI protocol.

Configure remote storage device and E-Series array

Before proceeding to the SANtricity System Manager to configure the Remote Storage Volumes feature, do the following:

1. Manually establish a cabled connection between the E-Series system and the remote storage system such that the two systems can be configured to communicate via iSCSI.
2. Configure the iSCSI ports such that the E-Series system and the remote storage system can communicate successfully with each other.
3. Obtain the IQN of the E-Series system.
4. Make the E-Series system visible to the remote storage system. If the remote storage system is an E-Series system, then create a host using the IQN of the destination E-Series system as the connection information for the host port.
5. If the remote storage device is in use by a host/application:
 - Stop I/O to the remote storage device.
 - Unmap/unmount the remote storage device.
6. Map the remote storage device to the host defined for the E-Series storage system.
7. Obtain the LUN number of the device used for the mapping.



Recommended: Back up the remote source volume before starting the import process.

Cable the storage arrays

As part of the setup process, the storage arrays and I/O host must be cabled to the iSCSI-compatible interface.

The following diagrams provide examples of how to cable the systems such that they perform Remote Storage Volume operations over an iSCSI connection.

Fabric Connection - Use Case 1



Fabric Connection - Use Case 2



Configure the iSCSI ports

You must configure the iSCSI ports to ensure communication between the target (local E-Series storage array) and source (remote storage array).

The iSCSI ports can be configured multiple ways based on your subnet. The following are a few examples on how to configure the iSCSI ports for use with the Remote Storage Volumes feature.

Source A	Source B	Target A	Target B
10.10.1.100/22	10.10.2.100/22	10.10.1.101/22	10.10.2.101/22

Source A	Source B	Target A	Target B
10.10.0.100/16	10.10.0.100/16	10.10.0.101/16	10.10.0.101/16

Import remote storage

To initiate a storage import from a remote system to a local E-Series storage system, use the Import Remote Storage wizard in the SANtricity System Manager user interface.

What you'll need

- The E-Series storage system must be configured to communicate with the remote storage system. See [Configure hardware](#).
- For the remote storage system, gather the following information:
 - iSCSI IQN
 - iSCSI IP addresses
 - LUN number of the remote storage device (source volume)
- For the local E-Series storage system, create or select a volume to be used for the data import. The target volume must meet the following requirements:
 - Matches the block size of the remote storage device (the source volume).
 - Has a capacity that is equal to or larger than the remote storage device.
 - Has a state of Optimal and is available. For a full list of requirements, see [Requirements and restrictions](#).
- Recommended: Back up volumes on the remote storage system before starting the import process.

About this task

In this task, you create a mapping between the remote storage device and a volume on the local E-Series storage system. When you finish the configuration, the import begins.



Because many variables can impact the import operation and its completion time, you should first perform smaller “test” imports. Use these tests to ensure that all connections work as expected and that the import operation completes in an appropriate amount of time.

Steps

1. From the SANtricity System Manager, click **Storage > Remote Storage**.

2. Click **Import Remote Storage**.

A wizard for importing remote storage is displayed.

3. In Step 1a of the Configure Source panel, enter connection information.

- a. Under the **Name** field, enter the name for the remote storage device.
- b. Under the **iSCSI connection properties**, enter the following for the remote storage device: IQN, IP address, and the port number (default is 3260).

If you want to add another iSCSI connection, click **+Add another IP address** to include an additional IP address for the remote storage. When you are done, click **Next**.

After you click Next, Step 1b of the Configure Source panel is displayed.

4. Under the **LUN** field, select the desired source LUN for the remote storage device, and then click **Next**.

The Configure Target panel opens and displays volume candidates to serve as the target for the import. Some volumes do not display in the list of candidates due to block size, capacity, or volume availability.

5. From the table, select a target volume on the E-Series storage system. If needed, use the slider to change the import priority. Click **Next**. Confirm the operation in the next dialog box by typing `continue`, and then clicking **Continue**.

If the target volume has a capacity that is larger than the source volume, that additional capacity is not reported to the host connected to the E-Series system. To use the new capacity, you must perform a file system expansion operation on the host after the import operation completes and is disconnected.

After you confirm the configuration in the dialog, the Review panel is displayed.

6. From the Review screen, verify the specified remote storage device, target, and import settings are accurate. Click **Finish** to complete the creation of the remote storage.

Another dialog box opens asking if you want to initiate another import.

7. If needed, click **Yes** to create another remote storage import. Clicking Yes returns to Step 1a of the Configure Source panel, where you can select the existing configuration or add a new one. If you do not want to create another import, click **No** to exit the dialog.

Once the import process begins, the entire target volume is overwritten with the copied data. If the host writes any new data to the target volume during this process, that new data is propagated back to the remote device (source volume).

8. View the progress of the operation in the View Operations dialog under the Remote Storage panel.

The time required to complete the import operation depends on the size of the remote storage system, the priority setting for the import, and the amount of I/O load on both storage systems and their associated volumes. Once the import is complete, the local volume is a duplicate of the remote storage device.

9. When you are ready to break the relationship between the two volumes, select **Disconnect** on the import object from the Operations in Progress view. Once the relationship is disconnected, performance of the local volume returns to normal and is no longer impacted by the remote connection.

Manage import progress

After the import process begins, you can view and take action on its progress.

For each import operation, the Operations in Progress page displays a percentage of completion and estimated time remaining. Actions include changing the import priority, stopping and resuming operations, and disconnecting from the operation.



You can also view Operations in Progress from the Home page (**Home > Show operations in progress**).

Steps

1. In SANtricity System Manager, go to the Remote Storage page and select **View Operations**.

The Operations in Progress dialog is displayed.

2. If desired, use the links in the Actions column to stop and resume, change priority, or disconnect from an operation.
 - **Change Priority** – Select **Change Priority** to change the processing priority of an operation that is in progress or pending. Apply a priority to the operation and then click **OK**.
 - **Stop** – Select **Stop** to pause the copying of data from the remote storage device. The relationship between the import pair is still intact, and you can select **Resume** when you are ready to continue the import operation.
 - **Resume** – Select **Resume** to begin a stopped or failed process from where it left off. Next, apply a priority to the Resume operation, and then click **OK**.

The Resume operation does **not** restart the import from the beginning. If you want to restart the process from the beginning, you must select **Disconnect**, and then re-create the import through the Import Remote Storage wizard.

- **Disconnect** – Select **Disconnect** to break the relationship between the source and destination volumes for an import operation that has stopped, completed, or failed.

Modify remote storage connection settings

You can edit, add, or delete connection settings for any remote storage configuration through the View/Edit Settings option.

Making changes to connection properties will affect in-progress imports. To avoid disruptions, only make changes to connection properties when imports are not running.

Steps

1. From the Remote Storage screen of the SANtricity System Manager, select the desired Remote Storage object under the result list section.
2. Click **View/Edit Settings**.

The Remote Storage Settings screen is displayed.

3. Click the **Connection Properties** tab.

The configured IP address and port settings for the remote storage import are displayed.

4. Perform one of the following actions:

- **Edit** – Click **Edit** next to the corresponding line item for the remote storage object. Enter the revised IP address and/or port information in the fields.
- **Add** – Click **Add**, and then enter the new IP address and port information in the fields provided. Click **Add** to confirm, and then the new connection appears in the list of remote storage objects.
- **Delete** – Select the desired connection from the list and then click **Delete**. Confirm the operation by typing `delete` in the provided field and then click **Delete**. The connection is removed from the list of remote storage objects.

5. Click **Save**.

The modified connection settings are applied to the remote storage object.

Remove remote storage object

After an import completes, you can remove a remote storage object if you no longer want data copied between the local and remote devices.

Steps

1. Make sure that no imports are associated with the remote storage object you plan to remove.
2. From the Remote Storage screen of the SANtricity System Manager, select the desired Remote Storage object under the result list section.
3. Click **Remove**.

The Confirm Remove Remote Storage Connection dialog is displayed.

4. Confirm the operation by typing `remove` and then clicking **Remove**.

The selected Remote Storage object is removed.

Storage plugin for vCenter

Overview of the Storage Plugin for vCenter

The SANtricity Storage Plugin for vCenter provides integrated management of E-Series storage arrays from within a VMware vSphere Client session.

Available tasks

You can use the plugin to perform the following tasks:

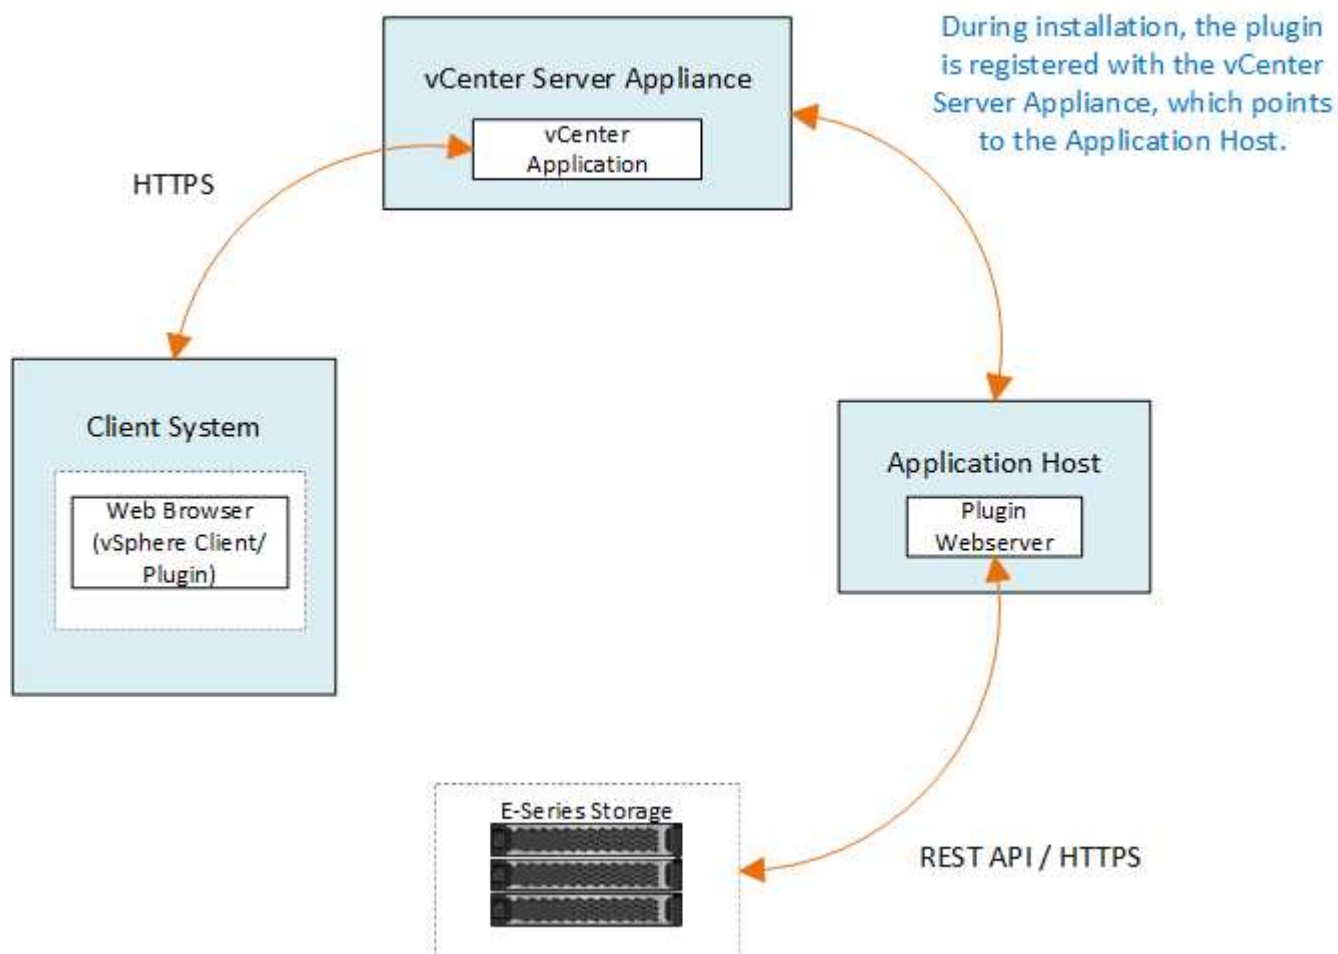
- View and manage discovered storage arrays in the network.
- Perform batch operations on groups of multiple storage arrays.
- Perform upgrades on the software OS.
- Import settings from one storage array to another.
- Configure volumes, SSD cache, hosts, host clusters, pools, and volume groups.
- Launch the System Manager interface for additional management tasks on an array.



The plugin is not a direct replacement for the System Manager interface, which is embedded on each controller for a storage array. System Manager provides additional management features; if desired, you can open System Manager by selecting a storage array in the main view of the plugin and then clicking **Launch**.

The plugin requires a VMware vCenter Server Appliance deployed in the VMware environment and an application host to install and run the plugin webserver.

Refer to the following figure for more information on communications in the vCenter environment.



Interface overview

When you log in to the plugin, the main page opens to **Manage - All**. From this page, you can view and manage all discovered storage arrays in your network.

Navigation sidebar

The navigation sidebar displays the following:

- **Manage** — Discover storage arrays in your network, launch System Manager for an array, import settings from one array to multiple arrays, manage array groups, upgrade the SANtricity OS, and provision storage.
- **Certificate Management** — Manage certificates to authenticate between browsers and clients.
- **Operations** — View the progress of batch operations, such as importing settings from one array to another.



Some operations are not available when a storage array has a non-optimal status.

- **Support** — View technical support options, resources, and contacts.

Supported browsers

The Storage Plugin for vCenter can be accessed from several types of browsers. The following browsers and versions are supported.

- Google Chrome 89 or later
- Mozilla Firefox 80 or later
- Microsoft Edge 90 or later

User roles and permissions

To access tasks in the Storage Plugin for vCenter, the user must have read-write permissions. By default, all defined VMware vCenter user IDs have no permissions to perform tasks in the plugin.

Configuration overview

Configuration involves the following steps:

1. [Install and register the plugin.](#)
2. [Configure plugin access permissions.](#)
3. [Log in to the plugin interface.](#)
4. [Discover storage arrays.](#)
5. [Provision storage.](#)

Find more information

For more information about managing datastores in the vSphere Client, see [VMware vSphere Documentation](#).

Get started

Review installation and upgrade requirements

Before installing or upgrading the SANtricity Storage Plugin for vCenter, review the installation requirements and upgrade considerations.

Installation requirements

You can install and configure the Storage Plugin for vCenter on a Windows host system. Plugin installation includes the following requirements.

Requirement	Description
Supported versions	<ul style="list-style-type: none"> • VMware vCenter Server Appliance supported versions: 6.7U3J, 7.0U1, 7.0U2, 7.0U3, and 8.0. • NetApp SANtricity OS version: 11.60.2 or higher • Supported application host versions: Windows 2016, Windows 2019, Windows 2022. <p>For more information on compatibility, see the NetApp Interoperability Matrix Tool.</p>
Multiple instances	You can install only one instance of Storage Plugin for vCenter on a Windows host and can only register it to one vCSA.
Capacity planning	<p>Storage Plugin for vCenter requires adequate space for execution and logging. Make sure that your system meets the following available disk space requirements:</p> <ul style="list-style-type: none"> • Required installation space — 275 MB • Storage space — 275 MB + 200 MB (logging) • System memory — 512 MB
License	The Storage Plugin for vCenter is a free, standalone product that does not require a license key. However, applicable copyrights and terms of service apply.

Upgrade considerations

If you are upgrading from a previous version, be aware that the plugin must be unregistered from the vCSA prior to the upgrade.

- During the upgrade, most of the plugin's previous configuration settings are preserved. These settings include user passwords, all discovered storage systems, server certificates, trusted certificates, and server runtime configuration.
- The upgrade process does not preserve the **vcenter.properties** files, so you must unregister the plugin prior to the upgrade. Once the upgrade is successful, you can then register the plugin again to the vCSA.
- All SANtricity OS files previously loaded in the repository are removed during the upgrade.

Install or upgrade the Storage Plugin for vCenter

Follow these steps to install the Storage Plugin for vCenter and verify the plugin registration. You can also upgrade the plugin using these instructions.

Review installation prerequisites

Be sure that your systems meet the requirements in [Review installation and upgrade requirements](#).



The upgrade process does not preserve the **vcenter.properties** files. If you are upgrading, you must unregister the plugin prior to the upgrade. Once the upgrade is successful, you can then register the plugin again to the vCSA.

Install the plugin software

To install the plugin software:

1. Copy the installer file to the host that will be used as the application server, and then access the folder where you downloaded the installer.
2. Double-click the installation file:

```
santricity_storage_vcenterplugin-windows_x64-- nn.nn.nn.nnnn.exe
```

In the above filename, `nn.nn.nn.nnnn` represents the version number.

3. When the installation starts, follow the on-screen prompts to enable several features and enter some configuration parameters. If necessary, you can change any of these selections later in the configuration files.



During an upgrade, you are not prompted for configuration parameters.



During installation, you are prompted for certificate validation. Keep the checkbox selected if you want to enforce certificate validation between the plugin and the storage arrays. With this enforcement, the storage array certificates are checked to be trusted against the plugin. If the certificates are not trusted, then they are not allowed to be added to the plugin. If you want to override certificate validation, deselect the checkbox so that all storage arrays can be added to the plugin using self-signed certificates. To learn more about certificates, refer to the online help available from the plugin interface.

4. When the Webserver Started message appears, click **OK** to complete the installation, and then click **Done**.
5. Verify that the application server was installed successfully by running the **services.msc** command.
6. Verify that the Application Server (vCP) service, **NetApp SANtricity Storage Plugin for vCenter**, was installed and the service has started.



If necessary, you can change the Certificate Validation and Web Service Port settings after installation. From the installation directory, open the `wsconfig.xml` file. To remove the Certificate Validation on storage arrays, change the `env` key, `trust.all.arrays`, to `true`. To change the Web Services port, modify the `sslport` value to the desired port value ranging from 0-65535. Ensure that the port number used is not binding to another process. When you are done, save the changes and restart the plugin webserver. If the port value of the plugin webserver is changed after registering the plugin to a vCSA, then you must unregister and re-register the plugin so the vCSA is communicating on the changed port to the plugin webserver.

Register the plugin with a vCenter Server Appliance

After the plugin software is installed, register the plugin with a vCSA.



The plugin can only be registered to one vCSA at a time. To register to a different vCSA, you must unregister the plugin from the current vCSA and uninstall it from the application host. You can then re-install the plugin and register it to the other vCSA.

1. Open a prompt through the command line and navigate to the following directory:

```
<install directory>\vcenter-register\bin
```

2. Execute the **vcenter-register.bat** file:

```
vcenter-register.bat ^  
-action registerPlugin ^  
-vcenterHostname <vCenter FQDN> ^  
-username <Administrator username> ^
```

3. Verify that the script was successful.

The logs are saved to %install_dir%/working/logs/vc-registration.log.

Verify the plugin registration

After the plugin is installed and the registration script has executed, verify that the plugin successfully registered with the vCenter Server Appliance.

1. Open the vSphere Client to the vCenter Server Appliance.
2. On the menu bar, select **Administrator > Client Plugins**.
3. Make sure the Storage Plugin for vCenter is listed as **Enabled**.

If the plugin is listed as Disabled with an error message stating that it cannot communicate with the application server, verify that the port number defined for the application server is enabled to pass through any firewalls that might be in use. The default application server Transmission Control Protocol (TCP) port number is 8445.

Configure plugin access permissions

You can configure access permissions for the Storage Plugin for vCenter, which includes users, roles, and privileges.

Review required vSphere privileges

To access the plugin within the vSphere Client, you must be assigned to a role that has the appropriate vSphere privileges. Users with the “Configure datastore” vSphere privilege have read-write access to the plugin, while users with the “Browse datastore” privilege have read-only access. If a user has neither of these privileges, the plugin displays an “Insufficient Privileges” message.

Plugin access type	vSphere privilege required
Read-Write (Configure)	Datastore.Configure
Read-Only (View)	Datastore.Browse

Configure Storage Administrator roles

To provide read/write privileges for plugin users, you can create, clone, or edit a role. For more information about configuring roles in the vSphere Client, see the following topic in the VMware Doc Center:

- [Create a Custom Role](#)

Access role actions

1. From the home page of the vSphere Client, select **Administrator** from the access control area.
2. Click **Roles** from the access control area.
3. Perform one of the following actions:
 - **Create new role:** Click on the **Create Role** action icon.
 - **Clone role:** Select an existing role and click on the **Clone Role** action icon.
 - **Edit existing role:** Select an existing role and click on the **Edit Role** action icon.



The Administrator role is not editable.

The appropriate wizard appears, depending on the above selection.

Create a new role

1. In the Privileges list, select the access permissions to assign to this role.

To allow Read-Only access to the plugin, select **Datastore** › **Browse datastore**. To allow Read-Write access, select **Datastore** › **Configure datastore**.

2. Assign other privileges for the list if needed, and then click **Next**.
3. Name the role and provide a description.
4. Click **Finish**.

Clone a role

1. Name the role and provide a description.
2. Click **OK** to finish the wizard.
3. Select the cloned role from the list, and then click on **Edit Role**.
4. In the Privileges list, select the access permissions to assign to this role.

To allow Read-Only access to the plugin, select **Datastore** › **Browse datastore**. To allow Read-Write access, select **Datastore** › **Configure datastore**.

5. Click **Next**.
6. Update the name and description, if desired.
7. Click **Finish**.

Edit an existing role

1. In the Privileges list, select the access permissions to assign to this role.

To allow Read-Only access to the plugin, select **Datastore > Browse datastore**. To allow Read-Write access, select **Datastore > Configure datastore**.

2. Click **Next**.
3. Update the name or description, if desired.
4. Click **Finish**.

Set permissions for vCenter Server Appliance

After setting privileges for a role, you must then add a permission to the vCenter Server Appliance. This permission allows a given user or group access to the plugin.

1. From the menu dropdown list, select **Hosts and Clusters**.
2. Select the **vCenter Server Appliance** from the access control area.
3. Click the **Permissions** tab.
4. Click the **Add Permission** action icon.
5. Select the appropriate domain and user/group.
6. Select the role created that allows for the read/write plugin privilege.
7. Enable the **Propagate to Children** option, if needed.
8. Click **OK**.



You can select an existing permission and modify it to use the created role. **However, be aware that the role must have the same privileges along with read/write plugin privileges as to avoid a regress in permissions.**

To access the plugin, you must log in to the vSphere Client under the user account that has the read/write privileges for the plugin.

For more information about managing permissions, see the following topics in the VMware Doc Center:

- [Managing Permissions for vCenter Components](#)
- [Best Practices for Roles and Permissions](#)

Log in and navigate the Storage Plugin for vCenter

You can log in to the Storage Plugin for vCenter to navigate the user interface.

1. Before you log in to the plugin, make sure you are using one of the following browsers:
 - Google Chrome 89 or later
 - Mozilla Firefox 80 or later
 - Microsoft Edge 90 or later
2. Log in to the vSphere Client under the user account that has read/write privileges for the plugin.
3. From the vSphere Client Home page, click **SANtricity Storage Plugin for vCenter**.

The plugin opens within a vSphere Client window. The plugin's main page opens to **Manage-All**.

4. Access storage management tasks from the navigation sidebar on the left:

- **Manage** – Discover storage arrays in your network, open System Manager for an array, import settings from one array to multiple arrays, manage array groups, upgrade the OS software, and provision storage.
- **Certificate Management** – Manage certificates to authenticate between browsers and clients.
- **Operations** – View the progress of batch operations, such as importing settings from one array to another.
- **Support** – View technical support options, resources, and contacts.



Some operations are not available when a storage array has a non-optimal status.

Discover storage arrays in the plugin

To display and manage storage resources, you must use the Storage Plugin for vCenter interface to discover the IP addresses of arrays in your network.

Before you begin

- You must know the network IP addresses (or range of addresses) of the array controllers.
- The storage arrays must be correctly set up and configured, and you must know the storage array login credentials (user name and password).

Step 1: Enter network addresses for discovery

Steps

1. From the Manage page, select **Add/Discover**.

The Enter Network Address Range dialog box appears.

2. Do one of the following:
 - To discover one array, select the **Discover a single storage array** radio button, and then enter the IP address for one of the controllers in the storage array.
 - To discover multiple storage arrays, select the **Discover all storage arrays within a network range** radio button, and then enter the starting network address and ending network address to search across your local sub-network.
3. Click **Start Discovery**.

As the discovery process begins, the dialog box displays the storage arrays as they are discovered. The discovery process might take several minutes to complete.

If no manageable arrays are discovered, verify that the storage arrays are properly connected to your network and their assigned addresses are within range. Click **New Discovery Parameters** to return to the Add/Discover page.

4. Select the checkbox next to any storage array that you want to add to your management domain.

The system performs a credential check on each array that you are adding to the management domain. You might need to resolve any issues with untrusted certificates before proceeding.

5. Click **Next** to proceed to the next step in the wizard.

If the storage arrays have valid certificates, go to [Step 3: Provide passwords](#).

If any storage arrays do not have valid certificates, the Resolve Self-Signed Certificates dialog box appears. Go to [Step 2: Resolve untrusted certificates during discovery](#).

If you want to import CA-signed certificates, cancel out of the discovery wizard and click **Certificate Management** from the left panel. Refer to the online help for further instructions.

Step 2: Resolve untrusted certificates during discovery

You must resolve any certificate issues before proceeding with the discovery process.

1. If the Resolve Self-Signed Certificates dialog box opens, review the information displayed for the untrusted certificates. For more information, you can also click the ellipses at the far end of the table and select **View** from the pop-up menu.
2. Do one of the following:
 - If you trust the connections to the discovered storage arrays, click **Next** and then click **Yes** to confirm and continue to the next dialog in the wizard. The self-signed certificates are marked as trusted and the storage arrays will be added to the plugin.
 - If you do not trust the connections to the storage arrays, select **Cancel** and validate each storage array's security certificate strategy before adding any of them.
3. Click **Next** to proceed to the next step in the wizard.

Step 3: Provide passwords

As the last step for discovery, you must enter the passwords for the storage arrays that you want to add to your management domain.

1. For each discovered array, enter its admin password in the fields.
2. Click **Finish**.

It can take several minutes for the system to connect to the specified storage arrays. When the process is finished, the storage arrays are added to your management domain and associated with the selected group (if specified).

Provision storage in the plugin

To provision storage, you create volumes, assign volumes to hosts, and then assign volumes to datastores.

Step 1: Create volumes

Volumes are data containers that manage and organize the storage space on your storage array. You create volumes from the storage capacity available on your storage array, which helps organize your system's resources. The concept of "volumes" is similar to using folders/directories on a computer to organize files for quick access.

Volumes are the only data layer visible to hosts. In a SAN environment, volumes are mapped to logical unit numbers (LUNs). These LUNs hold the user data that is accessible using one or more of the host access protocols supported by the storage array.

Steps

1. From the Manage page, select the storage array.

2. Select **Provisioning > Manage Volumes**.

3. Select **Create > Volumes**.

The Select Host dialog box appears.

4. From the drop-down list, select a specific host or host cluster to which you want to assign volumes, or choose to assign the host or host cluster at a later time.

5. To continue the volume creation sequence for the selected host or host cluster, click **Next**.

The Select Workload dialog box appears. A workload contains volumes with similar characteristics, which are optimized based on the type of application the workload supports. You can define a workload or you can select existing workloads.

6. Do one of the following:

- Select the **Create volumes for an existing workload** option and then select the workload from the drop-down list.
- Select the **Create a new workload** option to define a new workload for a supported application or for “Other” applications, and then following these steps:
 - a. From the drop-down list, select the name of the application you want to create the new workload for. Select one of the “Other” entries if the application you intend to use on this storage array is not listed.
 - b. Enter a name for the workload you want to create.

7. Click **Next**. If your workload is associated with a supported application type, enter the information requested; otherwise, go to the next step.

The Add/Edit Volumes dialog box appears. In this dialog, you create volumes from eligible pools or volume groups. For each eligible pool and volume group, the number of drives available and the total free capacity appears. For some application-specific workloads, each eligible pool or volume group shows the proposed capacity based on the suggested volume configuration and shows the remaining free capacity in GiB. For other workloads, the proposed capacity appears as you add volumes to a pool or volume group and specify the reported capacity.

8. Before you begin adding volumes, read the guidelines in the following table.

Field	Description
Free capacity	Because volumes are created from pools or volume groups, the pool or volume group you select must have sufficient free capacity.

Field	Description
Data Assurance (DA)	<p>To create a DA-enabled volume, the host connection you are planning to use must support DA.</p> <ul style="list-style-type: none"> • If you want to create a DA-enabled volume, select a pool or volume group that is DA capable (look for Yes next to "DA" in the pool and volume group candidates table). • DA capabilities are presented at the pool and volume group level. DA protection checks for and corrects errors that might occur as data is transferred through the controllers down to the drives. Selecting a DA-capable pool or volume group for the new volume ensures that any errors are detected and corrected. • If any of the host connections on the controllers in your storage array do not support DA, the associated hosts cannot access data on DA-enabled volumes.
Drive security	<p>To create a secure-enabled volume, a security key must be created for the storage array.</p> <ul style="list-style-type: none"> • If you want to create a secure-enabled volume, select a pool or volume group that is secure capable (look for Yes next to "Secure-capable" in the pool and volume group candidates table). • Drive security capabilities are presented at the pool and volume group level. Secure-capable drives prevent unauthorized access to the data on a drive that is physically removed from the storage array. A secure-enabled drive encrypts data during writes and decrypts data during reads using a unique encryption key. • A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.
Resource provisioning	<p>To create a resource-provisioned volume, all drives must be NVMe drives with the Deallocated or Unwritten Logical Block Error (DULBE) option.</p>

9. Choose one of these actions based on whether you selected "Other" or an application-specific workload in the previous step:

- **Other** – Click **Add new volume** in each pool or volume group that you want to use to create one or more volumes.
- **Application-specific workload** – Either click **Next** to accept the system-recommended volumes and characteristics for the selected workload, or click **Edit Volumes** to change, add, or delete the system-recommended volumes and characteristics for the selected workload.

The following fields appear.

Field	Description
Volume Name	A volume is assigned a default name during the volume creation sequence. You can either accept the default name or provide a more descriptive one indicating the type of data stored in the volume.
Reported Capacity	<p>Define the capacity of the new volume and the capacity units to use (MiB, GiB, or TiB). For thick volumes, the minimum capacity is 1 MiB, and the maximum capacity is determined by the number and capacity of the drives in the pool or volume group.</p> <p>Capacity in a pool is allocated in 4-GiB increments. Any capacity that is not a multiple of 4 GiB is allocated but not usable. To make sure that the entire capacity is usable, specify the capacity in 4-GiB increments. If unusable capacity exists, the only way to regain it is to increase the capacity of the volume.</p>
Volume Type	If you selected "Application-specific workload," the Volume Type field appears. This indicates the type of volume that was created for an application-specific workload.
Volume Block Size (EF300 and EF600 only)	<p>Shows the block sizes that can be created for the volume:</p> <ul style="list-style-type: none"> • 512 – 512 bytes • 4K – 4,096 bytes

Field	Description
Segment Size	<p>Shows the setting for segment sizing, which only appears for volumes in a volume group. You can change the segment size to optimize performance.</p> <p>Allowed segment size transitions – The system determines the segment size transitions that are allowed. Segment sizes that are inappropriate transitions from the current segment size are unavailable on the drop-down list. Allowed transitions usually are double or half of the current segment size. For example, if the current volume segment size is 32 KiB, a new volume segment size of either 16 KiB or 64 KiB is allowed.</p> <p>SSD Cache-enabled volumes – You can specify a 4-KiB segment size for SSD Cache-enabled volumes. Make sure you select the 4-KiB segment size only for SSD Cache-enabled volumes that handle small-block I/O operations (for example, 16 KiB I/O block sizes or smaller). Performance might be impacted if you select 4 KiB as the segment size for SSD Cache-enabled volumes that handle large block sequential operations.</p> <p>Amount of time to change segment size – The amount of time to change a volume's segment size depends on these variables:</p> <ul style="list-style-type: none"> • The I/O load from the host • The modification priority of the volume • The number of drives in the volume group • The number of drive channels • The processing power of the storage array controllers <p>When you change the segment size for a volume, I/O performance is affected, but your data remains available.</p>
Secure-capable	<p>Yes appears next to "Secure-capable" only if the drives in the pool or volume group are encryption-capable.</p> <p>Drive Security prevents unauthorized access to the data on a drive that is physically removed from the storage array. This option is available only when the Drive Security feature has been enabled, and a security key is set up for the storage array.</p> <p>A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.</p>
DA	<p>Yes appears next to "DA" only if the drives in the pool or volume group support Data Assurance (DA).</p> <p>DA increases data integrity across the entire storage system. DA enables the storage array to check for errors that might occur as data is transferred through the controllers down to the drives. Using DA for the new volume ensures that any errors are detected.</p>

10. To continue the volume creation sequence for the selected application, click **Next**.

11. In the last step, review a summary of the volumes you intend to create and make any necessary changes. To make changes, click **Back**. When you are satisfied with your volume configuration, click **Finish**.

Step 2: Create host access and assign volumes

A host can be created automatically or manually:

- **Automatic** — Automatic host creation for SCSI-based (not NVMe-oF) hosts is initiated by the Host Context Agent (HCA). The HCA is a utility that you can install on each host attached to the storage array. Each host that has the HCA installed pushes its configuration information to the storage array controllers through the I/O path. Based on the host information, the controllers automatically create the host and the associated host ports and set the host type. If needed, you can make any additional changes to the host configuration. After the HCA performs its automatic detection, the host is automatically configured with the following attributes:
 - The host name derived from the system name of the host.
 - The host identifier ports that are associated with the host.
 - The Host Operating System Type of the host.



Host Context Agent software for Linux and Windows is available from [NetApp Support - Downloads](#).



Hosts are created as stand-alone hosts; the HCA does not automatically create or add to host clusters.

- **Manual** – During manual host creation, you associate host port identifiers by selecting them from a list or manually entering them. After you create a host, you can assign volumes to it or add it to a host cluster if you plan to share access to volumes.

Using the HCA to auto-discover the host

You can allow the Host Context Agent (HCA) to automatically detect the hosts, and then verify that the information is correct.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning > Configure Hosts**.

The Configure Hosts page opens.

3. Select **Storage > Hosts**.

The table lists the automatically created hosts.

4. Verify that the information provided by the HCA is correct (name, host type, host port identifiers).
5. If you need to change any of the information, select the host, and then click **View/Edit Settings**.

Manually creating the host

Before you begin

Read the following guidelines:

- You must already have added or discovered storage arrays within your environment.
- You must define the host identifier ports that are associated with the host.

- Make sure that you provide the same name as the host's assigned system name.
- This operation does not succeed if the name you choose is already in use.
- The length of the name cannot exceed 30 characters.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning › Configure Hosts**.

The Configure Hosts page opens.

3. Click **Create › Host**.

The Create Host dialog box appears.

4. Select the settings for the host as appropriate.

Field	Description
Name	Type a name for the new host.
Host operating system type	Select the operating system that is running on the new host from the drop-down list.
Host interface type	(Optional) If you have more than one type of host interface supported on your storage array, select the host interface type that you want to use.
Host ports	<p>Do one of the following:</p> <ul style="list-style-type: none"> • Select I/O Interface — Generally, the host ports should have logged in and be available from the drop-down list. You can select the host port identifiers from the list. • Manual add — If a host port identifier is not displayed in the list, it means that the host port has not logged in. An HBA utility or the iSCSI initiator utility may be used to find the host port identifiers and associate them with the host. <p>You can manually enter the host port identifiers or copy/paste them from the utility (one at a time) into the Host ports field.</p> <p>You must select one host port identifier at a time to associate it with the host, but you can continue to select as many identifiers that are associated with the host. Each identifier is displayed in the Host ports field. If necessary, you also can remove an identifier by selecting the X next to it.</p>

Field	Description
Set CHAP initiator secret	<p>(Optional) If you selected or manually entered a host port with an iSCSI IQN, and if you want to require a host that tries to access the storage array to authenticate using Challenge Handshake Authentication Protocol (CHAP), select the Set CHAP initiator secret checkbox. For each iSCSI host port you selected or manually entered, do the following:</p> <ul style="list-style-type: none"> • Enter the same CHAP secret that was set on each iSCSI host initiator for CHAP authentication. If you are using mutual CHAP authentication (two-way authentication that enables a host to validate itself to the storage array and for a storage array to validate itself to the host), you also must set the CHAP secret for the storage array at initial setup or by changing settings. • Leave the field blank if you do not require host authentication. <p>Currently, the only iSCSI authentication method used is CHAP.</p>

5. Click **Create**.

6. If you need to update the host information, select the host from the table and click **View/Edit Settings**.

After the host is successfully created, the system creates a default name for each host port configured for the host (user label). The default alias is <Hostname_Port Number>. For example, the default alias for the first port created for host IPT is IPT_1.

7. Next, you must assign a volume to a host or a host cluster so it can be used for I/O operations. Select **Provisioning > Configure Hosts**.

The Configure Hosts page opens.

8. Select the host or host cluster to which you want to assign volumes, and then click **Assign Volumes**.

A dialog box appears that lists all the volumes that can be assigned. You can sort any of the columns or type something in the Filter box to make it easier to find particular volumes.

9. Select the check box next to each volume that you want to assign or select the check box in the table header to select all volumes.

10. Click **Assign** to complete the operation.

The system performs the following actions:

- The assigned volume receives the next available LUN number. The host uses the LUN number to access the volume.
- The user-supplied volume name appears in volume listings associated to the host. If applicable, the factory-configured access volume also appears in volume listings associated to the host.

Step 3: Create a datastore in vSphere Client

To create a datastore in the vSphere Client, see the following topic in the VMware Doc Center:

- [Create a VMFS Datastore in the vSphere Client](#)

Increase capacity of existing datastore by increasing volume capacity

You can increase the reported capacity (the capacity reported to hosts) of a volume by using the free capacity that is available in the pool or volume group.

Before you begin

Make sure that:

- Enough free capacity is available in the volume's associated pool or volume group.
- The volume is Optimal and not in any state of modification.
- No hot spare drives are in use in the volume. (Applies only to volumes in volume groups.)



Increasing the capacity of a volume is supported only on certain operating systems. If you increase the volume capacity on a host operating system that does not support LUN expansion, the expanded capacity is unusable, and you cannot restore the original volume capacity.

Steps

1. Navigate to the plugin within vSphere Client.
2. Within the plugin, select the desired storage array.
3. Click on **Provisioning** and select **Manage Volumes**.
4. Select the volume for which you want to increase capacity, and then select **Increase Capacity**.

The Confirm Increase Capacity dialog box appears.

5. Select **Yes** to continue.

The Increase Reported Capacity dialog box appears.

This dialog box displays the volume's current reported capacity and the free capacity available in the volume's associated pool or volume group.

6. Use the **Increase reported capacity by adding...** box to add capacity to the current available reported capacity. You can change the capacity value to display in either mebibytes (MiB), gibibytes (GiB), or tebibytes (TiB).
7. Click **Increase**.
8. View the Recent Tasks pane for the progress of the increase capacity operation that is currently running for the selected volume. This operation can be lengthy and could affect system performance.
9. After the volume capacity is complete, you must manually increase the VMFS size to match as described in the following topic:
 - [Increase VMFS Datastore Capacity in the vSphere Client](#)

Increase capacity of existing datastore by adding volumes

1. You can increase the capacity of a datastore by adding volumes. Follow the steps in [Step 1: Create volumes](#).
2. Next, assign the volumes to the desired host to increase the datastore's capacity. See the following topic:
 - [Increase VMFS Datastore Capacity in the vSphere Client](#)

View status

You can view system status from the Storage Plugin for vCenter or from the vSphere Client.

1. Open the plugin from within the vSphere Client.
2. View status from the following panels:
 - **Storage array status** — Go to the **Manage-All** panel. For each discovered array, the row provides a Status column.
 - **Operations in progress** — Click **Operations** on the side panel to view all long-running tasks, such as importing settings. You can also view long-running operations from the Provisioning drop-down. For each operation listed on the Operations in Progress dialog, a percentage of completion and estimated time remaining to complete the operation are shown. In some cases, you can stop an operation or place it at a higher or lower priority. If desired, use the links in the Actions column to stop or change priority for an operation.



Read all cautionary text provided in the dialog boxes, particularly when stopping an operation.

Operations that might appear for the plugin are listed in the following table. Additional operations might also appear in the System Manager interface.

Operation	Possible status of the operation	Actions you can take
Volume create (thick pool volumes larger than 64TiB only)	In progress	none
Volume delete (thick pool volumes larger than 64TiB only)	In progress	none
Add capacity to pool or volume group	In progress	none
Change a RAID level for a volume	In progress	none
Reduce capacity for a pool	In progress	none
Check the time remaining on an instant availability format (IAF) operation for pool volumes	In progress	none
Check the data redundancy of a volume group	In progress	none
Initialize a volume	In progress	none
Increase capacity for a volume	In progress	none
Change segment size for a volume	In progress	none

Manage certificates

Certificates overview

Certificate Management in the Storage Plugin for vCenter allows you to create certificate signing requests (CSRs), import certificates, and manage existing certificates.

What are certificates?

Certificates are digital files that identify online entities, such as websites and servers, for secure communications on the internet. They ensure that web communications are transmitted in encrypted form, privately and unaltered, only between the specified server and client. Using the Storage Plugin for vCenter, you can manage certificates for the browser on a host management system and the controllers in the discovered storage arrays.

A certificate can be signed by a trusted authority, or it can be self-signed. "Signing" simply means that someone validated the owner's identity and determined that their devices can be trusted.

Storage arrays ship with an automatically generated self-signed certificate on each controller. You can continue to use the self-signed certificates, or you can obtain CA-signed certificates for a more secure connection between the controllers and the host systems.



Although CA-signed certificates provide better security protection (for example, preventing man-in-the-middle attacks), they also require fees that can be expensive if you have a large network. In contrast, self-signed certificates are less secure, but they are free. Therefore, self-signed certificates are most often used for internal testing environments, not in production environments.

Signed certificates

A signed certificate is validated by a certificate authority (CA), which is a trusted third-party organization. Signed certificates include details about the owner of the entity (typically, a server or website), date of certificate issue and expiration, valid domains for the entity, and a digital signature composed of letters and numbers.

When you open a browser and enter a web address, your system performs a certificate-checking process in the background to determine if you are connecting to a website that includes a valid, CA-signed certificate. Generally, a site that is secured with a signed certificate includes a padlock icon and an https designation in the address. If you attempt to connect to a website that does not contain a CA-signed certificate, your browser displays a warning that the site is not secure.

The CA takes steps to verify your identity during the application process. They might send an email to your registered business, verify your business address, and perform an HTTP or DNS verification. When the application process is complete, the CA sends you digital files to load on a host management system. Typically, these files include a chain of trust, as follows:

- **Root** — At the top of the hierarchy is the root certificate, which contains a private key used to sign other certificates. The root identifies a particular CA organization. If you use the same CA for all your network devices, you need only one root certificate.
- **Intermediate** — Branching off from the root are the intermediate certificates. The CA issues one or more intermediate certificates to act as middlemen between a protected root and server certificates.
- **Server** — At the bottom of the chain is the server certificate, which identifies your specific entity, such as a website or other device. Each controller in a storage array requires a separate server certificate.

Self-signed certificates

Each controller in the storage array includes a pre-installed, self-signed certificate. A self-signed certificate is similar to a CA-signed certificate, except that it is validated by the owner of the entity instead of a third party. Like a CA-signed certificate, a self-signed certificate contains its own private key, and also ensures that data is encrypted and sent over an HTTPS connection between a server and client.

Self-signed certificates are not “trusted” by browsers. Each time you attempt to connect to a website that contains only a self-signed certificate, the browser displays a warning message. You must click a link in the warning message that allows you to proceed to the website; by doing so, you are essentially accepting the self-signed certificate.

Management certificate

When you open the plugin, the browser attempts to verify that the management host is a trusted source by checking for a digital certificate. If the browser does not locate a CA-signed certificate, it opens a warning message. From there, you can continue to the website to accept the self-signed certificate for that session. You can also obtain signed, digital certificates from a CA so you no longer see the warning message.

Trusted certificates

During a plugin session, you might see additional security messages when you attempt to access a controller that does not have a CA-signed certificate. In this event, you can permanently trust the self-signed certificate or you can import the CA-signed certificates for the controllers so the plugin can authenticate incoming client requests from these controllers.

Use CA-signed certificates

You can obtain and import CA-signed certificates for secure access to the management system hosting the Storage Plugin for vCenter.

Using CA-signed certificates is a three-step procedure:

- [Step 1: Complete a CSR file.](#)
- [Step 2: Submit CSR file.](#)
- [Step 3: Import management certificates.](#)

Step 1: Complete a CSR file

You must first generate a certificate signing request (CSR) file, which identifies your organization and the host system where the plugin is running. Alternatively, you can generate a CSR file using a tool such as OpenSSL and skip to [Step 2: Submit CSR file.](#)

Steps

1. Select **Certificate Management**.
2. From the **Management** tab, select **Complete CSR**.
3. Enter the following information, and then click **Next**:
 - **Organization** — The full, legal name of your company or organization. Include suffixes, such as Inc. or Corp.
 - **Organizational unit (optional)** — The division of your organization that is handling the certificate.
 - **City/Locality** — The city where your host system or business is located.
 - **State/Region (optional)** — The state or region where your host system or business is located.
 - **Country ISO code** — Your country's two-digit ISO (International Organization for Standardization) code, such as US.
4. Enter the following information about the host system where the plugin is running:
 - **Common name** — The IP address or DNS name of the host system where the plugin is running. Make

sure this address is correct; it must match exactly what you enter to access the plugin in the browser. Do not include `http://` or `https://`. The DNS name cannot begin with a wildcard.

- **Alternate IP addresses** — If the common name is an IP address, you can optionally enter any additional IP addresses or aliases for the host system. For multiple entries, use a comma-delimited format.
- **Alternate DNS names** — If the common name is a DNS name, enter any additional DNS names for the host system. For multiple entries, use a comma-delimited format. If there are no alternate DNS names, but you entered a DNS name in the first field, copy that name here. The DNS name cannot begin with a wildcard.

5. Make sure that the host information is correct. If it is not, the certificates returned from the CA will fail when you try to import them.

6. Click **Finish**.

Step 2: Submit CSR file

After you create a certificate signing request (CSR) file, you send the generated CSR file to a CA to receive signed, management certificates for the system hosting the plugin.

E-Series systems require PEM format (Base64 ASCII encoding) for signed certificates, which includes the following file types: `.pem`, `.crt`, `.cer`, or `.key`.

Steps

1. Locate the downloaded CSR file.

The folder location of the download depends on your browser.

2. Submit the CSR file to a CA (for example, Verisign or DigiCert), and request signed certificates in PEM format.



After you submit a CSR file to the CA, do NOT regenerate another CSR file.

Whenever you generate a CSR, the system creates a private and public key pair. The public key is part of the CSR, while the private key is kept in the system's keystore. When you receive the signed certificates and import them, the system ensures that both the private and public keys are the original pair. If the keys do not match, the signed certificates will not work and you must request new certificates from the CA.

Step 3: Import management certificates

After you receive signed certificates from the Certificate Authority (CA), import the certificates into the host system where the plugin is installed.

Before you begin

- You must have the signed certificates from the CA. These files include the root certificate, one or more intermediate certificates, and the server certificate.
- If the CA provided a chained certificate file (for example, a `.p7b` file), you must unpack the chained file into individual files: the root certificate, one or more intermediate certificates, and the server certificate. You can use the Windows `certmgr` utility to unpack the files (right-click and select **All Tasks > Export**). Base-64 encoding is recommended. When the exports are complete, a CER file is shown for each certificate file in the chain.
- You must copy the certificate files to the host system where the plugin is running.

Steps

1. Select **Certificate Management**.
2. From the **Management** tab, select **Import**.

A dialog box opens for importing the certificate files.

3. Click **Browse** to first select the root and intermediate certificate files, and then select the server certificate. If you generated the CSR from an external tool, you must also import the private key file that was created along with the CSR.

The filenames are displayed in the dialog box.

4. Click **Import**.

Result

The files are uploaded and validated. The certificate information displays on the Certificate Management page.

Reset management certificates

For the management system hosting the Storage Plugin for vCenter, you can revert the management certificate to the original, factory self-signed state.

About this task

This task deletes the current management certificate from the host system where the Storage Plugin for vCenter is running. After the certificate is reset, the host system reverts to using the self-signed certificate.

Steps

1. Select **Certificate Management**.
2. From the **Management** tab, select **Reset**.

A Confirm Reset Management Certificate dialog box opens.

3. Type reset in the field, and then click **Reset**.

After your browser refreshes, the browser might block access to the destination site and report that the site is using HTTP Strict Transport Security. This condition arises when you switch back to self-signed certificates. To clear the condition that is blocking access to the destination, you must clear the browsing data from the browser.

Result

The system reverts to using the self-signed certificate from the server. As a result, the system prompts users to manually accept the self-signed certificate for their sessions.

Import certificates for arrays

If necessary, you can import certificates for the storage arrays so they can authenticate with the system hosting the Storage Plugin for vCenter. Certificates can be signed by a certificate authority (CA) or can be self-signed.

Before you begin

If you are importing trusted certificates, the certificates must be imported for the storage array controllers using

System Manager.

Steps

1. Select **Certificate Management**.
2. Select the **Trusted** tab.

This page shows all certificates reported for the storage arrays.

3. Select either **Import › Certificates** to import a CA certificate or **Import › Self-signed storage array certificates** to import a self-signed certificate.
4. To limit the view, you can use the **Show certificates that are...** filtering field or you can sort the certificate rows by clicking one of the column heads.
5. In the dialog box, select the certificate and then click **Import**.

The certificate is uploaded and validated.

View certificates

You can view summary information for a certificate, which includes the organization using the certificate, the authority that issued the certificate, the period of validity, and the fingerprints (unique identifiers).

Steps

1. Select **Certificate Management**.
2. Select one of the following tabs:
 - **Management** — Shows the certificate for the system hosting the plugin. A management certificate can be self-signed or approved by a certificate authority (CA). It allows secure access to the plugin.
 - **Trusted** — Shows certificates that the plugin can access for storage arrays and other remote servers, such as an LDAP server. The certificates can be issued from a certificate authority (CA) or can be self-signed.
3. To see more information about a certificate, select its row, select the ellipses at the end of the row, and then click **View** or **Export**.

Export certificates

You can export a certificate to view its complete details.

Before you begin

To open the exported file, you must have a certificate viewer application.

Steps

1. Select **Certificate Management**.
2. Select one of the following tabs:
 - **Management** — Shows the certificate for the system hosting the plugin. A management certificate can be self-signed or approved by a certificate authority (CA). It allows secure access to the plugin.
 - **Trusted** — Shows certificates that the plugin can access for storage arrays and other remote servers, such as an LDAP server. The certificates can be issued from a certificate authority (CA) or can be self-signed.

3. Select a certificate from the page, and then click the ellipses at the end of the row.
4. Click **Export**, and then save the certificate file.
5. Open the file in your certificate viewer application.

Delete trusted certificates

You can delete one or more certificates that are no longer needed, such as an expired certificate.

Before you begin

Import the new certificate before deleting the old one.



Be aware that deleting a root or intermediate certificate can impact multiple storage arrays, since these arrays can share the same certificate files.

Steps

1. Select **Certificate Management**.
2. Select the **Trusted** tab.
3. Select one or more certificates in the table, and then click **Delete**.



The Delete function is not available for pre-installed certificates.

The Confirm Delete Trusted Certificate dialog box opens.

4. Confirm the deletion, and then click **Delete**.

The certificate is removed from the table.

Resolve untrusted certificates

From the Certificate page, you can resolve untrusted certificates by importing a self-signed certificate from the storage array or by importing a certificate authority (CA) certificate that has been issued by a trusted third party.

Before you begin

If you plan to import a CA-signed certificate, make sure that:

- You have generated a certificate signing request (.CSR file) for each controller in the storage array and sent it to the CA.
- The CA returned trusted certificate files.
- The certificate files are available on your local system.

About this task

Untrusted certificates occur when a storage array attempts to establish a secure connection to the plugin, but the connection fails to confirm as secure. You might need to install additional trusted CA certificates if any of the following are true:

- You recently added a storage array.

- One or both certificates are expired or revoked.
- One or both certificates are missing a root or intermediate certificate.

Steps

1. Select **Certificate Management**.
2. Select the **Trusted** tab.

This page shows all certificates reported for the storage arrays.

3. Select either **Import › Certificates** to import a CA certificate or **Import › Self-Signed storage array certificates** to import a self-signed certificate.
4. To limit the view, you can use the **Show certificates that are...** filtering field or you can sort the certificate rows by clicking one of the column heads.
5. In the dialog box, select the certificate, and then click **Import**.

The certificate is uploaded and validated.

Manage arrays

Array management overview

Use the Add/Discover feature to find and add the storage arrays you want to manage in the Storage plugin for vCenter. From the Manage page, you can also rename, remove, and provide new passwords for these discovered arrays.

Considerations for discovering arrays

For the plugin to display and manage storage resources, you must discover the storage arrays you want to manage in your organization's network. You can discover and then add a single array or multiple arrays.

Multiple storage arrays

If you choose to discover multiple arrays, you enter a network IP address range and then the system attempts individual connections to each IP address in that range. Any storage array successfully reached appears in the plugin and you can then add them to your management domain.

Single storage array

If you choose to discover a single array, you enter the single IP address for one of the controllers in the storage array and then add that array to your management domain.



The plugin discovers and displays only the single IP address or IP address within a range assigned to a controller. If there are alternate controllers or IP addresses assigned to these controllers that fall outside of this single IP address or IP address range, then the plugin will not discover or display them. However, once you add the storage array, all associated IP addresses will be discovered and displayed in the Manage view.

User credentials

You must supply the administrator password for each storage array you want to add.

Certificates

As part of the discovery process, the system verifies that the discovered storage arrays are using certificates by a trusted source. The system uses two types of certificate-based authentication for all connections that it establishes with the browser:

- **Trusted certificates** — You might need to install additional trusted certificates supplied by the Certificate Authority if one or both controller certificates are expired, revoked, or missing a certificate in its chain.
- **Self-signed certificates** — Arrays can also use self-signed certificates. If you attempt to discover arrays without importing signed certificates, the plugin provides an additional step that allows you to accept the self-signed certificate. The storage array's self-signed certificate will be marked as trusted and the storage array will be added to the plugin. If you do not trust the connections to the storage array, select **Cancel** and validate the storage array's security certificate strategy before adding the storage array to the plugin.

Storage array status

When you open the Storage Plugin for vCenter, communication with each storage array is established and the status for each storage array is displayed.

From the **Manage - All** page, you can view the status of the storage array and the status of the storage array connection.

Status	Indicates
Optimal	The storage array is in an optimal state. There are no certificate issues and the password is valid.
Invalid Password	An invalid storage array password was provided.
Untrusted Certificate	One or more connections with the storage array is untrusted because the HTTPS certificate is either self-signed and has not been imported, or the certificate is CA-signed and the root and intermediate CA certificates have not been imported.
Needs Attention	There is a problem with the storage array that requires your intervention to correct it.
Lockdown	The storage array is in a locked-down state.
Unknown	The storage array has never been contacted. This can happen when the plugin is starting up and has not yet made contact with the storage array, or the storage array is offline and has never been contacted since the plugin was started.
Offline	The plugin had previously contacted the storage array, but now has lost all connection to it.

Plugin interface compared to System Manager

You can use Storage Plugin for vCenter for basic operating tasks on your storage array; however, there might be times when you need to launch System Manager to perform tasks not available in the plugin.

System Manager is an embedded application on the storage array's controller, which is connected to the network through an Ethernet management port. System Manager includes all array-based functions.

The following table helps you decide whether you can use the plugin interface or the System Manager interface for a particular storage array task.

Function	Plugin interface	System Manager interface
Batch operations on groups of multiple storage arrays	Yes	No. Operations are performed on a single array.
Upgrades for the SANtricity OS firmware	Yes. One or more arrays in a batch operation.	Yes. Single array only.
Import settings from one array to multiple arrays	Yes	No
Host and host cluster management (create, assign volumes, update, and delete)	Yes	Yes
Pools and volume group management (create, update, enable security, and delete)	Yes	Yes
Volume management (create, resize, update, and delete)	Yes	Yes
SSD Cache management (create, update, and delete)	Yes	Yes
Mirroring and snapshot management	No	Yes
Hardware management (view controller status, configure port connections, take controller offline, enable hot spares, erase drives, etc.)	No	Yes
Manage alerts (email, SNMP, and syslog)	No	Yes
Security key management	No	Yes
Certificate management for controllers	No	Yes
Access management for controllers (LDAP, SAML, etc.)	No	Yes
AutoSupport management	No	Yes

Discover storage arrays

To display and manage storage resources in the Storage Plugin for vCenter, you must discover the IP addresses of arrays in your network.

Before you begin

- You must know the network IP addresses (or range of addresses) of the array controllers.
- The storage arrays must be correctly set up and configured.

- Storage array passwords must be set up using System Manager's Access Management tile.

About this task

Array discovery is a multi-step procedure:

- [Step 1: Enter network addresses for discovery](#)
- [Step 2: Resolve untrusted certificates during discovery](#)
- [Step 3: Provide passwords](#)

Step 1: Enter network addresses for discovery

As the first step to discovering storage arrays, you enter a single IP address or a range of IP addresses to search across the local sub-network. The Add/Discover feature opens a wizard that guides you through the process of discovery.

Steps

1. From the **Manage** page, select **Add/Discover**.

The Enter Network Address Range dialog box appears.

2. Do one of the following:
 - To discover one array, select the **Discover a single storage array** radio button, and then enter the IP address for one of the controllers in the storage array.
 - To discover multiple storage arrays, select the **Discover all storage arrays within a network range** radio button, and then enter the starting network address and ending network address to search across your local sub-network.
3. Click **Start Discovery**.

As the discovery process begins, the dialog box displays the storage arrays as they are discovered. The discovery process might take several minutes to complete.



If no manageable arrays are discovered, verify that the storage arrays are properly connected to your network and their assigned addresses are within range. Click **New Discovery Parameters** to return to the Add/Discover page.

4. Select the checkbox next to any storage array that you want to add to your management domain.

The system performs a credential check on each array you are adding to the management domain. You might need to resolve any issues with untrusted certificates before proceeding.

5. Click **Next** to proceed to the next step in the wizard.
6. If the storage arrays have valid certificates, go to [Step 3: Provide passwords](#). If any storage arrays do not have valid certificates, the Resolve Self-Signed Certificates dialog box appears; go to [Step 2: Resolve untrusted certificates during discovery](#). If you want to import CA-signed certificates, cancel out of the discovery dialogs and go to [Import certificates for arrays](#).

Step 2: Resolve untrusted certificates during discovery

If necessary, you must resolve any certificate issues before proceeding with the discovery process.

During discovery, if any storage arrays show an "Untrusted Certificates" status, the Resolve Self-Signed

Certificates dialog box appears. You can resolve untrusted certificates in this dialog, or you can import CA certificates (see [Import certificates for arrays](#)).

Steps

1. If the Resolve Self-Signed Certificates dialog box opens, review the information displayed for the untrusted certificates. For more information, you can also click the ellipses at the far end of the table and select **View** from the pop-up menu.
2. Do one of the following:
 - If you trust the connections to the discovered storage arrays, click **Next** and then click **Yes** to confirm and continue to the next card in the wizard. The self-signed certificates will be marked as trusted and the storage arrays will be added to the plugin.
 - If you do not trust the connections to the storage arrays, select **Cancel** and validate each storage array's security certificate strategy before adding any of them to the plugin.

Step 3: Provide passwords

As the last step for discovery, you must enter the passwords for the storage arrays that you want to add to your management domain.

Steps

1. Optionally, if you have previously configured groups for the arrays, you can use the drop-down to select a group for the discovered arrays.
2. For each discovered array, enter its admin password in the fields.
3. Click **Finish**.



It can take several minutes for the system to connect to the specified storage arrays.

Result

The storage arrays are added to your management domain and associated with the selected group (if specified).



You can use the Launch option to open the browser-based System Manager for one or more storage arrays when you want to perform management operations.

Rename storage array

You can change the storage array's name displayed on the Manage page of the Storage Plugin for vCenter.

Steps

1. From the **Manage** page, select the checkbox to the left of the storage array name.
2. Select the ellipses at the far right of the row, and then select **Rename storage array** from the pop-up menu.
3. Enter the new name and click **Save**.

Change storage array passwords

You can update the passwords used for viewing and accessing storage arrays in the Storage Plugin for vCenter.

Before you begin

You must know the current password for the storage array, which is set in System Manager.

About this task

In this task, you enter the current password for a storage array so you can access it in the plugin. This might be necessary if the array password was changed in System Manager.

Steps

1. From the **Manage** page, select one or more storage arrays.
2. Select **Uncommon Tasks** > **Provide storage array passwords**.
3. Enter the password or passwords for each storage array, and then click **Save**.

Remove storage arrays

You can remove one or more storage arrays if you no longer want to manage it from the Storage Plugin for vCenter.

About this task

You cannot access any of the storage arrays you remove. You can, however, establish a connection to any of the removed storage arrays by pointing a browser directly to its IP address or host name.

Removing a storage array does not affect the storage array or its data in any way. If a storage array is accidentally removed, it can be added again.

Steps

1. From the **Manage** page, select one or more storage arrays that you want to remove.
2. Select **Uncommon Tasks** > **Remove storage arrays**.

The storage array is removed from all the views in the plugin interface.

Launch System Manager

To manage a single array, use the Launch option to open SANtricity System Manager in a new browser window.

System Manager is an embedded application on the storage array's controller, which is connected to the network through an Ethernet management port. System Manager includes all array-based functions. To access System Manager, you must have an out-of-band connection to a network management client with a web browser.

Steps

1. From the **Manage** page, select one or more storage arrays that you want to manage.
2. Click **Launch**.

The system opens a new tab in the browser, and then displays the System Manager login page.

3. Enter your username and password, and then click **Log in**.

Import settings

Import settings overview

The Import Settings feature is a batch operation that allows you to replicate the settings in a single storage array (the source) to multiple arrays (the targets) in the Storage Plugin for vCenter.

Settings available for import

The following configurations can be imported from one array to another:

- **Alerts** — Alerting methods to send important events to administrators using email, a syslog server, or an SNMP server.
- **AutoSupport** — A feature that monitors the health of a storage array and sends automatic dispatches to technical support.
- **Directory services** — A method of user authentication that is managed through an LDAP (Lightweight Directory Access Protocol) server and directory service, such as Microsoft's Active Directory.
- **System settings** — Configurations relating to the following:
 - Media scan settings for a volume
 - SSD settings
 - Automatic load balancing (does not include host connectivity reporting)
- **Storage configuration** — Configurations relating to the following:
 - Volumes (thick and non-repository volumes only)
 - Volume groups and pools
 - Hot spare drive assignments

Configuration workflow

To import settings, follow this workflow:

1. On a storage array to be used as the source, configure the settings using System Manager.
2. On the storage arrays to be used as the targets, back up their configuration using System Manager.
3. From the plugin interface, go to the **Manage** page and import the settings.
4. From the Operations page, review the results of the Import Settings operation.

Requirements for replicating storage configurations

Before importing a storage configuration from one storage array to another, review the requirements and guidelines.

Shelves

- The shelves where the controllers reside must be identical on the source and target arrays.
- Shelf IDs must be identical on the source and target arrays.
- Expansion shelves must be populated in the same slots with the same drive types (if the drive is used in the configuration, the location of unused drives does not matter).

Controllers

- The controller type can be different between the source and target arrays, but the RBOD enclosure type must be identical.
- The HICs, including the DA capabilities of the host, must be identical between the source and target arrays.
- Importing from a duplex to simplex configuration is not supported; however, importing from simplex to duplex is allowed.
- FDE settings are not included in the import process.

Status

- The target arrays must be in Optimal status.
- The source array does not need to be in Optimal status.

Storage

- Drive capacity may vary between the source and target arrays, as long as the volume capacity on the target is larger than the source. (A target array might have newer, larger capacity drives that would not be fully configured into volumes by the replication operation.)
- Disk pool volumes 64 TB or larger on the source array will prevent the import process on the targets.

Import alert settings

You can import alert configurations from one storage array to other storage arrays. This batch operation saves time when you need to configure multiple arrays in the network.

Before you begin

Make sure that:

- Alerts are configured in System Manager (**Settings** > **Alerts**) for the storage array you want to use as the source.
- The existing configuration for the target storage arrays are backed up in System Manager (**Settings** > **System** > **Save Storage Array Configuration**).
- You have reviewed the requirements for replicating storage configurations in [Import settings overview](#).

About this task

You can select email, SNMP, or syslog alerts for the import operation:

- **Email alerts** — A mail server address and the email addresses of the alert recipients.
- **Syslog alerts** — A syslog server address and a UDP port.
- **SNMP alerts** — A community name and IP address for the SNMP server.

Steps

1. From the Manage page, click **Actions** > **Import Settings**.

The Import Settings wizard opens.

2. In the Select Settings dialog, select either **Email alerts**, **SNMP alerts**, or **Syslog alerts**, and then click **Next**.

A dialog box opens for selecting the source array.

3. In the Select Source dialog, select the array with the settings you want to import, and then click **Next**.
4. In the Select Targets dialog, select one or more arrays to receive the new settings.



Storage arrays with firmware below 8.50 are not available for selection. In addition, an array does not appear in this dialog if the plugin cannot communicate with that array (for example, if it is offline or if it has certificate, password, or networking problems).

5. Click **Finish**.

The Operations page displays the results of the import operation. If the operation fails, you can click on its row to see more information.

Result

The target storage arrays are now configured to send alerts to administrators through email, SNMP, or syslog.

Import AutoSupport settings

You can import an AutoSupport configuration from one storage array to other storage arrays. This batch operation saves time when you need to configure multiple arrays in the network.

Before you begin

Make sure that:

- AutoSupport is configured in System Manager (**Support › Support Center**) for the storage array you want to use as the source.
- The existing configuration for the target storage arrays are backed up in System Manager (**Settings › System › Save Storage Array Configuration**).
- You have reviewed the requirements for replicating storage configurations in [Import settings overview](#).

About this task

Imported settings include the separate features (Basic AutoSupport, AutoSupport OnDemand, and Remote Diagnostics), the maintenance window, delivery method, and dispatch schedule.

Steps

1. From the Manage page, click **Actions › Import Settings**.

The Import Settings wizard opens.

2. In the Select Settings dialog, select **AutoSupport** and then click **Next**.

A dialog box opens for selecting the source array.

3. In the Select Source dialog, select the array with the settings you want to import, and then click **Next**.
4. In the Select Targets dialog, select one or more arrays to receive the new settings.



Storage arrays with firmware below 8.50 are not available for selection. In addition, an array does not appear in this dialog if the plugin cannot communicate with that array (for example, if it is offline or if it has certificate, password, or networking problems).

5. Click **Finish**.

The Operations page displays the results of the import operation. If the operation fails, you can click on its row to see more information.

Result

The target storage arrays are now configured with the same AutoSupport settings as the source array.

Import directory services settings

You can import a directory services configuration from one storage array to other storage arrays. This batch operation saves time when you need to configure multiple arrays in the network.

Before you begin

Make sure that:

- Directory services are configured in System Manager (**Settings** > **Access Management**) for the storage array you want to use as the source.
- The existing configuration for the target storage arrays are backed up in System Manager (**Settings** > **System** > **Save Storage Array Configuration**).
- You have reviewed the requirements for replicating storage configurations in [Import settings overview](#).

About this task

Imported settings include the domain name and URL of an LDAP (Lightweight Directory Access Protocol) server, along with the mappings for the LDAP server's user groups to the storage array's predefined roles.

Steps

1. From the Manage page, click **Actions** > **Import Settings**.

The Import Settings wizard opens.

2. In the Select Settings dialog, select **Directory services** and then click **Next**.

A dialog box opens for selecting the source array.

3. In the Select Source dialog, select the array with the settings you want to import, and then click **Next**.
4. In the Select Targets dialog, select one or more arrays to receive the new settings.



Storage arrays with firmware below 8.50 are not available for selection. In addition, an array does not appear in this dialog if the plugin cannot communicate with that array (for example, if it is offline or if it has certificate, password, or networking problems).

5. Click **Finish**.

The Operations page displays the results of the import operation. If the operation fails, you can click on its

row to see more information.

Result

The target storage arrays are now configured with the same directory services as the source array.

Import system settings

You can import the system settings from one storage array to other storage arrays. This batch operation saves time when you need to configure multiple arrays in the network.

Before you begin

Make sure that:

- System settings are configured in System Manager for the storage array you want to use as the source.
- The existing configuration for the target storage arrays are backed up in System Manager (**Settings** > **System** > **Save Storage Array Configuration**).
- You have reviewed the requirements for replicating storage configurations in [Import settings overview](#).

About this task

Imported settings include media scan settings for a volume, SSD settings for controllers, and automatic load balancing (does not include host connectivity reporting).

Steps

1. From the Manage page, click **Actions** > **Import Settings**.

The Import Settings wizard opens.

2. In the Select Settings dialog, select **System** and then click **Next**.

A dialog box opens for selecting the source array.

3. In the Select Source dialog, select the array with the settings you want to import, and then click **Next**.
4. In the Select Targets dialog, select one or more arrays to receive the new settings.



Storage arrays with firmware below 8.50 are not available for selection. In addition, an array does not appear in this dialog if the plugin cannot communicate with that array (for example, if it is offline or if it has certificate, password, or networking problems).

5. Click **Finish**.

The Operations page displays the results of the import operation. If the operation fails, you can click on its row to see more information.

Result

The target storage arrays are now configured with the same system settings as the source array.

Import storage configuration settings

You can import the storage configuration from one storage array to other storage arrays. This batch operation saves time when you need to configure multiple arrays in the

network.

Before you begin

Make sure that:

- Storage is configured in System Manager for the storage array you want to use as the source.
- The existing configuration for the target storage arrays are backed up in System Manager (**Settings > System > Save Storage Array Configuration**).
- You have reviewed the requirements for replicating storage configurations in [Import settings overview](#).
- The source and target arrays must meet these requirements:
 - The shelves where the controllers reside must be identical.
 - Shelf IDs must be identical.
 - Expansion shelves must be populated in the same slots with the same drive types.
 - The RBOD enclosure type must be identical.
 - The HICs, including the Data Assurance capabilities of the host, must be identical.
 - The target arrays must be in Optimal status.
 - The volume capacity on the target array is larger than the source array's capacity.
- You understand the following restrictions:
 - Importing from a duplex to simplex configuration is not supported; however, importing from simplex to duplex is allowed.
 - Disk pool volumes 64 TB or larger on the source array will prevent the import process on the targets.

About this task

Imported settings include configured volumes (thick and non-repository volumes only), volume groups, pools, and hot spare drive assignments.

Steps

1. From the Manage page, click **Actions > Import Settings**.

The Import Settings wizard opens.

2. In the Select Settings dialog, select **Storage configuration** and then click **Next**.

A dialog box opens for selecting the source array.

3. In the Select Source dialog, select the array with the settings you want to import, and then click **Next**.
4. In the Select Targets dialog, select one or more arrays to receive the new settings.



Storage arrays with firmware below 8.50 are not available for selection. In addition, an array does not appear in this dialog if the plugin cannot communicate with that array (for example, if it is offline or if it has certificate, password, or networking problems).

5. Click **Finish**.

The Operations page displays the results of the import operation. If the operation fails, you can click on its row to see more information.

Result

The target storage arrays are now configured with the same storage configuration as the source array.

Manage array groups

Array groups overview

You can manage your physical and virtualized infrastructure in the Storage Plugin for vCenter by grouping a set of storage arrays. You might want to group storage arrays to make it easier to run monitoring or reporting jobs.

Types of storage array groups:

- **All group** — The All group is the default group and includes all the storage arrays discovered in your organization. The All group can be accessed from the main view.
- **User-created group** — A user-created group includes the storage arrays that you manually select to add to that group. User-created groups can be accessed from the main view.

Create storage array group

You create storage groups, and then add storage arrays to the groups. The storage group defines which drives provide the storage that makes up the volume.

Steps

1. From the Manage page, select **Manage Groups > Create storage array group**.
2. In the **Name** field, type a name for the new group.
3. Select the storage arrays that you want to add to the new group.
4. Click **Create**.

Add storage array to group

You can add one or more storage arrays to a user-created group.

Steps

1. From the main view, select **Manage**, and then select the group that you want to add storage arrays to.
2. Select **Manage Groups > Add storage arrays to group**.
3. Select the storage arrays that you want to add to the group.
4. Click **Add**.

Rename storage array group

You can change the name of a storage array group when the current name is no longer meaningful or applicable.

About this task

Keep these guidelines in mind.

- A name can consist of letters, numbers, and the special characters underscore (_), hyphen (-), and pound (#). If you choose any other characters, an error message appears. You are prompted to choose another name.
- Limit the name to 30 characters. Any leading and trailing spaces in the name are deleted.
- Use a unique, meaningful name that is easy to understand and remember.
- Avoid arbitrary names or names that would quickly lose their meaning in the future.

Steps

1. From the main view, select **Manage**, and then select the storage array group you want to rename.
2. Select **Manage Groups** › **Rename storage array group**.
3. In the **Group Name** field, type a new name for the group.
4. Click **Rename**.

Remove storage arrays from group

You can remove one or more managed storage arrays from a group if you no longer want to manage it from a specific storage group.

About this task

Removing storage arrays from a group does not affect the storage array or its data in any way. If your storage array is managed by System Manager, you can still manage it using your browser. If a storage array is accidentally removed from a group, it can be added again.

Steps

1. From the Manage page, select **Manage Groups** › **Remove storage arrays from group**.
2. From the drop-down, select the group that contains the storage arrays you want to remove, and then click the check box next to each storage array that you want to remove from the group.
3. Click **Remove**.

Delete storage array group

You can remove one or more storage array groups that are no longer needed.

About this task

This operation deletes only the storage array group. Storage arrays associated with the deleted group remain accessible through the Manage All view or any other group it is associated with.

Steps

1. From the Manage page, select **Manage Groups** › **Delete storage array group**.
2. Select one or more storage array groups that you want to delete.
3. Click **Delete**.

Upgrade OS software

Upgrade overview

In the Storage Plugin for vCenter, you can manage SANtricity software and NVSRAM

upgrades for multiple storage arrays of the same type.

Upgrade workflow

The following steps provide a high-level workflow for performing software upgrades:

1. You download the latest SANtricity OS file from the support site (a link is available from the Support page). Save the file on the management host system (the host where you access the plugin in a browser), and then unzip the file.
2. In the plugin, you can load the SANtricity OS software file and the NVSRAM file into the repository (an area of the server where files are stored).
3. After the files are loaded in the repository, you can then select the file to be used in the upgrade. From the Upgrade SANtricity OS Software page, you select the OS software file and the NVSRAM file. After you select a software file, a list of compatible storage arrays appear on this page. You then select the storage arrays that you want to upgrade with the new software. (You cannot select incompatible arrays.)
4. You can then begin an immediate software transfer and activation, or you can choose to stage the files for activation at a later time. During the upgrade process, the plugin performs the following tasks:
 - Performs a health check on the storage arrays to determine if any conditions exist that might prevent the upgrade from completing. If any arrays fail the health check, you can skip that particular array and continue the upgrade for the others, or you can stop the entire process and troubleshoot the arrays that did not pass.
 - Transfers the upgrade files to each controller.
 - Reboots the controllers and activates the new OS software, one controller at a time. During activation, the existing OS file is replaced with the new file.



You can also specify that the software is activated at a later time.

Upgrade considerations

Before you upgrade multiple storage arrays, review the key considerations as part of your planning.

Current versions

You can view the current SANtricity OS software versions from the Manage page of the Storage Plugin for vCenter for each discovered storage array. The version is shown in the SANtricity OS Software column. The controller firmware and NVSRAM information is available in a pop-up dialog box when you click on the OS version in each row.

Other components requiring upgrade

As part of the upgrade process, you might also need to upgrade the host's multipath/failover driver or the HBA driver so that the host can interact with the controllers correctly. For compatibility information, refer to the [Interoperability Matrix Tool](#).

Dual controllers

If a storage array contains two controllers and you have a multipath driver installed, the storage array can continue to process I/O while the upgrade occurs. During the upgrade, the following process occurs:

1. Controller A fails over all its LUNs to controller B.
2. Upgrade occurs on controller A.

3. Controller A takes back its LUNs and all of controller B's LUNs.
4. Upgrade occurs on controller B.

After the upgrade completes, you might need to manually redistribute volumes between the controllers to ensure volumes return to the correct owning controller.

Perform pre-upgrade health check

A health check runs as part of the upgrade process, but you also can run a health check separately before you begin. The health check assesses components of the storage array to make sure that the upgrade can proceed.

Steps

1. From the main view, select **Manage**, and then select **Upgrade Center** > **Pre-Upgrade Health Check**.

The Pre-Upgrade Health Check dialog box opens and lists all the discovered storage systems.

2. If needed, filter or sort the storage systems in the list, so you can view all systems that are not currently in the Optimal state.
3. Select the check boxes for the storage systems that you want to run through the health check.
4. Click **Start**.

The progress is shown in the dialog box while the health check is performed.

5. When the health check completes, you can click on the ellipses (...) to the right of each row to view more information and perform other tasks.



If any arrays fail the health check, you can skip that particular array and continue the upgrade for the others, or you can stop the entire process and troubleshoot the arrays that did not pass.

Upgrade SANtricity OS

Upgrade one or more storage arrays with the latest software and NVSRAM to make sure that you have all the latest features and bug fixes. Controller NVSRAM is a controller file that specifies the default settings for the controllers.

Before you begin

Make sure that:

- The latest SANtricity OS files are available on the host system where the plugin is running.
- You know whether you want to activate your software upgrade now or later. You might choose to activate later for these reasons:
 - **Time of day** — Activating the software can take a long time, so you might want to wait until I/O loads are lighter. The controllers fail over during activation, so performance might be lower than usual until the upgrade completes.
 - **Type of package** — You might want to test the new OS software on one storage array before you upgrade the files on other storage arrays.



Risk of data loss or risk of damage to the storage array — Do not make changes to the storage array while the upgrade is occurring. Maintain power to the storage array.

Steps

1. If your storage array contains only one controller or a multipath driver is not in use, stop I/O activity to the storage array to prevent application errors. If your storage array has two controllers and you have a multipath driver installed, you do not need to stop I/O activity.
2. From the main view, select **Manage**, and then select one or more storage arrays that you want to upgrade.
3. Select **Upgrade Center > Upgrade > SANtricity OS > Software**.

The Upgrade SANtricity OS software page appears.

4. Download the latest SANtricity OS software package from the Support site to your local machine.
 - a. Click Add new file to software repository
 - b. Click the link for finding the latest SANtricity OS downloads.
 - c. Click the **Download Latest Release** link.
 - d. Follow the remaining instructions to download the OS file and the NVSRAM file to your local machine.



Digitally signed firmware is required in version 8.42 and above. If you attempt to download unsigned firmware, an error is displayed and the download is aborted.

5. Select the OS software file and the NVSRAM file that you want to use to upgrade the controllers:
 - a. From the drop-down, select the OS file that you downloaded to your local machine.

If there are multiple files available, the files are sorted from newest date to oldest date.



The software repository lists all software files associated with the plugin. If you do not see the file that you want to use, you can click the link, **Add new file to software repository**, to browse to the location where the OS file that you want to add resides.

- b. From the **Select an NVSRAM file** drop-down, select the controller file that you want to use.

If there are multiple files, the files are sorted from newest date to oldest date.

6. In the Compatible Storage Array table, review the storage arrays that are compatible with the OS software file that you selected, and then select the arrays you want to upgrade.
 - The storage arrays that you selected in the Manage view and that are compatible with the selected firmware file are selected by default in the Compatible Storage Array table.
 - The storage arrays that cannot be updated with the selected firmware file are not selectable in the Compatible Storage Array table as indicated by the status **Incompatible**.
7. (Optional) To transfer the software file to the storage arrays without activating them, select the **Transfer the OS software to the storage arrays, mark it as staged, and activate at a later time** check box.
8. Click **Start**.
9. Depending on whether you chose to activate now or later, do one of the following:
 - Type **TRANSFER** to confirm that you want to transfer the proposed OS software versions on the arrays you selected to upgrade, and then click **Transfer**. To activate the transferred software, select **Upgrade Center > Activate Staged SANtricity OS Software**.

- Type **UPGRADE** to confirm that you want to transfer and activate the proposed OS software versions on the arrays you selected to upgrade, and then click **Upgrade**.

The system transfers the software file to each storage array you selected to upgrade and then activates that file by initiating a reboot.

The following actions occur during the upgrade operation:

- A pre-upgrade health check runs as part of the upgrade process. The pre-upgrade health check assesses all storage array components to make sure that the upgrade can proceed.
 - If any health check fails for a storage array, the upgrade stops. You can click the ellipsis (...) and select **Save Log** to review the errors. You can also choose to override the health check error and then click **Continue** to proceed with the upgrade.
 - You can cancel the upgrade operation after the pre-upgrade health check.
10. (Optional) Once the upgrade has completed, you can see a list of what was upgraded for a specific storage array by clicking the ellipsis (...) and then selecting **Save Log**.

The file is saved in the Downloads folder for your browser with the name `upgrade_log-<date>.json`.

Activate staged OS software

You can choose to activate the software file immediately or wait until a more convenient time. This procedure assumes you chose to activate the software file at a later time.

About this task

You can transfer the firmware files without activating them. You might choose to activate later for these reasons:

- **Time of day** — Activating the software can take a long time, so you might want to wait until I/O loads are lighter. The controllers reboot and fail over during activation so performance might be lower than usual until the upgrade completes.
- **Type of package** — You might want to test the new software and firmware on one storage array before upgrading the files on other storage arrays.



You cannot stop the activation process after it starts.

Steps

1. From the main view, select **Manage**. If necessary, click the **Status** column to sort, at the top of the page, all storage arrays with a status of "OS Upgrade (awaiting activation)."
2. Select one or more storage arrays that you want to activate software for, and then select **Upgrade Center** > **Activate Staged SANtricity Software**.

The following actions occur during the upgrade operation:

- A pre-upgrade health check runs as part of the activate process. The pre-upgrade health check assesses all storage array components to make sure that the activation can proceed.
- If any health check fails for a storage array, the activation stops. You can click the ellipsis (...) and select **Save Log** to review the errors. You can also choose to override the health check error and then click **Continue** to proceed with the activation.

- You can cancel the activate operation after the pre-upgrade health check.

On successful completion of the pre-upgrade health check, activation occurs. The time it takes to activate depends on your storage array configuration and the components that you are activating.

3. (Optional) After the activation is complete, you can see a list of what was activated for a specific storage array by clicking the ellipsis (...) and then selecting **Save Log**.

The file is saved in the Downloads folder for your browser with the name `activate_log-<date>.json`.

Clear staged OS software

You can remove staged OS software to ensure that a pending version is not inadvertently activated at a later time. Removing the staged OS software does not affect the current version that is running on the storage arrays.

Steps

1. From the main view, select **Manage**, and then select **Upgrade Center > Clear Staged SANtricity Software**.

The Clear Staged SANtricity Software dialog box opens and lists all the discovered storage systems with pending software or NVSRAM.

2. If needed, filter or sort the storage systems in the list, so you can view all systems that have staged software.
3. Select the check boxes for the storage systems with pending software that you want cleared.
4. Click **Clear**.

The status of the operation is shown in the dialog box.

Manage software repository

You can view and manage a software repository, which lists all software files associated with the Storage Plugin for vCenter.

Before you begin

If you are using the repository to add SANtricity OS files, make sure that the OS files are available on your local system.

About this task

You can use the Manage SANtricity OS Software Repository option to import one or more OS files to the host system where the plugin is running. You can also choose to delete one or more OS files that are available in the software repository.

Steps

1. From the main view, select **Manage**, and then select **Upgrade Center > Manage SANtricity Software Repository**.

The Manage SANtricity OS Software Repository dialog appears.

2. Perform one of the following actions:

- **Import:**
 - a. Click **Import**.
 - b. Click **Browse**, and then navigate to the location where the OS files you want to add reside. OS files have a filename similar to `N2800-830000-000.dlp`.
 - c. Select one or more OS files that you want to add, and then click **Import**.
- **Delete:**
 - a. Select one or more OS files that you want to remove from the software repository.
 - b. Click **Delete**.

Result

If you selected import, the file(s) are uploaded and validated. If you selected delete, the files are removed from the software repository.

Provision storage

Provisioning overview

In the Storage Plugin for vCenter, you can create data containers, called volumes, so the host can access storage on the array.

Volume types and characteristics

Volumes are data containers that manage and organize the storage space on your storage array.

You create volumes from the storage capacity available on your storage array, which helps organize your system's resources. The concept of "volumes" is similar to using folders/directories on a computer to organize files for quick access.

Volumes are the only data layer visible to hosts. In a SAN environment, volumes are mapped to logical unit numbers (LUNs). These LUNs hold the user data that is accessible using one or more of the host access protocols supported by the storage array, including FC, iSCSI, and SAS.

Each volume in a pool or volume group can have its own individual characteristics based on what type of data will be stored in it. Some of these characteristics include:

- **Segment size** — A segment is the amount of data in kilobytes (KiB) that is stored on a drive before the storage array moves to the next drive in the stripe (RAID group). The segment size is equal to or less than the capacity of the volume group. The segment size is fixed and cannot be changed for pools.
- **Capacity** — You create a volume from the free capacity available in either a pool or volume group. Before you create a volume, the pool or volume group must already exist, and it must have enough free capacity to create the volume.
- **Controller ownership** — All storage arrays can have either one or two controllers. On a single- controller array, a volume's workload is managed by a single controller. On a dual-controller array, a volume will have a preferred controller (A or B) that "owns" the volume. In a dual- controller configuration, volume ownership is automatically adjusted using the Automatic Load Balancing feature to correct any load balance issues when workloads shift across the controllers. Automatic load balancing provides automated I/O workload balancing and ensures that incoming I/O traffic from the hosts is dynamically managed and balanced across both controllers.

- **Volume assignment** — You can give hosts access to a volume either when you create the volume or at a later time. All host access is managed through a logical unit number (LUN). Hosts detect LUNs that are, in turn, assigned to volumes. If you are assigning a volume to multiple hosts, use clustering software to make sure that the volume is available to all of the hosts.

The host type can have specific limits on how many volumes the host can access. Keep this limitation in mind when you create volumes for use by a particular host.

- **Resource provisioning** — For EF600 or EF300 storage arrays, you can specify that volumes be put in use immediately with no background initialization process. A resource-provisioned volume is a thick volume in an SSD volume group or pool, where drive capacity is allocated (assigned to the volume) when the volume is created, but the drive blocks are deallocated (unmapped).
- **Descriptive name** — You can name a volume whatever name you like, but we recommend making the name descriptive.

During volume creation, each volume is allocated capacity and is assigned a name, segment size (volume groups only), controller ownership, and volume-to-host assignment. Volume data is automatically load balanced across controllers, as needed.

Capacity for volumes

The drives in your storage array provide the physical storage capacity for your data. Before you can begin storing data, you must configure the allocated capacity into logical components known as pools or volume groups. You use these storage objects to configure, store, maintain, and preserve data on your storage array.

Capacity to create and expand volumes

You can create volumes from either the unassigned capacity or free capacity in a pool or volume group.

- When you create a volume from unassigned capacity, you can create a pool or volume group and the volume at the same time.
- When you create a volume from free capacity, you are creating an additional volume on an already existing pool or volume group. After you expand the volume capacity, you must manually increase the file system size to match. How you do this depends on the file system you are using. See your host operating system documentation for details.



The plugin interface does not provide an option to create thin volumes.

Reported capacity for volumes

The reported capacity of the volume is equal to the amount of physical storage capacity allocated. The entire amount of physical storage capacity must be present. The physically allocated space is equal to the space that is reported to the host.

You normally set the volume's reported capacity to be the maximum capacity to which you think the volume will grow. Volumes provide high and predictable performance for your applications mainly because all of the user capacity is reserved and allocated upon creation.

Capacity limits

The minimum capacity for a volume is 1 MiB, and the maximum capacity is determined by the number and capacity of the drives in the pool or volume group.

When increasing reported capacity for a volume, keep the following guidelines in mind:

- You can specify up to three decimal places (for example, 65.375 GiB).
- Capacity needs to be less than (or equal to) the maximum available in the volume group. When you create a volume, some additional capacity is pre-allocated for Dynamic Segment Size (DSS) migration. DSS migration is a feature of the software that allows you to change the segment size of a volume.
- Volumes larger than 2 TiB are supported by some host operating systems (maximum reported capacity is determined by the host operating system). In fact, some host operating systems support up to 128 TiB volumes. Refer to your host operating system documentation for additional details.

Application-specific workloads

When creating a volume, you select a workload to customize the storage array configuration for a specific application.

A workload is a storage object that supports an application. You can define one or more workloads, or instances, per application. For some applications, the system configures the workload to contain volumes with similar underlying volume characteristics. These volume characteristics are optimized based on the type of application the workload supports. For example, if you create a workload that supports a Microsoft SQL Server application and then subsequently create volumes for that workload, the underlying volume characteristics are optimized to support Microsoft SQL Server.

During volume creation, the system prompts you to answer questions about a workload's use. For example, if you are creating volumes for Microsoft Exchange, you are asked how many mailboxes you need, what your average mailbox capacity requirements are, and how many copies of the database you want. The system uses this information to create an optimal volume configuration for you, which can be edited as needed. Optionally, you can skip this step in the volume creation sequence.

Types of workloads

You can create two types of workloads: application-specific and other.

- **Application-specific** — When you are creating volumes using an application-specific workload, the system may recommend an optimized volume configuration to minimize contention between application workload I/O and other traffic from your application instance. Volume characteristics like I/O type, segment size, controller ownership, and read and write cache are automatically recommended and optimized for workloads that are created for the following application types.
 - Microsoft SQL Server
 - Microsoft Exchange Server
 - Video Surveillance applications
 - VMware ESXi (for volumes to be used with Virtual Machine File System)

You can review the recommended volume configuration and edit, add, or delete the system-recommended volumes and characteristics using the Add/Edit Volumes dialog box.

- **Other (or applications without specific volume creation support)** — Other workloads use a volume configuration that you must manually specify when you want to create a workload that is not associated with a specific application, or if the system does not have built-in optimization for the application you intend to use on the storage array. You must manually specify the volume configuration using the Add/Edit Volumes dialog box.

Application and workload views

To view applications and workloads, launch System Manager. From that interface, you can view information associated with an application-specific workload in a couple of different ways:

- You can select the Applications & Workloads tab in the Volumes tile to view the storage array's volumes grouped by workload and the application type the workload is associated with.
- You can select the Applications & Workloads tab in the Performance tile to view performance metrics (latency, IOPS, and MBs) for logical objects. The objects are grouped by application and associated workload. By collecting this performance data at regular intervals, you can establish baseline measurements and analyze trends, which can help as you investigate problems related to I/O performance.

Create storage

In the Storage Plugin for vCenter, you create storage by first creating a workload for a specific application type. Next, you add storage capacity to the workload by creating volumes with similar underlying volume characteristics.

Step 1: Create workloads

A workload is a storage object that supports an application. You can define one or more workloads, or instances, per application.

About this task

For some applications, the system configures the workload to contain volumes with similar underlying volume characteristics. These volume characteristics are optimized based on the type of application the workload supports. For example, if you create a workload that supports a Microsoft SQL Server application and then subsequently create volumes for that workload, the underlying volume characteristics are optimized to support Microsoft SQL Server.

The system recommends an optimized volume configuration only for the following application types:

- Microsoft SQL Server
- Microsoft Exchange Server
- Video Surveillance
- VMware ESXi (for volumes to be used with Virtual Machine File System)

Steps

1. From the Manage page, select the storage array.
2. Select **Provisioning > Manage Volumes**.
3. Select **Create > Workload**.

The Create Application Workload dialog box appears.

4. Use the drop-down list to select the type of application that you want to create the workload for and then type a workload name.
5. Click **Create**.

Step 2: Create volumes

You create volumes to add storage capacity to an application-specific workload, and to make the created volumes visible to a specific host or host cluster.

About this task

Most application types default to a user-defined volume configuration, while other types have a smart configuration applied at volume creation. For example, if you are creating volumes for a Microsoft Exchange application, you are asked how many mailboxes you need, what your average mailbox capacity requirements are, and how many copies of the database you want. The system uses this information to create an optimal volume configuration for you, which can be edited as needed.

You can create volumes from **Provisioning › Manage Volumes › Create › Volumes** or from **Provisioning › Configure Pools and Volume Groups › Create › Volumes**. The procedure is the same for either selection.

The process to create a volume is a multi-step procedure.

Step 2a: Select host for a volume

In the first step, you can select a specific host or host cluster for the volume, or you can choose to assign the host later.

Before you begin

Make sure that:

- Valid hosts or host clusters have been defined (go to **Provisioning › Configure Hosts**).
- Host port identifiers have been defined for the host.
- The host connection must support Data Assurance (DA) if you plan to create DA-enabled volumes. If any of the host connections on the controllers in your storage array do not support DA, the associated hosts cannot access data on DA-enabled volumes.

About this task

Keep these guidelines in mind when you assign volumes:

- A host's operating system can have specific limits on how many volumes the host can access. Keep this limitation in mind when you create volumes for use by a particular host.
- You can define one assignment for each volume in the storage array.
- Assigned volumes are shared between controllers in the storage array.
- The same logical unit number (LUN) cannot be used twice by a host or a host cluster to access a volume. You must use a unique LUN.
- If you want to speed the process for creating volumes, you can skip the host assignment step so that newly created volumes are initialized offline.



Assigning a volume to a host will fail if you try to assign a volume to a host cluster that conflicts with an established assignment for a host in the host clusters.

Steps

1. From the Manage page, select the storage array.
2. Select **Provisioning › Manage Volumes**.

3. Select **Create > Volumes**.

The Select Host dialog box appears.

4. From the drop-down list, select a specific host or host cluster to which you want to assign volumes, or choose to assign the host or host cluster at a later time.
5. To continue the volume creation sequence for the selected host or host cluster, click **Next**.

The Select Workload dialog box appears.

Step 2b: Select a workload for a volume

In the second step, you select a workload to customize the storage array configuration for a specific application, such as VMware.

About this task

This task describes how to create volumes for a workload. Typically, a workload contains volumes with similar characteristics, which are optimized based on the type of application the workload supports. You can define a workload in this step or you can select existing workloads.

Keep these guidelines in mind:

- When using an application-specific workload, the system recommends an optimized volume configuration to minimize contention between application workload I/O and other traffic from your application instance. You can review the recommended volume configuration, and then edit, add, or delete the system-recommended volumes and characteristics using the Add/Edit Volumes dialog box (available in the next step).
- When using other application types, you manually specify the volume configuration using the Add/Edit Volumes dialog box (available in the next step).

Steps

1. Do one of the following:
 - Select the **Create volumes for an existing workload** option and then select the workload from the drop-down list.
 - Select the **Create a new workload** option to define a new workload for a supported application or for "Other" applications, and then following these steps:
 - From the drop-down list, select the name of the application you want to create the new workload for. Select one of the "Other" entries if the application you intend to use on this storage array is not listed.
 - Enter a name for the workload you want to create.
2. Click **Next**.
3. If your workload is associated with a supported application type, enter the information requested; otherwise, go to the next step.

Step 2c: Add or edit volumes

In the third step, you define the volume configuration.

Before you begin

- The pools or volume groups must have sufficient free capacity.

- The maximum number of volumes allowed in a volume group is 256.
- The maximum number of volumes allowed in a pool depends on the storage system model:
 - 2,048 volumes (EF600 and E5700 series)
 - 1,024 volumes (EF300)
 - 512 volumes (E2800 series)
- To create a Data Assurance (DA)-enabled volume, the host connection you are planning to use must support DA.
 - If you want to create a DA-enabled volume, select a pool or volume group that is DA capable (look for **Yes** next to "DA" in the pool and volume group candidates table).
 - DA capabilities are presented at the pool and volume group level. DA protection checks for and corrects errors that might occur as data is transferred through the controllers down to the drives. Selecting a DA-capable pool or volume group for the new volume ensures that any errors are detected and corrected.
 - If any of the host connections on the controllers in your storage array do not support DA, the associated hosts cannot access data on DA-enabled volumes.
- To create a secure-enabled volume, a security key must be created for the storage array.
 - If you want to create a secure-enabled volume, select a pool or volume group that is secure capable (look for Yes next to "Secure-capable" in the pool and volume group candidates table).
 - Drive security capabilities are presented at the pool and volume group level. Secure-capable drives prevent unauthorized access to the data on a drive that is physically removed from the storage array. A secure-enabled drive encrypts data during writes and decrypts data during reads using a unique encryption key.
 - A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.
- To create a resource-provisioned volume, all drives must be NVMe drives with the Deallocated or Unwritten Logical Block Error (DULBE) option.

About this task

You create volumes from eligible pools or volume groups, which are shown in the Add/Edit Volumes dialog box. For each eligible pool and volume group, the number of drives available and the total free capacity appears.

For some application-specific workloads, each eligible pool or volume group shows the proposed capacity based on the suggested volume configuration and shows the remaining free capacity in GiB. For other workloads, the proposed capacity appears as you add volumes to a pool or volume group and specify the reported capacity.

Steps

1. Choose one of these actions based on whether you selected Other or an application-specific workload in the previous step:
 - **Other** — Click **Add new volume** in each pool or volume group that you want to use to create one or more volumes.

Field Details

Field	Description
Volume Name	A volume is assigned a default name during the volume creation sequence. You can either accept the default name or provide a more descriptive one indicating the type of data stored in the volume.
Reported Capacity	<p>Define the capacity of the new volume and the capacity units to use (MiB, GiB, or TiB). For Thick volumes, the minimum capacity is 1 MiB, and the maximum capacity is determined by the number and capacity of the drives in the pool or volume group.</p> <p>Keep in mind that storage capacity is also required for copy services (snapshot images, snapshot volumes, volume copies, and remote mirrors); therefore, do not allocate all of the capacity to standard volumes.</p> <p>Capacity in a pool is allocated in 4GiB increments. Any capacity that is not a multiple of 4GiB is allocated but not usable. To make sure that the entire capacity is usable, specify the capacity in 4GiB increments. If unusable capacity exists, the only way to regain it is to increase the capacity of the volume.</p>
Volume Block Size (EF300 and EF600 only)	<p>Shows the block sizes that can be created for the volume:</p> <ul style="list-style-type: none">• 512 – 512 bytes• 4K – 4,096 bytes

Field	Description
Segment Size	<p>Shows the setting for segment sizing, which only appears for volumes in a volume group. You can change the segment size to optimize performance.</p> <p>Allowed segment size transitions — The system determines the segment size transitions that are allowed. Segment sizes that are inappropriate transitions from the current segment size are unavailable on the drop-down list. Allowed transitions usually are double or half of the current segment size. For example, if the current volume segment size is 32 KiB, a new volume segment size of either 16 KiB or 64 KiB is allowed.</p> <p>SSD Cache-enabled volumes — You can specify a 4-KiB segment size for SSD Cache-enabled volumes. Make sure you select the 4-KiB segment size only for SSD Cache-enabled volumes that handle small-block I/O operations (for example, 16 KiB I/O block sizes or smaller). Performance might be impacted if you select 4 KiB as the segment size for SSD Cache-enabled volumes that handle large block sequential operations.</p> <p>Amount of time to change segment size — The amount of time to change a volume's segment size depends on these variables:</p> <ul style="list-style-type: none"> • The I/O load from the host • The modification priority of the volume • The number of drives in the volume group • The number of drive channels • The processing power of the storage array controllers <p>When you change the segment size for a volume, I/O performance is affected, but your data remains available.</p>
Secure-capable	<p>Yes appears next to “Secure-capable” only if the drives in the pool or volume group are secure-capable.</p> <p>Drive Security prevents unauthorized access to the data on a drive that is physically removed from the storage array. This option is available only when the Drive Security feature has been enabled, and a security key is set up for the storage array.</p> <p>A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.</p>
DA	<p>Yes appears next to “DA” only if the drives in the pool or volume group support Data Assurance (DA).</p> <p>DA increases data integrity across the entire storage system. DA enables the storage array to check for errors that might occur as data is transferred through the controllers down to the drives. Using DA for the new volume ensures that any errors are detected.</p>

Field	Description
Resource provisioned (EF300 and EF600 only)	Yes appears next to “Resource provisioned” only if the drives support this option. Resource Provisioning is a feature available in the EF300 and EF600 storage arrays, which allows volumes to be put in use immediately with no background initialization process.

- **Application-specific workload** — Either click **Next** to accept the system-recommended volumes and characteristics for the selected workload, or click **Edit Volumes** to change, add, or delete the system-recommended volumes and characteristics for the selected workload.

Field Details

Field	Description
Volume Name	A volume is assigned a default name during the volume creation sequence. You can either accept the default name or provide a more descriptive one indicating the type of data stored in the volume.
Reported Capacity	<p>Define the capacity of the new volume and the capacity units to use (MiB, GiB, or TiB). For Thick volumes, the minimum capacity is 1 MiB, and the maximum capacity is determined by the number and capacity of the drives in the pool or volume group.</p> <p>Keep in mind that storage capacity is also required for copy services (snapshot images, snapshot volumes, volume copies, and remote mirrors); therefore, do not allocate all of the capacity to standard volumes.</p> <p>Capacity in a pool is allocated in 4-GiB increments. Any capacity that is not a multiple of 4 GiB is allocated but not usable. To make sure that the entire capacity is usable, specify the capacity in 4-GiB increments. If unusable capacity exists, the only way to regain it is to increase the capacity of the volume.</p>
Volume Type	Volume type indicates the type of volume that was created for an application- specific workload.
Volume Block Size (EF300 and EF600 only)	<p>Shows the block sizes that can be created for the volume:</p> <ul style="list-style-type: none">• 512 — 512 bytes• 4K — 4,096 bytes

Field	Description
Segment Size	<p>Shows the setting for segment sizing, which only appears for volumes in a volume group. You can change the segment size to optimize performance.</p> <p>Allowed segment size transitions — The system determines the segment size transitions that are allowed. Segment sizes that are inappropriate transitions from the current segment size are unavailable on the drop-down list. Allowed transitions usually are double or half of the current segment size. For example, if the current volume segment size is 32 KiB, a new volume segment size of either 16 KiB or 64 KiB is allowed.</p> <p>SSD Cache-enabled volumes — You can specify a 4-KiB segment size for SSD Cache-enabled volumes. Make sure you select the 4-KiB segment size only for SSD Cache-enabled volumes that handle small-block I/O operations (for example, 16 KiB I/O block sizes or smaller). Performance might be impacted if you select 4 KiB as the segment size for SSD Cache-enabled volumes that handle large block sequential operations.</p> <p>Amount of time to change segment size — The amount of time to change a volume's segment size depends on these variables:</p> <ul style="list-style-type: none"> • The I/O load from the host • The modification priority of the volume • The number of drives in the volume group • The number of drive channels • The processing power of the storage array controllers <p>When you change the segment size for a volume, I/O performance is affected, but your data remains available.</p>
Secure-capable	<p>Yes appears next to “Secure-capable” only if the drives in the pool or volume group are secure-capable.</p> <p>Drive security prevents unauthorized access to the data on a drive that is physically removed from the storage array. This option is available only when the drive security feature has been enabled, and a security key is set up for the storage array.</p> <p>A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.</p>
DA	<p>Yes appears next to “DA” only if the drives in the pool or volume group support Data Assurance (DA).</p> <p>DA increases data integrity across the entire storage system. DA enables the storage array to check for errors that might occur as data is transferred through the controllers down to the drives. Using DA for the new volume ensures that any errors are detected.</p>

Field	Description
Resource provisioned (EF300 and EF600 only)	Yes appears next to “Resource Provisioned” only if the drives support this option. Resource Provisioning is a feature available in the EF300 and EF600 storage arrays, which allows volumes to be put in use immediately with no background initialization process.

2. To continue the volume creation sequence for the selected application, click **Next**.

Step 2d: Review volume configuration

In the last step, you review a summary of the volumes you intend to create and make any necessary changes.

Steps

1. Review the volumes you want to create. To make changes, click **Back**.
2. When you are satisfied with your volume configuration, click **Finish**.

After you finish

- In the vSphere Client, create datastores for the volumes.
- Perform any operating system modifications necessary on the application host so that the applications can use the volume.
- Run either the host-based `hot_add` utility or an operating system-specific utility (available from a third-party vendor), and then run the `SMdevices` utility to correlate volume names with host storage array names.

The `hot_add` utility and the `SMdevices` utility are included as part of the `SMutils` package. The `SMutils` package is a collection of utilities to verify what the host sees from the storage array. It is included as part of the SANtricity software installation.

Increase capacity of a volume

You can resize a volume to increase its reported capacity.

Before you begin

Make sure that:

- Enough free capacity is available in the volume’s associated pool or volume group.
- The volume is Optimal and not in any state of modification.
- No hot spare drives are in use in the volume. (Applies only to volumes in volume groups.)

About this task

This task describes how to increase the reported capacity (the capacity reported to hosts) of a volume by using the free capacity that is available in the pool or volume group. Be sure to consider any future capacity requirements that you might have for other volumes in this pool or volume group.



Increasing the capacity of a volume is supported only on certain operating systems. If you increase the volume capacity on a host operating system that is unsupported, the expanded capacity is unusable, and you cannot restore the original volume capacity.

Steps

1. From the **Manage** page, select the storage array that contains the volumes you want to resize.
2. Select **Provisioning > Manage Volumes**.
3. Select the volume for which you want to increase capacity, and then select **Increase Capacity**.

The Confirm Increase Capacity dialog box appears.

4. Select **Yes** to continue.

The Increase Reported Capacity dialog box appears. This dialog box displays the volume's current reported capacity and the free capacity available in the volume's associated pool or volume group.

5. Use the **Increase reported capacity by adding...** box to add capacity to the current available reported capacity. You can change the capacity value to display in either mebibytes (MiB), gibibytes (GiB), or tebibytes (TiB).
6. Click **Increase**.

The volume's capacity is increased based on your selection. Be aware that this operation can be lengthy and could affect system performance.

After you finish

After you expand the volume capacity, you must manually increase the file system size to match. How you do this depends on the file system you are using. See your host operating system documentation for details.

Change settings for a volume

You can change a volume's settings such as its name, host assignment, segment size, modification priority, caching, and so on.

Before you begin

Make sure that the volume you want to change is in Optimal status.

Steps

1. From the Manage page, select the storage array that contains the volumes you want to change.
2. Select **Provisioning > Manage Volumes**.
3. Select the volume that you want to change, and then select **View/Edit Settings**.

The Volume Settings dialog box appears. The configuration settings for the volume you selected appear in this dialog box.

4. Select the **Basic** tab to change the volume's name and host assignment.

Field Details

Setting	Description
Name	Displays the name of the volume. Change the name of a volume when the current name is no longer meaningful or applicable.
Capacities	Displays the reported and allocated capacity for the selected volume.
Pool / Volume group	Displays the name and RAID level of the pool or volume group. Indicates whether the pool or volume group is secure-capable and secure-enabled.
Host	<p>Displays the volume assignment. You assign a volume to a host or host cluster so it can be accessed for I/O operations. This assignment grants a host or host cluster access to a particular volume or to a number of volumes in a storage array.</p> <ul style="list-style-type: none"> • Assigned to — Identifies the host or host cluster that has access to the selected volume. • LUN — A logical unit number (LUN) is the number assigned to the address space that a host uses to access a volume. The volume is presented to the host as capacity in the form of a LUN. Each host has its own LUN address space. Therefore, the same LUN can be used by different hosts to access different volumes. <p>For NVMe interfaces, this column displays Namespace ID. A namespace is NVM storage that is formatted for block access. It is analogous to a logical unit in SCSI, which relates to a volume in the storage array. The namespace ID is the NVMe controller's unique identifier for the namespace, and can be set to a value between 1 and 255. It is analogous to a logical unit number (LUN) in SCSI.</p>
Identifiers	<p>Displays the identifiers for the selected volume.</p> <ul style="list-style-type: none"> • World-wide identifier (WWID). A unique hexadecimal identifier for the volume. • Extended unique identifier (EUI). An EUI-64 identifier for the volume. • Subsystem identifier (SSID). The storage array subsystem identifier of a volume.

5. Select the **Advanced** tab to change additional configuration settings for a volume in a pool or in a volume group.

Field Details

Setting	Description
Application & workload information	<p>During volume creation, you can create application-specific workloads or other workloads. If applicable, the workload name, application type, and volume type appears for the selected volume.</p> <p>You can change the workload name if desired.</p>
Quality of Service settings	<p>Permanently disable data assurance — This setting appears only if the volume is Data Assurance (DA)-enabled. DA checks for and corrects errors that might occur as data is transferred through the controllers down to the drives. Use this option to permanently disable DA on the selected volume. When disabled, DA cannot be re-enabled on this volume.</p> <p>Enable pre-read redundancy check — This setting appears only if the volume is a thick volume. Pre-read redundancy checks determine whether the data on a volume is consistent any time a read is performed. A volume that has this feature enabled returns read errors if the data is determined to be inconsistent by the controller firmware.</p>
Controller ownership	<p>Defines the controller that is designated to be the owning, or primary, controller of the volume.</p> <p>Controller ownership is very important and should be planned carefully. Controllers should be balanced as closely as possible for total I/Os.</p>
Segment sizing	<p>Shows the setting for segment sizing, which appears only for volumes in a volume group. You can change the segment size to optimize performance.</p> <p>Allowed segment size transitions — The system determines the segment size transitions that are allowed. Segment sizes that are inappropriate transitions from the current segment size are unavailable on the drop-down list. Allowed transitions usually are double or half of the current segment size. For example, if the current volume segment size is 32 KiB, a new volume segment size of either 16 KiB or 64 KiB is allowed.</p> <p>SSD Cache-enabled volumes — You can specify a 4-KiB segment size for SSD Cache-enabled volumes. Make sure you select the 4-KiB segment size only for SSD Cache-enabled volumes that handle small-block I/O operations (for example, 16 KiB I/O block sizes or smaller). Performance might be impacted if you select 4 KiB as the segment size for SSD Cache-enabled volumes that handle large block sequential operations.</p> <p>Amount of time to change segment size. The amount of time to change a volume's segment size depends on these variables:</p> <ul style="list-style-type: none"> • The I/O load from the host • The modification priority of the volume • The number of drives in the volume group • The number of drive channels • The processing power of the storage array controllers <p>When you change the segment size for a volume, I/O performance is affected, but your data remains available.</p>

Setting	Description
Modification priority	<p>Shows the setting for modification priority, which only appears for volumes in a volume group.</p> <p>The modification priority defines how much processing time is allocated for volume modification operations relative to system performance. You can increase the volume modification priority, although this might affect system performance.</p> <p>Move the slider bars to select a priority level.</p> <p>Modification priority rates — The lowest priority rate benefits system performance, but the modification operation takes longer. The highest priority rate benefits the modification operation, but system performance might be compromised.</p>
Caching	Shows the caching setting, which you can change to impact the overall I/O performance of a volume.
SSD Cache	<p>(This feature is not available on the EF600 or EF300 storage system.)</p> <p>Shows the SSD Cache setting, which you can enable on compatible volumes as a way to improve read-only performance. Volumes are compatible if they share the same drive security and data assurance capabilities.</p> <p>The SSD Cache feature uses a single or multiple solid-state disks (SSDs) to implement a read cache. Application performance is improved because of the faster read times for SSDs. Because the read cache is in the storage array, caching is shared across all applications using the storage array.</p> <p>Simply select the volume that you want to cache, and then caching is automatic and dynamic.</p>

6. Click **Save**.

Result

The volume settings are changed based on your selections.

Add volumes to workload

You can add unassigned volumes to an existing or new workload.

About this task

Volumes are not associated with a workload if they have been created using the command line interface (CLI) or if they have been migrated (imported/exported) from a different storage array.

Steps

1. From the Manage page, select the storage array that contains the volumes you want to add.
2. Select **Provisioning > Manage Volumes**.
3. Select the **Applications & Workloads** tab.

The Applications & Workloads view appears.

4. Select **Add to Workload**.

The Select Workload dialog box appears.

5. Do one of the following actions:

- **Add volumes to an existing workload** — Select this option to add volumes to an existing workload. Use the drop-down list to select a workload. The workload's associated application type is assigned to the volumes you add to this workload.
- **Add volumes to a new workload** — Select this option to define a new workload for an application type and add volumes to the new workload.

6. Select **Next** to continue with the add to workload sequence.

The Select Volumes dialog box appears.

7. Select the volumes you want to add to the workload.

8. Review the volumes that you want to add to the selected workload.

9. When you are satisfied with your workload configuration, click **Finish**.

Change workload settings

You can change the name for a workload and view its associated application type.

Steps

1. From the Manage page, select the storage array that contains the workload you want to change.
2. Select **Provisioning > Manage Volumes**.
3. Select the **Applications & Workloads** tab.

The Applications & Workloads view appears.

4. Select the workload that you want to change, and then select **View/Edit Settings**.

The Applications & Workloads Settings dialog box appears.

5. (Optional) Change the user-supplied name of the workload.
6. Click **Save**.

Initialize volumes

A volume is automatically initialized when it is first created. However, the Recovery Guru might advise that you manually initialize a volume to recover from certain failure conditions.

Use this option only under the guidance of technical support. You can select one or more volumes to initialize.

Before you begin

- All I/O operations have been stopped.
- Any devices or file systems on the volumes you want to initialize must be unmounted.
- The volume is in Optimal status and no modification operations are in progress on the volume.*Attention:
*You cannot cancel the operation after it starts. All volume data is erased. Do not try this operation unless the Recovery Guru advises you to do so. Contact technical support before you begin this procedure.

About this task

When you initialize a volume, the volume keeps its WWN, host assignments, allocated capacity, and reserved capacity settings. It also keeps the same Data Assurance (DA) settings and security settings.

The following types of volumes cannot be initialized:

- Base volume of a snapshot volume
- Primary volume in a mirror relationship
- Secondary volume in a mirror relationship
- Source volume in a volume copy
- Target volume in a volume copy
- Volume that already has an initialization in progress

This procedure applies only to standard volumes created from pools or volume groups.

Steps

1. From the Manage page, select the storage array that contains the volumes you want to initialize.
2. Select **Provisioning > Manage Volumes**.
3. Select any volume, and then select **More > Initialize volumes**.

The Initialize Volumes dialog box appears. All volumes on the storage array appear in this dialog box.

4. Select one or more volumes that you want to initialize and confirm that you want to perform the operation.

Results

The system performs the following actions:

- Erases all data from the volumes that were initialized.
- Clears the block indices, which causes unwritten blocks to be read as if they are zero-filled (the volume appears to be completely empty).

This operation can be lengthy and could affect system performance.

Redistribute volumes

You redistribute volumes to move volumes back to their preferred controller owners. Typically, multipath drivers move volumes from their preferred controller owner when a problem occurs along the data path between the host and storage array.

Before you begin

- The volumes you want to redistribute are not in use, or I/O errors will occur.
- A multipath driver is installed on all hosts using the volumes you want to redistribute, or I/O errors will occur. If you want to redistribute volumes without a multipath driver on the hosts, all I/O activity to the volumes while the redistribution operation is in progress must be stopped to prevent application errors.

About this task

Most host multipath drivers attempt to access each volume on a path to its preferred controller owner. However, if this preferred path becomes unavailable, the multipath driver on the host fails over to an alternate path. This failover might cause the volume ownership to change to the alternate controller. After you have

resolved the condition that caused the failover, some hosts might automatically move the volume ownership back to the preferred controller owner, but in some cases, you might need to manually redistribute the volumes.

Steps

1. From the Manage page, select the storage array that contains the volumes you want to redistribute.
2. Select **Provisioning > Manage Volumes**.
3. Select **More > Redistribute volumes**.

The Redistribute Volumes dialog box appears. All volumes on the storage array whose preferred controller owner does not match its current owner appear in this dialog box.

4. Select one or more volumes that you want to redistribute, and confirm that you want to perform the operation.

Result

The system moves the selected volumes to their preferred controller owners or you might see a Redistribute Volumes Unnecessary dialog box.

Change controller ownership of a volume

You can change the preferred controller ownership of a volume, so that I/O for host applications is directed through the new path.

Before you begin

If you do not use a multipath driver, any host applications that are currently using the volume must be shut down. This action prevents application errors when the I/O path changes.

About this task

You can change controller ownership for one or more volumes in a pool or volume group.

Steps

1. From the Manage page, select the storage array that contains the volumes for which you want to change the controller ownership.
2. Select **Provisioning > Manage Volumes**.
3. Select any volume, and then select **More > Change ownership**.

The Change Volume Ownership dialog box appears. All volumes on the storage array appear in this dialog box.

4. Use the **Preferred Owner** drop-down list to change the preferred controller for each volume that you want to change, and confirm that you want to perform the operation.

Results

- The system changes the controller ownership of the volume. I/O to the volume is now directed through this I/O path.
- The volume might not use the new I/O path until the multipath driver reconfigures to recognize the new path.

This action usually takes less than five minutes.

Change cache settings for a volume

You can change read cache and write cache settings to impact the overall I/O performance of a volume.

About this task

Keep these guidelines in mind when you change cache settings for a volume:

- After opening the Change Cache Settings dialog box, you might see an icon shown next to the selected cache properties. This icon indicates that the controller has temporarily suspended caching operations. This action might occur when a new battery is charging, when a controller has been removed, or if a mismatch in cache sizes has been detected by the controller. After the condition has cleared, the cache properties selected in the dialog box become active. If the selected cache properties do not become active, contact technical support.
- You can change the cache settings for a single volume or for multiple volumes on a storage array. You can change the cache settings for all volumes at the same time.

Steps

1. From the Manage page, select the storage array that contains the volumes for which you want to change cache settings.
2. Select **Provisioning > Manage Volumes**.
3. Select any volume, and then select **More > Change cache settings**.

The Change Cache Settings dialog box appears. All volumes on the storage array appear in this dialog box.

4. Select the **Basic** tab to change the settings for read caching and write caching.

Field Details

Cache setting	Description
Read Caching	The read cache is a buffer that stores data that has been read from the drives. The data for a read operation might already be in the cache from a previous operation, which eliminates the need to access the drives. The data stays in the read cache until it is flushed.
Write Caching	The write cache is a buffer that stores data from the host that has not yet been written to the drives. The data stays in the write cache until it is written to the drives. Write caching can increase I/O performance. Cache is automatically flushed after the Write caching is disabled for a volume.

5. Select the **Advanced** tab to change the advanced settings for thick volumes. The advanced cache settings are available only for thick volumes.

Field Details

Setting	Description
Dynamic Read Cache Prefetch	Dynamic Cache Read Prefetch allows the controller to copy additional sequential data blocks into the cache while it is reading data blocks from a drive to the cache. This caching increases the chance that future requests for data can be filled from the cache. Dynamic cache read prefetch is important for multimedia applications that use sequential I/O. The rate and amount of data that is prefetched into cache is self-adjusting based on the rate and request size of the host reads. Random access does not cause data to be prefetched into cache. This feature does not apply when read caching is disabled.
Write Caching without Batteries	<p>The Write Caching without Batteries setting enables write caching to continue even when the batteries are missing, failed, discharged completely, or not fully charged. Choosing write caching without batteries is not typically recommended, because data might be lost if power is lost. Typically, write caching is turned off temporarily by the controller until the batteries are charged or a failed battery is replaced.</p> <p>CAUTION: Possible loss of data — If you select this option and do not have a universal power supply for protection, you could lose data. In addition, you could lose data if you do not have controller batteries and you enable the Write caching without batteries option.</p>
Write Caching with Mirroring	Write Caching with Mirroring occurs when the data written to the cache memory of one controller is also written to the cache memory of the other controller. Therefore, if one controller fails, the other can complete all outstanding write operations. Write cache mirroring is available only if write caching is enabled and two controllers are present. Write caching with mirroring is the default setting at volume creation.

6. Click **Save** to change the cache settings.

Change media scan settings for a volume

A media scan is a background operation that scans all data and redundancy information in the volume. Use this option to enable or disable the media scan settings for one or more volumes, or to change the scan duration.

Before you begin

Understand the following:

- Media scans run continuously at a constant rate based on the capacity to be scanned and the scan duration. Background scans may be temporarily suspended by a higher priority background task (e.g. reconstruction), but will resume at the same constant rate.
- A volume is scanned only when the media scan option is enabled for the storage array and for that volume. If redundancy check is also enabled for that volume, redundancy information in the volume will be checked for consistency with data, provided that the volume has redundancy. Media scan with redundancy check is enabled by default for each volume when it is created.
- If an unrecoverable medium error is encountered during the scan, data will be repaired using redundancy

information, if available.

For example, redundancy information is available in optimal RAID 5 volumes, or in RAID 6 volumes that are optimal or only have one drive failed. If the unrecoverable error cannot be repaired using redundancy information, the data block will be added to the unreadable sector log. Both correctable and uncorrectable medium errors are reported to the event log.

- If the redundancy check finds an inconsistency between data and the redundancy information, it is reported to the event log.

About this task

Media scans detect and repair media errors on disk blocks that are infrequently read by applications. This can prevent data loss in the event of a drive failure, as data for failed drives is reconstructed using redundancy information and data from other drives in the volume group or pool.

You can perform the following actions:

- Enable or disable background media scans for the entire storage array
- Change the scan duration for the entire storage array
- Enable or disable media scan for one or more volumes
- Enable or disable the redundancy check for one or more volumes

Steps

1. From the Manage page, select the storage array that contains the volumes for which you want to change media scan settings.
2. Select **Provisioning > Manage Volumes**.
3. Select any volume, and then select **More > Change media scan settings**.

The Change Drive Media Scan Settings dialog box appears. All volumes on the storage array appear in this dialog box.

4. To enable the media scan, select the **Scan media over the course of...** check box. Disabling the media scan check box suspends all media scan settings.
5. Specify the number of days over which you want the media scan to run.
6. Select the **Media Scan** check box for each volume you want to perform a media scan on. The system enables the Redundancy Check option for each volume on which you choose to run a media scan. If there are individual volumes for which you do not want to perform a redundancy check, deselect the **Redundancy Check** check box.
7. Click **Save**.

Result

The system applies changes to background media scans based on your selection.

Delete volume

You can delete one or more volumes to increase the free capacity of a pool or volume group.

Before you begin

On the volumes that you plan to delete, make sure that:

- All data is backed up.
- All Input/Output (I/O) is stopped.
- Any devices and file systems are unmounted.

About this task

Typically, you delete volumes if the volumes were created with the wrong parameters or capacity, or it no longer meets storage configuration needs. Deleting a volume increases the free capacity in the pool or volume group.



Deleting a volume causes loss of all data on those volumes.

Be aware that you **cannot** delete a volume that has one of these conditions:

- The volume is initializing.
- The volume is reconstructing.
- The volume is part of a volume group that contains a drive that is undergoing a copyback operation.
- The volume is undergoing a modification operation, such as a change of segment size, unless the volume is now in Failed status.
- The volume is holding any type of persistent reservation.
- The volume is a source volume or a target volume in a Copy Volume that has a status of Pending, In Progress, or Failed.



When a volume exceeds a given size (currently 128 TB), the delete operation is performed in the background and the freed space might not be immediately available.

Steps

1. From the **Manage** page, select the storage array that contains the volumes you want to delete.
2. Select **Provisioning > Manage Volumes**.
3. Click **Delete**.

The Delete Volumes dialog box appears.

4. Select one or more volumes that you want to delete, and then confirm that you want to perform the operation.
5. Click **Delete**.

Configure hosts

Host creation overview

To manage storage with the Storage Plugin for vCenter, you must discover or define each host in the network. A host is a server that sends I/O to a volume on a storage array.

Automatic vs manual host creation

Creating a host is one of the steps required to let the storage array know which hosts are attached to it and to allow I/O access to the volumes. A host can be created automatically or manually.

- **Automatic** — Automatic host creation for SCSI-based (not NVMe-oF) hosts is initiated by the Host Context Agent (HCA). The HCA is a utility that you can install on each host attached to the storage array. Each host that has the HCA installed pushes its configuration information to the storage array controllers through the I/O path. Based on the host information, the controllers automatically create the host and the associated host ports and set the host type. If needed, you can make any additional changes to the host configuration. After the HCA performs its automatic detection, the host is automatically configured with the following attributes:

- The host name derived from the system name of the host.
- The host identifier ports that are associated with the host.
- The Host Operating System Type of the host.



Hosts are created as stand-alone hosts; the HCA does not automatically create or add to host clusters.

- **Manual** — During manual host creation, you associate host port identifiers by selecting them from a list or manually entering them. After you create a host, you can assign volumes to it or add it to a host cluster if you plan to share access to volumes.

How volumes are assigned

For a host to send I/O to a volume, you must assign the volume to it. You can select either a host or host cluster when you create a volume or you can assign a volume to a host or host cluster later. A host cluster is a group of hosts. You create a host cluster to make it easy to assign the same volumes to multiple hosts.

Assigning volumes to hosts is flexible, allowing you to meet your particular storage needs.

- **Stand-alone host, not part of a host cluster** — You can assign a volume to an individual host. The volume can be accessed only by the one host.
- **Host cluster** — You can assign a volume to a host cluster. The volume can be accessed by all the hosts in the host cluster.
- **Host within a host cluster** — You can assign a volume to an individual host that is part of a host cluster. Even though the host is part of a host cluster, the volume can be accessed only by the individual host and not by any other hosts in the host cluster.

When volumes are created, logical unit numbers (LUNs) are assigned automatically. The LUN serves as the address between the host and the controller during I/O operations. You can change LUNs after the volume is created.

Create host access

To manage storage with the Storage Plugin for vCenter, you must discover or define each host in the network.

About this task

By creating a host, you define the host parameters to provide connection to the storage array and I/O access to the volumes.

You can allow the Host Context Agent (HCA) to automatically detect the hosts, and then verify that the information is correct by selecting **View/Edit Settings** from the Configure Hosts page. However, the HCA is not available on all supported operating systems and you must create the host manually.

When you create a host, keep these guidelines in mind:

- You must define the host identifier ports that are associated with the host.
- Make sure that you provide the same name as the host's assigned system name.
- This operation does not succeed if the name you choose is already in use.
- The length of the name cannot exceed 30 characters.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning** › **Configure Hosts**.

The Configure Hosts page opens.

3. Click **Create** › **Host**.

The Create Host dialog box appears.

4. Select the settings for the host as appropriate.

Field Details

Setting	Description
Name	Type a name for the new host.
Host operating system type	Select the operating system that is running on the new host from the drop-down list.
Host interface type	(Optional) If you have more than one type of host interface supported on your storage array, select the host interface type that you want to use.
Host ports	<p>Do one of the following:</p> <ul style="list-style-type: none">• Select I/O Interface — Generally, the host ports should have logged in and be available from the drop-down list. You can select the host port identifiers from the list.• Manual add — If a host port identifier is not displayed in the list, it means that the host port has not logged in. An HBA utility or the iSCSI initiator utility may be used to find the host port identifiers and associate them with the host. You can manually enter the host port identifiers or copy/paste them from the utility (one at a time) into the Host ports field. You must select one host port identifier at a time to associate it with the host, but you can continue to select as many identifiers that are associated with the host. Each identifier is displayed in the Host ports field. If necessary, you also can remove an identifier by selecting the X next to it.
Set CHAP initiator secret	<p>(Optional) If you selected or manually entered a host port with an iSCSI IQN, and if you want to require a host that tries to access the storage array to authenticate using Challenge Handshake Authentication Protocol (CHAP), select the “Set CHAP initiator secret” checkbox. For each iSCSI host port you selected or manually entered, do the following:</p> <ul style="list-style-type: none">• Enter the same CHAP secret that was set on each iSCSI host initiator for CHAP authentication. If you are using mutual CHAP authentication (two-way authentication that enables a host to validate itself to the storage array and for a storage array to validate itself to the host), you also must set the CHAP secret for the storage array at initial setup or by changing settings.• Leave the field blank if you do not require host authentication. Currently, the only iSCSI authentication method used is CHAP.

5. Click **Create**.

6. If you need to update the host information, select the host from the table and click **View/Edit Settings**.

Result

After the host is successfully created, the system creates a default name for each host port configured for the host (user label). The default alias is <Hostname_Port Number>. For example, the default alias for the first

port created for host IPT is `IPT_1`.

After you finish

You must assign a volume to a host so it can be used for I/O operations. Go to [Assign volumes to hosts](#).

Create host cluster

When two or more hosts require I/O access to the same volumes, you can create a host cluster.

About this task

Keep these guidelines in mind when you create a host cluster:

- This operation does not start unless there are two or more hosts available to create the cluster.
- Hosts in host clusters can have different operating systems (heterogeneous).
- NVMe hosts in host clusters cannot be mixed with non- NVMe hosts.
- To create a Data Assurance (DA)-enabled volume, the host connection you are planning to use must support DA.

If any of the host connections on the controllers in your storage array do not support DA, the associated hosts cannot access data on DA-enabled volumes.

- This operation does not succeed if the name you choose is already in use.
- The length of the name cannot exceed 30 characters.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning** › **Configure Hosts**.

The Configure Hosts page opens.

3. Select **Create** › **Host cluster**.

The Create Host Cluster dialog box appears.

4. Select the settings for the host cluster as appropriate.

Setting	Description
Name	Type the name for the new host cluster.
Select hosts to share volume access	Select two or more hosts from the drop-down list. Only those hosts that are not already part of a host cluster appear in the list.

5. Click **Create**.

If the selected hosts are attached to interface types that have different Data Assurance (DA) capabilities, a dialog appears with the message that DA will be unavailable on the host cluster. This unavailability prevents DA-enabled volumes from being added to the host cluster. Select **Yes** to continue or **No** to cancel.

DA increases data integrity across the entire storage system. DA enables the storage array to check for

errors that might occur when data is moved between the hosts and the drives. Using DA for the new volume ensures that any errors are detected.

Result

The new host cluster appears in the table with the assigned hosts in the rows beneath.

After you finish

You must assign a volume to a host cluster so it can be used for I/O operations. Go to [Assign volumes to hosts](#).

Assign volumes to hosts

You must assign a volume to a host or a host cluster so it can be used for I/O operations.

Before you begin

Keep these guidelines in mind when you assign volumes to hosts:

- You can assign a volume to only one host or host cluster at a time.
- Assigned volumes are shared between controllers in the storage array.
- The same logical unit number (LUN) cannot be used twice by a host or a host cluster to access a volume. You must use a unique LUN.
- For new volume groups, if you wait until all volumes are created and initialized before you assign them to a host, the volume initialization time is reduced. Keep in mind that once a volume associated with the volume group is mapped, all volumes will revert to the slower initialization.

About this task

A volume assignment grants a host or host cluster access to that volume in a storage array.

All unassigned volumes are displayed during this task, but functions for hosts with or without Data Assurance (DA) apply as follows:

- For a DA-capable host, you can select volumes that are either DA-enabled or not DA-enabled.
- For a host that is not DA-capable, if you select a volume that is DA-enabled, a warning states that the system must automatically turn off DA on the volume before assigning the volume to the host.

Assigning a volume fails under these conditions:

- All volumes are assigned.
- The volume is already assigned to another host or host cluster. The ability to assign a volume is unavailable under these conditions:
- No valid hosts or host clusters exist.
- No host port identifiers have been defined for the host.
- All volume assignments have been defined.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning** > **Configure Hosts**.

The Configure Hosts page opens.

3. Select the host or host cluster to which you want to assign volumes, and then click **Assign Volumes**.

A dialog box appears that lists all the volumes that can be assigned. You can sort any of the columns or type something in the Filter box to make it easier to find particular volumes.

4. Select the check box next to each volume that you want to assign or select the check box in the table header to select all volumes.
5. Click **Assign** to complete the operation.

Results

After successfully assigning a volume or volumes to a host or a host cluster, the system performs the following actions:

- The assigned volume receives the next available LUN number. The host uses the LUN number to access the volume.
- The user-supplied volume name appears in volume listings associated to the host. If applicable, the factory-configured access volume also appears in volume listings associated to the host.

Unassign volumes

If you no longer need I/O access to a volume, you can unassign it from the host or host cluster.

About this task

Keep these guidelines in mind when you unassign a volume:

- If you are removing the last assigned volume from a host cluster, and the host cluster also has hosts with specific assigned volumes, make sure that you remove or move those assignments before removing the last assignment for the host cluster.
- If a host cluster, a host, or a host port is assigned to a volume that is registered to the operating system, you must clear this registration before you can remove these nodes.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning > Configure Hosts**.

The Configure Hosts page opens.

3. Select the host or host cluster that you want to edit, and then click **Unassign Volumes**.

A dialog box appears that shows all the volumes that are currently assigned.

4. Select the check box next to each volume that you want to unassign or select the check box in the table header to select all volumes.
5. Click **Unassign**.

Results

- The volumes that were unassigned are available for a new assignment.
- Until the changes are configured on the host, the volume is still recognized by the host operating system.

Change the settings for a host

You can change the name, host operating system type, and associated host clusters for a host or host cluster.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning** › **Configure Hosts**.

The Configure Hosts page opens.

3. Select the host that you want to edit, and then click **View/Edit Settings**.

A dialog box appears that shows the current host settings.


4. To change host properties, make sure the **Properties** tab is selected and then change the settings as appropriate.

Field Details

Setting	Description
Name	You can change the user-supplied name of the host. Specifying a name for the host is required.
Associated host cluster	You can choose one of the following options: <ul style="list-style-type: none">• None — The host remains a standalone host. If the host was associated to a host cluster, the system removes the host from the cluster.• <Host Cluster> — The system associates the host to the selected cluster.
Host operating system type	You can change the type of operating system running on the host you defined.

5. To change port settings, click the **Host Ports** tab and then change the settings as appropriate.

Field Details

Setting	Description
Host Port	<p>You can choose one of the following options:</p> <ul style="list-style-type: none">• Add — Use Add to associate a new host port identifier to the host. The length of the host port identifier name is determined by the host interface technology. Fibre Channel and Infiniband host port identifier names must have 16 characters. iSCSI host port identifier names have a maximum of 223 characters. The port must be unique. A port number that has already been configured is not allowed.• Delete — Use Delete to remove (unassociate) a host port identifier. The Delete option does not physically remove the host port. This option removes the association between the host port and the host. Unless you remove the host bus adapter or the iSCSI initiator, the host port is still recognized by the controller. <div><p>If you delete a host port identifier, it is no longer associated with this host. Also, the host loses access to any of its assigned volumes through this host port identifier.</p></div>
Label	<p>To change the port label name, click the Edit icon (pencil). The port label name must be unique. A label name that has already been configured is not allowed.</p>
CHAP Secret	<p>Appears only for iSCSI hosts. You can set or change the CHAP secret for the initiators (iSCSI hosts).</p> <p>The system uses the Challenge Handshake Authentication Protocol (CHAP) method, which validates the identity of targets and initiators during the initial link. Authentication is based on a shared security key called a CHAP secret.</p>

6. Click **Save**.

Delete host or host cluster

You can remove a host or host cluster so that volumes are no longer associated with that host.

About this task

Keep these guidelines in mind when you delete a host or a host cluster:

- Any specific volume assignments are deleted, and the associated volumes are available for a new assignment.
- If the host is part of a host cluster that has its own specific assignments, the host cluster is unaffected. However, if the host is part of a host cluster that does not have any other assignments, the host cluster and any other associated hosts or host port identifiers inherit any default assignments.
- Any host port identifiers that were associated with the host become undefined.

Steps

1. From the Manage page, select the storage array with the host connection.
2. Select **Provisioning** › **Configure Hosts**.

The Configure Hosts page opens.

3. Select the host or host cluster that you want to delete, and then click **Delete**.

The confirmation dialog box appears.

4. Confirm that you want to perform the operation, and then click **Delete**.

Results

If you deleted a host, the system performs the following actions:

- Deletes the host and, if applicable, removes it from the host cluster.
- Removes access to any assigned volumes.
- Returns the associated volumes to an unassigned state.
- Returns any host port identifiers associated with the host to an unassociated state. If you deleted a host cluster, the system performs the following actions:
 - Deletes the host cluster and its associated hosts (if any).
 - Removes access to any assigned volumes.
 - Returns the associated volumes to an unassigned state.
 - Returns any host port identifiers associated with the hosts to an unassociated state.

Configure pools and volume groups

Pools and volume group overview

To provision storage in the Storage Plugin for vCenter, you create either a pool or volume group that will contain the Hard Disk Drives (HDD) or Solid State Disk (SSD) drives that you want to use in your storage array.

Provisioning

Physical hardware is provisioned into logical components so that data can be organized and easily retrieved. There are two types of groupings supported:

- Pools
- Volume groups

The pools and volume groups are the top-level units of storage in a storage array: they divide the capacity of drives into manageable divisions. Within these logical divisions are the individual volumes or LUNs where data is stored.

When a storage system is deployed, the first step is to present the available drive capacity to the various hosts by:

- Creating pools or volume groups with sufficient capacity

- Adding the number of drives required to meet performance requirements to the pool or volume group
- Selecting the desired level of RAID protection (if using volume groups) to meet specific business requirements

You can have pools or volume groups on the same storage system, but a drive cannot be part of more than one pool or volume group. Volumes that are presented to hosts for I/O are then created, using the space on the pool or volume group.

Pools

Pools are designed to aggregate physical hard disk drives into a large storage space and to provide enhanced RAID protection for it. A pool creates many virtual RAID sets from the total number of drives assigned to the pool, and it spreads the data out evenly among all participating drives. If a drive is lost or added, the system dynamically re-balances the data across all the active drives.

Pools function as another RAID level, virtualizing the underlying RAID architecture to optimize performance and flexibility when performing tasks such as rebuilding, drive expansion, and handling drive loss. The system automatically sets the RAID level at 6 in an 8+2 configuration (eight data disks plus two parity disks).

Drive matching

You can choose from either HDD or SSDs for use in pools; however, as with volume groups, all drives in the pool must use the same technology. The controllers automatically select which drives to include, so you must make sure that you have a sufficient number of drives for the technology you choose.

Managing failed drives

Pools have a minimum capacity of 11 drives; however, one drive's worth of capacity is reserved for spare capacity in the event of a drive failure. This spare capacity is called "preservation capacity."

When pools are created, a certain amount of capacity is preserved for emergency use. This capacity is expressed in terms of a number of drives, but the actual implementation is spread across the entire pool of drives. The default amount of capacity that is preserved is based on the number of drives in the pool.

After the pool is created, you can change the preservation capacity value to more or less capacity, or even set it to no preservation capacity (0 drive's worth). The maximum amount of capacity that can be preserved (expressed as a number of drives) is 10, but the capacity that is available might be less, based on the total number of drives in the pool.

Volume groups

Volume groups define how capacity is allotted in the storage system to volumes. Disk drives are organized into RAID groups and volumes reside across the drives in a RAID group. Therefore, volume group configuration settings identify which drives are part of the group and what RAID level is used.

When you create a volume group, controllers automatically select the drives to include in the group. You must manually choose the RAID level for the group. The capacity of the volume group is the total of the number of drives that you select, multiplied by their capacity.

Drive matching

You must match the drives in the volume group for size and performance. If there are smaller and larger drives in the volume group, all drives are recognized as the smallest capacity size. If there are slower and faster drives in the volume group, all drives are recognized at the slowest speed. These factors affect the

performance and overall capacity of the storage system.

You cannot mix different drive technologies (HDD and SSD drives). RAID 3, 5, and 6 are limited to a maximum of 30 drives. RAID 1 and RAID 10 uses mirroring, so these volume groups must have an even number of disks.

Managing failed drives

Volume groups use hot spare drives as a standby in case a drive fails in RAID 1/10, RAID 3, RAID 5, or RAID 6 volumes contained in a volume group. A hot spare drive contains no data and adds another level of redundancy to your storage array.

If a drive fails in the storage array, the hot spare drive is automatically substituted for the failed drive without requiring a physical swap. If the hot spare drive is available when a drive fails, the controller uses redundancy data to reconstruct the data from the failed drive to the hot spare drive.

Decide whether to use pools or volume groups

Choose a pool

- If you need faster drive rebuilds and simplified storage administration, and/or have a highly random workload.
- If you want to distribute the data for each volume randomly across a set of drives that comprise the pool. You cannot set or change the RAID level of pools or the volumes in the pools. Pools use RAID level 6.

Choose a volume group

- If you need maximum system bandwidth, the ability to tune storage settings, and a highly sequential workload.
- If you want to distribute the data across the drives based on a RAID level. You can specify the RAID level when you create the volume group.
- If you want to write the data for each volume sequentially across the set of drives that comprise the volume group.



Because pools can co-exist with volume groups, a storage array can contain both pools and volume groups.

Automatic versus manual pool creation

Depending on your storage configuration, you can allow the system to create pools automatically or you can manually create them yourself. A pool is a set of logically grouped drives.

Before you create and manage pools, review the following sections for how pools are automatically created and when you might need to manually create them.

Automatic creation

When the system detects unassigned capacity in the storage array, it initiates automatic pool creation. It is initiated when the system detects unassigned capacity in a storage array. It automatically prompts you to create one or more pools, or add the unassigned capacity to an existing pool, or both.

Automatic pool creation occurs when one of these conditions is true:

- Pools do not exist in the storage array, and there are enough similar drives to create a new pool.

- New drives are added to a storage array that has at least one pool. Each drive in a pool must be of the same drive type (HDD or SSD) and have similar capacity. The system will prompt you to complete the following tasks:
- Create a single pool if there are a sufficient number of drives of those types.
- Create multiple pools if the unassigned capacity consists of different drive types.
- Add the drives to the existing pool if a pool is already defined in the storage array, and add new drives of the same drive type to the pool.
- Add the drives of the same drive type to the existing pool, and use the other drive types to create different pools if the new drives are of different drive types.

Manual creation

You might want to create a pool manually when automatic creation cannot determine the best configuration. This situation can occur for one of the following reasons:

- The new drives could potentially be added to more than one pool.
- One or more of the new pool candidates can use shelf loss protection or drawer loss protection.
- One or more of the current pool candidates cannot maintain their shelf loss protection or drawer loss protection status. You might also want to create a pool manually if you have multiple applications on your storage array and do not want them competing for the same drive resources. In this case, you might consider manually creating a smaller pool for one or more of the applications. You can assign just one or two volumes instead of assigning the workload to a large pool that has many volumes across which to distribute the data. Manually creating a separate pool that is dedicated to the workload of a specific application can allow storage array operations to perform more rapidly, with less contention.

Create pool automatically

You can create pools automatically when the system detects at least 11 unassigned drives or it detects one unassigned drive that is eligible for an existing pool. A pool is a set of logically grouped drives.

Before you begin

You can launch the Pool Auto-Configuration dialog box when one of these conditions are true:

- At least one unassigned drive has been detected that can be added to an existing pool with similar drive types.
- Eleven (11) or more unassigned drives have been detected that can be used to create a new pool (if they cannot be added to an existing pool due to dissimilar drive types).

About this task

You can use automatic pool creation to easily configure all unassigned drives in the storage array into one pool and to add drives into existing pools.

Keep in mind the following:

- When you add drives to a storage array, the system automatically detects the drives and prompts you to create a single pool or multiple pools based on the drive type and the current configuration.
- If pools were previously defined, the system automatically prompts you with the option of adding the compatible drives to an existing pool. When new drives are added to an existing pool, the system automatically redistributes the data across the new capacity, which now includes the new drives that you

added.

- When configuring an EF600 or EF300 storage array, make sure each controller has access to an equal number of drives in the first 12 slots and an equal number of drives in the last 12 slots. This configuration helps the controllers use both drive-side PCIe buses more effectively. For pool creation, you should use all drives in the storage array.

Steps

1. From the Manage page, select the storage array for the pool.
2. Select **Provisioning** › **Configure Pools and Volume Groups**.
3. Select **More** › **Launch pool auto-configuration**.

The results table lists new pools, existing pools with drives added, or both. A new pool is named with a sequential number by default.

Notice that the system does the following:

- Creates a single pool if there are a sufficient number of drives with the same drive type (HDD or SSD) and have similar capacity.
 - Creates multiple pools if the unassigned capacity consists of different drive types.
 - Adds the drives to an existing pool if a pool is already defined in the storage array, and you add new drives of the same drive type to the pool.
 - Adds the drives of the same drive type to the existing pool, and use the other drive types to create different pools if the new drives are of different drive types.
4. To change the name of a new pool, click the **Edit** icon (the pencil).
 5. To view additional characteristics of the pool, position the cursor over or touch the Details icon (the page).

Information about the drive type, security capability, data assurance (DA) capability, shelf loss protection, and drawer loss protection appears.

For EF600 and EF300 storage arrays, settings are also displayed for resource provisioning and volume block sizes.

6. Click **Accept**.

Create pool manually

You can create a pool manually if your setup does not meet the requirements for automatic pool configuration. A pool is a set of logically grouped drives.

Before you begin

- You must have a minimum of 11 drives with the same drive type (HDD or SSD).
- Shelf loss protection requires that the drives comprising the pool are located in at least six different drive shelves and there are no more than two drives in a single drive shelf.
- Drawer loss protection requires that the drives comprising the pool are located in at least five different drawers and the pool includes an equal number of drive shelves from each drawer.
- When configuring an EF600 or EF300 storage array, make sure each controller has access to an equal number of drives in the first 12 slots and an equal number of drives in the last 12 slots. This configuration helps the controllers use both drive-side PCIe buses more effectively. For pool creation, you should use all drives in the storage array.

About this task

During pool creation you determine its characteristics, such as drive type, security capability, data assurance (DA) capability, shelf loss protection, and drawer loss protection.

For EF600 and EF300 storage arrays, settings also include resource provisioning and volume block sizes.

Steps

1. From the Manage page, select the storage array for the pool.
2. Select **Provisioning › Configure Pools and Volume Groups**.
3. Click **Create › Pool**.


The Create Pool dialog box appears.

4. Type a name for the pool.
5. (Optional) If you have more than one type of drive in your storage array, select the drive type that you want to use.

The results table lists all the possible pools that you can create.

6. Select the pool candidate that you want to use based on the following characteristics, and then click **Create**.

Field Details

Characteristic	Use
Free Capacity	Shows the free capacity of the pool candidate in GiB. Select a pool candidate with the capacity for your application's storage needs. Preservation (spare) capacity is also distributed throughout the pool and is not part of the free capacity amount.
Total Drives	Shows the number of drives available in the pool candidate. The system automatically reserves as many drives as possible for preservation capacity (for every six drives in a pool, the system reserves one drive for preservation capacity). When a drive failure occurs, the preservation capacity is used to hold the reconstructed data.
Drive Block Size (EF300 and EF600 only)	Shows the block size (sector size) that the drives in the pool can write. Values may include: <ul style="list-style-type: none"> • 512 — 512-byte sector size. • 4K — 4,096-byte sector size.
Secure-Capable	Indicates whether this pool candidate is comprised entirely of secure-capable drives, which can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives. <ul style="list-style-type: none"> • You can protect your pool with Drive Security, but all drives must be secure-capable to use this feature. • If you want to create an FDE-only pool, look for Yes - FDE in the Secure-Capable column. If you want to create a FIPS-only pool, look for Yes - FIPS or Yes - FIPS (Mixed). "Mixed" indicates a mixture of 140-2 and 140-3 level drives. If you use a mixture of these levels, be aware that the pool will then operate at the lower level of security (140-2). • You can create a pool comprised of drives that may or may not be secure-capable or are a mix of security levels. If the drives in the pool include drives that are not secure-capable, you cannot make the pool secure.
Enable Security?	Provides the option for enabling the Drive Security feature with secure-capable drives. If the pool is secure-capable and you have created a security key, you can enable security by selecting the check box. <div>  <p>The only way to remove Drive Security after it is enabled is to delete the pool and erase the drives.</p> </div>

Characteristic	Use
DA Capable	<p>Indicates if Data Assurance (DA) is available for this pool candidate. DA checks for and corrects errors that might occur as data is transferred through the controllers down to the drives.</p> <p>If you want to use DA, select a pool that is DA capable. This option is available only when the DA feature has been enabled.</p> <p>A pool can contain drives that are DA-capable or not DA-capable, but all drives must be DA capable for you to use this feature.</p>
Resource Provisioning Capable (EF300 and EF600 only)	<p>Shows if Resource Provisioning is available for this pool candidate. Resource Provisioning is a feature available in the EF300 and EF600 storage arrays, which allows volumes to be put in use immediately with no background initialization process.</p>
Shelf Loss Protection	<p>Shows if shelf loss protection is available.</p> <p>Shelf loss protection guarantees accessibility to the data on the volumes in a pool if a total loss of communication occurs with a single drive shelf.</p>
Drawer Loss Protection	<p>Shows if drawer loss protection is available, which is provided only if you are using a drive shelf that contains drawers.</p> <p>Drawer loss protection guarantees accessibility to the data on the volumes in a pool if a total loss of communication occurs with a single drawer in a drive shelf.</p>
Volume Block Sizes Supported (EF300 and EF600 only)	<p>Shows the block sizes that can be created for the volumes in the pool:</p> <ul style="list-style-type: none"> • 512n — 512 bytes native. • 512e — 512 bytes emulated. • 4K — 4,096 bytes.

Create a volume group

You can create a volume group for one or more volumes that are accessible to the host. A volume group is a container for volumes with shared characteristics such as RAID level and capacity.

Before you begin

Review the following guidelines:

- You need at least one unassigned drive.
- Limits exist as to how much drive capacity you can have in a single volume group. These limits vary according to your host type.
- To enable shelf/drawer loss protection, you must create a volume group that uses drives located in at least three shelves or drawers, unless you are using RAID 1, where two shelves/drawers is the minimum.
- When configuring an EF600 or EF300 storage array, make sure each controller has access to an equal number of drives in the first 12 slots and an equal number of drives in the last 12 slots. This configuration

helps the controllers use both drive-side PCIe buses more effectively. The system currently allows for drive selection under the Advanced feature when creating a volume group.

Review how your choice of RAID level affects the resulting capacity of the volume group.

- If you select RAID 1, you must add two drives at a time to make sure that a mirrored pair is selected. Mirroring and striping (known as RAID 10 or RAID 1+0) is achieved when four or more drives are selected.
- If you select RAID 5, you must add a minimum of three drives to create the volume group.
- If you select RAID 6, you must add a minimum of five drives to create the volume group.

About this task

During volume group creation you determine the group characteristics, such as the number of drives, security capability, data assurance (DA) capability, shelf loss protection, and drawer loss protection.

For EF600 and EF300 storage arrays, settings also include resource provisioning, drive block sizes, and volume block sizes.



With larger capacity drives and the ability to distribute volumes across controllers, creating more than one volume per volume group is a good way to make use of your storage capacity and to protect your data.

Steps

1. From the Manage page, select the storage array for the volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Click **Create > Volume group**.

The Create Volume Group dialog box appears.

4. Type a name for the volume group.
5. Select the RAID level that best meets your requirements for data storage and protection. The volume group candidate table appears and displays only the candidates that support the selected RAID level.
6. (Optional) If you have more than one type of drive in your storage array, select the drive type that you want to use.

The volume group candidate table appears and displays only the candidates that support the selected drive type and RAID level.

7. (Optional) You can select either the automatic method or manual method to define which drives to use in the volume group. The Automatic method is the default selection.



Do not use the Manual method unless you are an expert who understands drive redundancy and optimal drive configurations.

To select drives manually, click the **Manually select drives (advanced)** link. When clicked, it changes to **Automatically select drives (advanced)**.

The Manual method lets you select which specific drives comprise the volume group. You can select specific unassigned drives to obtain the capacity that you require. If the storage array contains drives with different media types or different interface types, you can choose only the unconfigured capacity for a single drive type to create the new volume group.

8. Based on the displayed drive characteristics, select the drives you want to use in the volume group, and then click **Create**.

The drive characteristics displayed depend on whether you selected the automatic method or manual method. For more information, see the SANtricity System Manager documentation, [Create a volume group](#).

Add capacity to a pool or volume group

You can add drives to expand the free capacity in an existing pool or volume group.

Before you begin

- Drives must be in an Optimal status.
- Drives must have the same drive type (HDD or SSD).
- The pool or volume group must be in an Optimal status.
- If the pool or volume group contains all secure-capable drives, add only drives that are secure-capable to continue to use the encryption abilities of the secure-capable drives.

Secure-capable drives can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives.

About this task

In this task, you can add free capacity to be included in the pool or volume group. You can use this free capacity to create additional volumes. The data in the volumes remains accessible during this operation.

For pools, you can add a maximum of 60 drives at a time. For volume groups, you can add a maximum of two drives at a time. If you need to add more than the maximum number of drives, repeat the procedure. (A pool cannot contain more drives than the maximum limit for a storage array.)



With the addition of drives, your preservation capacity may need to be increased. You should consider increasing your reserved capacity after an expansion operation.



Avoid using drives that are Data Assurance (DA) capable to add capacity to a pool or volume group that is not DA capable. The pool or volume group cannot take advantage of the capabilities of the DA-capable drive. Consider using drives that are not DA capable in this situation.

Steps

1. From the Manage page, select the storage array with the pool or volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select the pool or volume group to which you want to add drives, and then click **Add Capacity**.

The Add Capacity dialog box appears. Only the unassigned drives that are compatible with the pool or volume group appear.

4. Under **Select drives to add capacity...**, select one or more drives that you want to add to the existing pool or volume group.

The controller firmware arranges the unassigned drives with the best options listed at the top. The total free capacity that is added to the pool or volume group appears below the list in **Total capacity selected**.

Field Details

Field	Description
Shelf	Indicates the shelf location of the drive.
Bay	Indicates the bay location of the drive
Capacity (GiB)	<p>Indicates the drive capacity.</p> <ul style="list-style-type: none">• Whenever possible, select drives that have a capacity equal to the capacities of the current drives in the pool or volume group.• If you must add unassigned drives with a smaller capacity, be aware that the usable capacity of each drive currently in the pool or volume group is reduced. Therefore, the drive capacity is the same across the pool or volume group.• If you must add unassigned drives with a larger capacity, be aware that the usable capacity of the unassigned drives that you add is reduced so that they match the current capacities of the drives in the pool or volume group.
Secure-Capable	<p>Indicates whether the drive is secure-capable.</p> <ul style="list-style-type: none">• You can protect your pool or volume group with the Drive Security feature, but all drives must be secure-capable to use this feature.• It is possible to create a pool or volume group with a mix of secure-capable and non-secure-capable drives, but the Drive Security feature cannot be enabled.• A pool or volume group with all secure-capable drives cannot accept a non-secure-capable drive for sparing or expansion, even if the encryption capability is not in use.• Secure-capable drives can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives. A FIPS drive can be level 140-2 or 140-3, with level 140-3 as the higher level of security. If you select a mixture of 140-2 and 140-3 level drives, the pool or volume group will then operate at the lower level of security (140-2).

Field	Description
DA Capable	<p>Indicates whether the drive is Data Assurance (DA) capable.</p> <ul style="list-style-type: none"> Using drives that are not Data Assurance (DA) capable to add capacity to a DA-capable pool or volume group is not recommended. The pool or volume group no longer has DA capabilities, and you no longer have the option to enable DA on newly created volumes within the pool or volume group. Using drives that are Data Assurance (DA) capable to add capacity to a pool or volume group that is non DA-capable is not recommended, because that pool or volume group cannot take advantage of the capabilities of the DA-capable drive (the drive attributes do not match). Consider using drives that are not DA-capable in this situation.
DULBE Capable	<p>Indicates whether the drive has the option for Deallocated or Unwritten Logical Block Error (DULBE). DULBE is an option on NVMe drives that allows the EF300 or EF600 storage array to support resource-provisioned volumes.</p>

5. Click **Add**.

If you are adding drives to a pool or volume group, a confirmation dialog box appears if you selected a drive that causes the pool or volume group to no longer have one or more of the following attributes:

- Shelf loss protection
- Drawer loss protection
- Full Disk Encryption capability
- Data Assurance capability
- DULBE capability

6. To continue, click **Yes**; otherwise click **Cancel**.

Result

After you add the unassigned drives to a pool or volume group, the data in each volume of the pool or volume group is redistributed to include the additional drives.

Create SSD Cache

To dynamically accelerate system performance, you can use the SSD Cache feature to cache the most frequently accessed data ("hot" data) onto lower latency Solid State Drives (SSDs). SSD Cache is used exclusively for host reads.

Before you begin

Your storage array must contain some SSD drives.



SSD Cache is not available on the EF600 or EF300 storage system.

About this task

When you create SSD Cache, you can use a single drive or multiple drives. Because the read cache is in the storage array, caching is shared across all applications using the storage array. You select the volumes that you want to cache, and then caching is automatic and dynamic.

Follow these guidelines when you create SSD Cache.

- You can enable security on the SSD Cache only when you are creating it, not later.
- Only one SSD Cache is supported per storage array.
- The maximum usable SSD Cache capacity on a storage array is dependent on the controller's primary cache capacity.
- SSD Cache is not supported on snapshot images.
- If you import or export volumes that are SSD Cache enabled or disabled, the cached data is not imported or exported.
- Any volume assigned to use a controller's SSD Cache is not eligible for an automatic load balance transfer.
- If the associated volumes are secure-enabled, create a secure-enabled SSD Cache.

Steps

1. From the Manage page, select the storage array for the cache.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Click **Create > SSD Cache**.

The Create SSD Cache dialog box appears.

4. Type a name for the SSD Cache.
5. Select the SSD Cache candidate that you want to use based on the following characteristics.

Field Details

Characteristic	Use
Capacity	Shows the available capacity in GiB. Select the capacity for your application's storage needs. The maximum capacity for SSD Cache depends on the controller's primary cache capacity. If you allocate more than the maximum amount to SSD Cache, then any extra capacity is unusable. SSD Cache capacity counts towards your overall allocated capacity.
Total drives	Shows the number of drives available for this SSD cache. Select the SSD candidate with the number of drives that you want
Secure-capable	Indicates whether the SSD Cache candidate is comprised entirely of secure-capable drives, which can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives. If you want to create a secure-enabled SSD Cache, look for "Yes - FDE" or "Yes - FIPS" in the Secure-capable column.
Enable security?	Provides the option for enabling the Drive Security feature with secure-capable drives. If you want to create a secure-enabled SSD Cache, select the Enable Security check box. NOTE: Once enabled, security cannot be disabled. You can enable security on the SSD Cache only when you are creating it, not later.
DA capable	Indicates if Data Assurance (DA) is available for this SSD Cache candidate. Data Assurance (DA) checks for and corrects errors that might occur as data is transferred through the controllers down to the drives. If you want to use DA, select an SSD Cache candidate that is DA capable. This option is available only when the DA feature has been enabled. SSD Cache can contain both DA-capable and non-DA-capable drives, but all drives must be DA-capable for you to use DA.

6. Associate the SSD Cache with the volumes for which you want to implement SSD read caching. To enable SSD Cache on compatible volumes immediately, select the **Enable SSD Cache on existing compatible volumes that are mapped to hosts** check box.

Volumes are compatible if they share the same Drive Security and DA capabilities.

7. Click **Create**.

Change configuration settings for a pool

You can edit the settings for a pool, including its name, capacity alerts settings, modification priorities, and preservation capacity.

About this task

This task describes how to change configuration settings for a pool.



You cannot change the RAID level of a pool using the plugin interface. The plugin automatically configures pools as RAID 6.

Steps

1. From the Manage page, select the storage array with the pool.
2. Select **Provisioning** › **Configure Pools and Volume Groups**.
3. Select the pool that you want to edit, and then click **View/Edit Settings**.

The Pool Settings dialog box appears.

4. Select the **Settings** tab, and then edit the pool settings as appropriate.

Field Details

Setting	Description
Name	You can change the user-supplied name of the pool. Specifying a name for a pool is required.
Capacity alerts	<p>You can send alert notifications when the free capacity in a pool reaches or exceeds a specified threshold. When the data stored in the pool exceeds the specified threshold, the plugin sends a message, allowing you time to add more storage space or to delete unnecessary objects.</p> <p>Alerts are shown in the Notifications area on the Dashboard and can be sent from the server to administrators by email and SNMP trap messages. You can define the following capacity alerts:</p> <ul style="list-style-type: none"> • Critical alert — This critical alert notifies you when the free capacity in the pool reaches or exceeds the specified threshold. Use the spinner controls to adjust the threshold percentage. Select the check box to disable this notification. • Early alert — This early alert notifies you when the free capacity in a pool is reaching a specified threshold. Use the spinner controls to adjust the threshold percentage. Select the check box to disable this notification.
Modification priorities	<p>You can specify the priority levels for modification operations in a pool relative to system performance. A higher priority for modification operations in a pool causes an operation to complete faster, but can slow the host I/O performance. A lower priority causes operations to take longer, but host I/O performance is less affected.</p> <p>You can choose from five priority levels: lowest, low, medium, high, and highest. The higher the priority level, the larger is the impact on host I/O and system performance.</p> <ul style="list-style-type: none"> • Critical reconstruction priority — This slider bar determines the priority of a data reconstruction operation when multiple drive failures result in a condition where some data has no redundancy and an additional drive failure might result in loss of data. • Degraded reconstruction priority — This slider bar determines the priority of the data reconstruction operation when a drive failure has occurred, but the data still has redundancy and an additional drive failure does not result in loss of data. • Background operation priority — This slider bar determines the priority of the pool background operations that occur while the pool is in an optimal state. These operations include Dynamic Volume Expansion (DVE), Instant Availability Format (IAF), and migrating data to a replaced or added drive.

Setting	Description
Preservation capacity ("Optimization capacity" for the EF600 or EF300)	<p>Preservation capacity — You can define the number of drives to determine the capacity that is reserved on the pool to support potential drive failures. When a drive failure occurs, the preservation capacity is used to hold the reconstructed data. Pools use preservation capacity during the data reconstruction process instead of hot spare drives, which are used in volume groups.</p> <p>Use the spinner controls to adjust the number of drives. Based on the number of drives, the preservation capacity in the pool appears next to the spinner box.</p> <p>Keep the following information in mind about preservation capacity.</p> <ul style="list-style-type: none"> • Because preservation capacity is subtracted from the total free capacity of a pool, the amount of capacity that you reserve affects how much free capacity is available to create volumes. If you specify 0 for the preservation capacity, all of the free capacity on the pool is used for volume creation. • If you decrease the preservation capacity, you increase the capacity that can be used for pool volumes. <p>Additional optimization capacity (EF600 and EF300 arrays only) — When a pool is created, a recommended optimization capacity is generated that provides a balance of available capacity versus performance and drive wear life. You can adjust this balance by moving the slider to the right for better performance and drive wear life at the expense of increased available capacity, or by moving it to the left for increased available capacity at the expense of better performance and drive wear life. SSD drives will have longer life and better maximum write performance when a portion of their capacity is unallocated. For drives associated with a pool, unallocated capacity is comprised of a pool's preservation capacity, the free capacity (capacity not used by volumes), and a portion of the usable capacity set aside as additional optimization capacity. The additional optimization capacity ensures a minimum level of optimization capacity by reducing the usable capacity, and as such, is not available for volume creation.</p>

5. Click **Save**.

Change configuration settings for a volume group

You can edit the settings for a volume group, including its name and RAID level.

Before you begin

If you are changing the RAID level to accommodate the performance needs of the applications that are accessing the volume group, be sure to meet the following prerequisites:

- The volume group must be in Optimal status.
- You must have enough capacity in the volume group to convert to the new RAID level.

Steps

1. From the Manage page, select the storage array with the volume group.
2. Select **Provisioning** › **Configure Pools and Volume Groups**.
3. Select the volume group that you want to edit, and then click **View/Edit Settings**.

The Volume Group Settings dialog box appears.

4. Select the **Settings** tab, and then edit the volume group settings as appropriate.

Field Details

Setting	Description
Name	You can change the user-supplied name of the volume group. Specifying a name for a volume group is required.
RAID level	<p>Select the new RAID level from the drop-down menu.</p> <ul style="list-style-type: none"> • RAID 0 striping — Offers high performance but does not provide any data redundancy. If a single drive fails in the volume group, all the associated volumes fail, and all data is lost. A striping RAID group combines two or more drives into one large, logical drive. • RAID 1 mirroring — Offers high performance and the best data availability and is suitable for storing sensitive data on a corporate or personal level. Protects your data by automatically mirroring the contents of one drive to the second drive in the mirrored pair. It provides protection in the event of a single drive failure. • RAID 10 striping/mirroring — Provides a combination of RAID 0 (striping) and RAID 1 (mirroring) and is achieved when four or more drives are selected. RAID 10 is suitable for high volume transaction applications, such as a database, that require high performance and fault tolerance. • RAID 5 — Optimal for multi-user environments (such as database or file system storage) where typical I/O size is small and there is a high proportion of read activity. • RAID 6 — Optimal for environments requiring redundancy protection beyond RAID 5, but not requiring high write performance. RAID 3 can be assigned only to volume groups using the command line interface (CLI). When you change the RAID level, you cannot cancel this operation after it begins. During the change, your data remains available.
Optimization capacity (EF600 arrays only)	When a volume group is created, a recommended optimization capacity is generated that provides a balance of available capacity versus performance and drive wear life. You can adjust this balance by moving the slider to the right for better performance and drive wear life at the expense of increased available capacity, or by moving it to the left for increased available capacity at the expense of better performance and drive wear life. SSD drives will have longer life and better maximum write performance when a portion of their capacity is unallocated. For drives associated with a volume group, unallocated capacity is comprised of a group's free capacity (capacity not used by volumes) and a portion of the usable capacity set aside as additional optimization capacity. The additional optimization capacity ensures a minimum level of optimization capacity by reducing the usable capacity, and as such, is not available for volume creation.

5. Click **Save**.

A confirmation dialog box appears if capacity is reduced, volume redundancy is lost, or shelf/ drawer loss

protection is lost as a result of the RAID level change. Select **Yes** to continue; otherwise click **No**.

Result

If you change the RAID level for a volume group, the plugin changes the RAID levels of every volume that comprises the volume group. Performance might be slightly affected during the operation.

Change SSD Cache settings

You can edit the name of the SSD Cache and view its status, maximum and current capacity, Drive Security and Data Assurance status, and its associated volumes and drives.



This feature is not available on the EF600 or EF300 storage system.

Steps

1. From the Manage page, select the storage array with the SSD Cache.
2. Select **Provisioning** > **Configure Pools and Volume Groups**.
3. Select the SSD Cache that you want to edit, and then click **View/Edit Settings**.

The SSD Cache Settings dialog box appears.

4. Review or edit the SSD Cache settings as appropriate.

Field Details

Setting	Description
Name	Displays the name of the SSD Cache, which you can change. A name for the SSD Cache is required.
Characteristics	Shows the status for the SSD Cache. Possible statuses include: <ul style="list-style-type: none">• Optimal• Unknown• Degraded• Failed (A failed state results in a critical MEL event.)• Suspended
Capacities	Shows the current capacity and maximum capacity allowed for the SSD Cache. The maximum capacity allowed for the SSD Cache depends on the controller's primary cache size: <ul style="list-style-type: none">• Up to 1 GiB• 1 GiB to 2 GiB• 2 GiB to 4 GiB• More than 4 GiB
Security and DA	Shows the Drive Security and Data Assurance status for the SSD Cache. <ul style="list-style-type: none">• Secure-capable --Indicates whether the SSD Cache is comprised entirely of secure-capable drives. A secure-capable drive is a self-encrypting drive that can protect its data from unauthorized access.• Secure-enabled — Indicates whether security is enabled on the SSD Cache.• DA capable — Indicates whether the SSD Cache is comprised entirely of DA-capable drives. A DA-capable drive can check for and correct errors that might occur as data is communicated between the host and storage array.
Associated objects	Shows the volumes and drives associated with the SSD Cache.

5. Click **Save**.

View SSD Cache statistics

You can view statistics for the SSD Cache, such as reads, writes, cache hits, cache allocation percentage, and cache utilization percentage.



This feature is not available on the EF600 or EF300 storage system.

About this task

The nominal statistics, which are a subset of the detailed statistics, are shown on the View SSD Cache Statistics dialog box. You can view detailed statistics for the SSD Cache only when you export all SSD statistics to a .csv file.

As you review and interpret the statistics, keep in mind that some interpretations are derived by looking at a combination of statistics.

Steps

1. From the Manage page, select the storage array with the SSD Cache.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select the SSD Cache for which you want to view statistics, and then click **More > View SSD Cache statistics**.

The View SSD Cache Statistics dialog box appears and displays the nominal statistics for the selected SSD cache.

Field Details

Setting	Description
Reads	Shows the total number of host reads from the SSD Cache-enabled volumes. The greater the ratio of Reads to Writes, the better is the operation of the cache.
Writes	The total number of host writes to the SSD Cache-enabled volumes. The greater the ratio of Reads to Writes, the better is the operation of the cache.
Cache hits	Shows the number of cache hits.
Cache hits %	Shows the percentage of cache hits. This number is derived from Cache Hits / (reads + writes). The cache hit percentage should be greater than 50 percent for effective SSD Cache operation.
Cache allocation %	Shows the percentage of SSD Cache storage that is allocated, expressed as a percentage of the SSD Cache storage that is available to this controller and is derived from allocated bytes / available bytes.
Cache utilization %	Shows the percentage of SSD Cache storage that contains data from enabled volumes, expressed as a percentage of SSD Cache storage that is allocated. This amount represents the utilization or density of the SSD Cache. Derived from allocated bytes / available bytes.
Export All	Exports all SSD Cache statistics to a CSV format. The exported file contains all available statistics for the SSD Cache (both nominal and detailed).

4. Click **Cancel** to close the dialog box.

Check volume redundancy

Under the guidance of technical support or as instructed by the Recovery Guru, you can check the redundancy on a volume in a pool or volume group to determine whether the data on that volume is consistent.

Redundancy data is used to quickly reconstruct information on a replacement drive if one of the drives in the pool or volume group fails.

Before you begin

- The status of the pool or volume group must be Optimal.
- The pool or volume group must have no volume modification operations in progress.
- You can check redundancy on any RAID level except on RAID 0, because RAID 0 has no data redundancy. (Pools are configured only as RAID 6.)



Check volume redundancy only when instructed to do so by the Recovery Guru and under the guidance of technical support.

About this task

You can perform this check only on one pool or volume group at a time. A volume redundancy check performs the following actions:

- Scans the data blocks in a RAID 3 volume, a RAID 5 volume, or a RAID 6 volume, and checks the redundancy information for each block. (RAID 3 can only be assigned to volume groups using the command line interface.)
- Compares the data blocks on RAID 1 mirrored drives.
- Returns redundancy errors if the controller firmware determines that the data is inconsistent.



Immediately running a redundancy check on the same pool or volume group might cause an error. To avoid this problem, wait one to two minutes before running another redundancy check on the same pool or volume group.

Steps

1. From the Manage page, select the storage array with the pool or volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select **Uncommon Tasks > Check volume redundancy**.

The Check Redundancy dialog box appears.

4. Select the volumes you want to check, and then type check to confirm you want to perform this operation.
5. Click **Check**.

The check volume redundancy operation starts. The volumes in the pool or volume group are sequentially scanned, starting from the top of the table in the dialog box. These actions occur as each volume is scanned:

- The volume is selected in the volume table.
- The status of the redundancy check is shown in the Status column.

- The check stops on any media or parity error encountered, and then reports the error. The following table provide more information about the status of the redundancy check:

Field Details

Status	Description
Pending	This is the first volume to be scanned, and you have not clicked Start to start the redundancy check. -or- The redundancy check operation is being performed on other volumes in the pool or volume group.
Checking	The volume is undergoing the redundancy check.
Passed	The volume passed the redundancy check. No inconsistencies were detected in the redundancy information.
Failed	The volume failed the redundancy check. Inconsistencies were detected in the redundancy information.
Media error	The drive media is defective and is unreadable. Follow the instructions displayed in the Recovery Guru.
Parity error	The parity is not what it should be for a given portion of the data. A parity error is potentially serious and could cause a permanent loss of data.

6. Click **Done** after the last volume in the pool or volume group has been checked.

Delete pool or volume group

You can delete a pool or volume group to create more unassigned capacity, which you can reconfigure to meet your application storage needs.

Before you begin

- You must have backed up the data on all of the volumes in the pool or volume group.
- You must have stopped all input/output (I/O).
- You must unmount any file systems on the volumes.
- You must have deleted any mirror relationships in the pool or volume group.
- You must have stopped any volume copy operation in progress for the pool or volume group.
- The pool or volume group must not be participating in an asynchronous mirroring operation.
- The drives in the volume group must not have a persistent reservation.

Steps

1. From the Manage page, select the storage array with the pool or volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select one pool or volume group from the list.

You can select only one pool or volume group at a time. Scroll down the list to see additional pools or volume groups.

4. Select **Uncommon Tasks** > **Delete** and confirm.

Results

The system performs the following actions:

- Deletes all of the data in the pool or volume group.
- Deletes all the drives associated with the pool or volume group.
- Unassigns the associated drives, which allows you to reuse them in new or existing pools or volume groups.

Consolidate free capacity for a volume group

Use the Consolidate Free Capacity option to consolidate existing free extents on a selected volume group. By performing this action, you can create additional volumes from the maximum amount of free capacity in a volume group.

Before you begin

- The volume group must contain at least one free capacity area.
- All of the volumes in the volume group must be online and in Optimal status.
- Volume modification operations must not be in progress, such as changing the segment size of a volume.

About this task

You cannot cancel the operation after it begins. Your data remains accessible during the consolidation operation.

You can launch the Consolidate Free Capacity dialog box using any of the following methods:

- When at least one free capacity area is detected for a volume group, the Consolidate free capacity recommendation appears on the Home page in the Notification area. Click the **Consolidate free capacity** link to launch the dialog box.
- You can also launch the Consolidate Free Capacity dialog box from the Pools & Volume Groups page as described in the following task.

More about free capacity areas

A free capacity area is the free capacity that can result from deleting a volume or from not using all available free capacity during volume creation. When you create a volume in a volume group that has one or more free capacity areas, the volume's capacity is limited to the largest free capacity area in that volume group. For example, if a volume group has a total of 15 GiB free capacity, and the largest free capacity area is 10 GiB, the largest volume you can create is 10 GiB.

You consolidate free capacity on a volume group to improve write performance. Your volume group's free capacity will become fragmented over time as the host writes, modifies, and deletes files. Eventually, the available capacity will not be located in a single contiguous block, but will be scattered in small fragments across the volume group. This causes further file fragmentation, since the host must write new files as fragments to fit them into the available ranges of free clusters.

By consolidating free capacity on a selected volume group, you will notice improved file system performance whenever the host writes new files. The consolidation process will also help prevent new files from being fragmented in the future.

Steps

1. From the Manage page, select the storage array with the volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select the volume group with free capacity that you want to consolidate, and then select **Uncommon Tasks > Consolidate volume group free capacity**.

The Consolidate Free Capacity dialog box appears.

4. Type `consolidate` to confirm you want to perform this operation.
5. Click **Consolidate**.

Result

The system begins consolidating (defragmenting) the volume group's free capacity areas into one contiguous amount for subsequent storage configuration tasks.

After you finish

From the navigation sidebar, select **Operations** to view the progress of the Consolidate Free Capacity operation. This operation can be lengthy and could affect system performance.

Turn on locator lights

You can locate drives to physically identify all of the drives that comprise a selected pool, volume group, or SSD Cache. An LED indicator lights up on each drive in the selected pool, volume group, or SSD Cache.

Steps

1. From the Manage page, select the storage array.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select the pool, volume group, or SSD Cache you want to locate, and then click **More > Turn on locator lights**.

A dialog box appears that indicates the lights on the drives comprising the selected pool, volume group, or SSD Cache are turned on.

4. After you successfully locate the drives, click **Turn Off**.

Remove capacity

You can remove drives to decrease the capacity of an existing pool or SSD Cache.

After you remove drives, the data in each volume of the pool or SSD Cache is redistributed to the remaining drives. The removed drives become unassigned, and their capacity becomes part of the total free capacity of the storage array.

About this task

Follow these guidelines when you remove capacity:

- You cannot remove the last drive in an SSD Cache without first deleting the SSD Cache.
- You cannot reduce the number of drives in a pool to be less than 11 drives.

- You can remove a maximum of 12 drives at a time. If you need to remove more than 12 drives, repeat the procedure.
- You cannot remove drives if there is not enough free capacity in the pool or SSD Cache to contain the data, when that data is redistributed to the remaining drives in the pool or SSD Cache.

The following are potential performance impacts:

- Removing drives from a pool or SSD Cache might result in reduced volume performance.
- The preservation capacity is not consumed when you remove capacity from a pool or SSD Cache. However, the preservation capacity might decrease based on the number of drives remaining in the pool or SSD Cache.

The following are impacts to secure-capable drives:

- If you remove the last drive that is not secure-capable, the pool is left with all secure-capable drives. In this situation, you are given the option to enable security for the pool.
- If you remove the last drive that is not Data Assurance (DA) capable, the pool is left with all DA-capable drives.
- Any new volumes that you create on the pool will be DA-capable. If you want existing volumes to be DA-capable, you need to delete and then re-create the volume.

Steps

1. From the Manage page, select the storage array.

Select **Provisioning** › **Configure Pools and Volume Groups**.

2. Select the pool or SSD Cache, and then click **More** › **Remove capacity**.

The Remove Capacity dialog box appears.

3. Select one or more drives in the list.

As you select or de-select drives in the list, the Total capacity selected field updates. This field shows the total capacity of the pool or SSD Cache that results after you remove the selected drives.

4. Click **Remove**, and then confirm you want to remove the drives.

Result

The newly reduced capacity of the pool or SSD Cache is reflected in the Pools and Volume Groups view.

Enable security for a pool or volume group

You can enable Drive Security for a pool or volume group to prevent unauthorized access to the data on the drives contained in the pool or volume group.

Read and write access for the drives is only available through a controller that is configured with a security key.

Before you begin

- The Drive Security feature must be enabled.
- A security key must be created.
- The pool or volume group must be in an Optimal state.

- All of the drives in the pool or volume group must be secure-capable drives.

About this task

If you want to use Drive Security, select a pool or volume group that is secure-capable. A pool or volume group can contain both secure-capable and non-secure-capable drives, but all drives must be secure-capable to use their encryption capabilities.

After enabling security, you can only remove it by deleting the pool or volume group, and then erasing the drives.

Steps

1. From the Manage page, select the storage array with the pool or volume group.
2. Select **Provisioning > Configure Pools and Volume Groups**.
3. Select the pool or volume group on which you want to enable security, and then click **More > Enable security**.

The Confirm Enable Security dialog box appears.

4. Confirm that you want to enable security for the selected pool or volume group, and then click **Enable**.

Remove the Storage Plugin for vCenter

You can remove the plugin from the vCenter Server Appliance and uninstall the plugin webserver from the application host.

These are two distinct steps that you can perform in any order. However, if you choose to remove the plugin webserver from the application host before unregistering the plugin, the registration script is removed during that process and you cannot use Method 1 to unregister.

Unregister the plugin from a vCenter Server Appliance

To unregister the plugin from a vCenter Server Appliance, select one of these methods:

- [Method 1: Execute the registration script](#)
- [Method 2: Use the vCenter Server Mob pages](#)

Method 1: Execute the registration script

1. Open a prompt through the command line and navigate to the following directory:

```
<install directory>\vcenter-register\bin
```

2. Execute the vcenter-register.bat file:

```
vcenter-register.bat ^  
  
-action unregisterPlugin ^  
  
-vcenterHostname <vCenter FQDN> ^  
  
-username <Administrator Username> ^
```

3. Verify that the script is successful.

The logs are saved to `%install_dir%/working/logs/vc-registration.log`.

Method 2: Use the vCenter Server Mob pages

1. Open a web browser and enter the following url:

`https://<FQDN of vCenter Server>/mob`

2. Log in under the administrator credentials.
3. Look for the property name of `extensionManager` and click the link associated with that property.
4. Expand the properties list by clicking the **More...** link at the bottom of the list.
5. Verify that the extension `plugin.netapp.eseries` is in the list.
6. If it is present, then click the Method `UnregisterExtension`.
7. Enter the value `plugin.netapp.eseries` in the dialog and click **Invoke Method**.
8. Close the dialog and refresh the web browser.
9. Verify that the `plugin.netapp.eseries` extension is not on the list.



This procedure unregisters the plugin from the vCenter Server Appliance; however, it does not remove plugin package files from the server. To remove package files, use SSH to access the vCenter Server Appliance and navigate to the following directory: `etc/vmware/vsphere-ui/vc-packages/vsphere-client-serenity/`. Then remove the directory associated with the plugin.

Remove the plugin webserver from the Application host

To remove the plugin software from the application host, follow these steps:

1. From the application server, navigate to the **Control Panel**.
2. Go to **Apps & Features**, and then select **SANtricity Storage Plugin for vCenter**.
3. Click **Uninstall/Change**.

A confirmation dialog opens.

4. Click **Uninstall**.

A confirmation message is displayed when the uninstall is complete.

5. Click **Done**.

FAQs

What settings are imported?

The Import Settings feature is a batch operation that loads configurations from one storage array to multiple storage arrays.

The settings that are imported during this operation depend on how the source storage array is configured in System Manager. The following settings can be imported to multiple storage arrays:

- **Email alerts** — Settings include a mail server address and the email addresses of the alert recipients.
- **Syslog alerts** — Settings include a syslog server address and a UDP port.
- **SNMP alerts** — Settings include a community name and IP address for the SNMP server.
- **AutoSupport** — Settings include the separate features (Basic AutoSupport, AutoSupport OnDemand, and Remote Diagnostics), the maintenance window, delivery method, and dispatch schedule.
- **Directory services** — Configuration includes the domain name and URL of an LDAP (Lightweight Directory Access Protocol) server, along with the mappings for the LDAP server's user groups to the storage array's predefined roles.
- **Storage configuration** — Configurations include volumes (only thick and only non-repository volumes), volume groups, pools, and hot spare drive assignments.
- **System settings** — Configurations include media scan settings for a volume, SSD cache for controllers, and automatic load balancing (does not include host connectivity reporting).

Why do I not see all of my storage arrays?

During the Import Settings operation, some of your storage arrays might not be available in the target selection dialog box.

Storage arrays might not appear for the following reasons:

- The firmware version is below 8.50.
- The storage array is offline.
- The system cannot communicate with that array (for example, the array has certificate, password, or networking problems).

Why are these volumes not associated with a workload?

Volumes are not associated with a workload if they have been created using the command line interface (CLI) or if they have been migrated (imported/exported) from a different storage array.

How does my selected workload impact volume creation?

During volume creation, you are prompted for information about a workload's use. The system uses this information to create an optimal volume configuration for you, which can be edited as needed. Optionally, you can skip this step in the volume creation sequence.

A workload is a storage object that supports an application. You can define one or more workloads, or instances, per application. For some applications, the system configures the workload to contain volumes with similar underlying volume characteristics. These volume characteristics are optimized based on the type of application the workload supports. For example, if you create a workload that supports a Microsoft SQL Server application and then subsequently create volumes for that workload, the underlying volume characteristics are optimized to support Microsoft SQL Server.

- **Application-specific** — When you are creating volumes using an application-specific workload, the system may recommend an optimized volume configuration to minimize contention between application

workload I/O and other traffic from your application instance. Volume characteristics like I/O type, segment size, controller ownership, and read and write cache are automatically recommended and optimized for workloads that are created for the following application types.

- Microsoft SQL Server
- Microsoft Exchange Server
- Video surveillance applications
- VMware ESXi (for volumes to be used with Virtual Machine File System)

You can review the recommended volume configuration and edit, add, or delete the system-recommended volumes and characteristics using the Add/Edit Volumes dialog box.

- **Other (or applications without specific volume creation support)** — Other workloads use a volume configuration that you must manually specify when you want to create a workload that is not associated with a specific application, or if there is no built-in optimization for the application you intend to use on the storage array. You must manually specify the volume configuration using the Add/Edit Volumes dialog box.

Why do I not see all my volumes, hosts, or host clusters?

Snapshot volumes with a DA-enabled base volume are not eligible to be assigned to a host that is not Data Assurance (DA) capable. You must disable DA on the base volume before a snapshot volume can be assigned to a host that is not DA capable.

Consider the following guidelines for the host to which you are assigning the snapshot volume:

- A host is not DA capable if it is connected to the storage array through an I/O interface that is not DA capable.
- A host cluster is not DA capable if it has at least one host member that is not DA capable.



You cannot disable DA on a volume that is associated with snapshots (consistency groups, snapshot groups, snapshot images, and snapshot volumes), volume copies, and mirrors. All associated reserved capacity and snapshot objects must be deleted before DA can be disabled on the base volume.

Why can I not delete the selected workload?

This workload consists of a group of volumes that were created using the command line interface (CLI) or migrated (imported/exported) from a different storage array. As a result, the volumes in this workload are not associated with an application-specific workload, so the workload cannot be deleted.

How do application-specific workloads help me manage my storage array?

The volume characteristics of your application-specific workload dictate how the workload interacts with the components of your storage array and helps determine the performance of your environment under a given configuration.

An application is software such as SQL Server or Exchange. You define one or more workloads to support each application. For some applications, the system automatically recommends a volume configuration that optimizes storage. Characteristics such as I/O type, segment size, controller ownership, and read and write cache are included in the volume configuration.

What do I need to do to recognize the expanded capacity?

If you increase the capacity for a volume, the host might not immediately recognize the increase in volume capacity.

Most operating systems recognize the expanded volume capacity and automatically expand after the volume expansion is initiated. However, some might not. If your OS does not automatically recognize the expanded volume capacity, you might need to perform a disk rescan or reboot.

After you have expanded the volume capacity, you must manually increase the file system size to match. How you do this depends on the file system you are using.

Refer to your host operating system documentation for additional details.

When would I want to use the assign host later selection?

If you want to speed the process for creating volumes, you can skip the host assignment step so that newly created volumes are initialized offline.

Newly created volumes must be initialized. The system can initialize them using one of two modes – either an Immediate Available Format (IAF) background initialization process or an offline process.

When you map a volume to a host, it forces any initializing volumes in that group to transition to background initialization. This background initialization process allows for concurrent host I/O, which can sometimes be time-consuming.

When none of the volumes in a volume group are mapped, offline initialization is performed. The offline process is much faster than the background process.

What do I need to know about host block size requirements?

For EF300 and EF600 systems, a volume can be set to support a 512-byte or a 4KiB block size (also called "sector size"). You must set the correct value during volume creation. If possible, the system suggests the appropriate default value.

Before setting the volume block size, read the following limitations and guidelines.

- Some operating systems and virtual machines (notably VMware, at this time) require a 512-byte block size and do not support 4KiB, so make sure you know the host requirements before creating a volume. Typically, you can achieve the best performance by setting a volume to present a 4KiB block size; however, ensure that your host allows for 4KiB (or "4Kn") blocks.
- The type of drives you select for your pool or volume group also determines what volume block sizes are supported, as follows:
 - If you create a volume group using drives that write to 512-byte blocks, then you can only create volumes with 512-byte blocks.
 - If you create a volume group using drives that write to 4KiB blocks, then you can create volumes with either 512-byte or 4KiB blocks.
- If the array has an iSCSI host interface card, all volumes are limited to 512-byte blocks (regardless of volume group block size). This is due to a specific hardware implementation.
- You cannot change a block size once it is set. If you need to change a block size, you must delete the volume and re-create it.

Why would I need to create a host cluster?

You need to create a host cluster if you want to have two or more hosts share access to the same set of volumes. Normally, the individual hosts have clustering software installed on them to coordinate volume access.

How do I know which host operating system type is correct?

The Host Operating System Type field contains the operating system of the host. You can select the recommended host type from the drop-down list or allow the Host Context Agent (HCA) to configure the host and appropriate host operating system type.

The host types that appear in the drop-down list depend on the storage array model and the firmware version. The most recent versions display the most common options first, which are the most likely to be appropriate. Appearance on this list does not imply the option is fully supported.



For more information about host support, refer to the [NetApp Interoperability Matrix Tool](#).

Some of the following host types might appear in the list:

Host Operating System type	Operating System (OS) and multipath driver
Linux DM-MP (Kernel 3.10 or later)	Supports Linux operating systems using a Device Mapper multipath failover solution with a 3.10 or later Kernel.
VMware ESXi	Supports VMware ESXi operating systems running the Native Multipathing Plug-in (NMP) architecture using the VMware built-in Storage Array Type Policy module SATP_ALUA.
Windows (clustered or non-clustered)	Supports Windows clustered or non-clustered configurations that are not running the ATTO multipathing driver.
ATTO Cluster (all operating systems)	Supports all cluster configurations using the ATTO Technology, Inc., multipathing driver.
Linux (Veritas DMP)	Supports Linux operating systems using a Veritas DMP multipathing solution.
Linux (ATTO)	Supports Linux operating systems using an ATTO Technology, Inc., multipathing driver.
Mac OS	Supports Mac OS versions using an ATTO Technology, Inc., multipathing driver.
Windows (ATTO)	Supports Windows operating systems using an ATTO Technology, Inc., multipathing driver.
FlexArray (ALUA)	Supports a NetApp FlexArray system using ALUA for multipathing.
IBM SVC	Supports an IBM SAN Volume Controller configuration.

Host Operating System type	Operating System (OS) and multipath driver
Factory Default	Reserved for the initial start-up of the storage array. If your host operating system type is set to Factory Default, change it to match the host operating system and multipath driver running on the connected host.
Linux DM-MP (Kernal 3.9 or earlier)	Supports Linux operating systems using a Device Mapper multipath failover solution with a 3.9 or earlier Kernel.
Window Clustered (deprecated)	If your host operating system type is set to this value, use the Windows (clustered or non-clustered) setting instead.

After the HCA is installed and the storage is attached to the host, the HCA sends the host topology to the storage controllers through the I/O path. Based on the host topology, the storage controllers automatically define the host and the associated host ports, and then set the host type.



If the HCA does not select the recommended host type, you must manually set the host type.

How do I match the host ports to a host?

If you are manually creating a host, you first must use the appropriate host bus adapter (HBA) utility available on the host to determine the host port identifiers associated with each HBA installed in the host.

When you have this information, select the host port identifiers that have logged into the storage array from the list provided in the Create Host dialog.



Make sure you select the appropriate host port identifiers for the host you are creating. If you associate the wrong host port identifiers, you might cause unintended access from another host to this data.

If you are automatically creating hosts using the host context agent (HCA) installed on each host, the HCA should automatically associate the host port identifiers with each host and configure them appropriately.

What is the default cluster?

The default cluster is a system-defined entity that allows any unassociated host port identifier that has logged into the storage array to gain access to volumes assigned to the default cluster.

An unassociated host port identifier is a host port that is not logically associated with a particular host but is physically installed in a host and logged into the storage array.



If you want hosts to have specific access to certain volumes in the storage array, you must not use the default cluster. Instead, you must associate the host port identifiers with their corresponding hosts. This task can be done either manually during the Create Host operation or automatically using the host context agent (HCA) installed on each host. Then, you assign volumes either to an individual host or to a host cluster.

You should only use the default cluster in special situations where your external storage environment is conducive to allowing all the hosts and all the logged-in host port identifiers connected to the storage array have access to all of the volumes (all-access mode) without specifically making the hosts known to the storage array or the user interface.

Initially, you can assign volumes only to the default cluster through the command line interface (CLI). However, after you assign at least one volume to the default cluster, this entity (called Default Cluster) is displayed in the user interface where you can then manage this entity.

What is redundancy check?

A redundancy check determines whether the data on a volume in a pool or volume group is consistent. Redundancy data is used to quickly reconstruct information on a replacement drive if one of the drives in the pool or volume group fails.

You can perform this check only on one pool or volume group at a time. A volume redundancy check performs the following actions:

- Scans the data blocks in a RAID 3 volume, a RAID 5 volume, or a RAID 6 volume, and then checks the redundancy information for each block. (RAID 3 can only be assigned to volume groups using the command line interface.)
- Compares the data blocks on RAID 1 mirrored drives.
- Returns redundancy errors if the data is determined to be inconsistent by the controller firmware.



Immediately running a redundancy check on the same pool or volume group might cause an error. To avoid this problem, wait one to two minutes before running another redundancy check on the same pool or volume group.

What is preservation capacity?

Preservation capacity is the amount of capacity (number of drives) that is reserved in a pool to support potential drive failures.

When a pool is created, the system automatically reserves a default amount of preservation capacity depending on the number of drives in the pool.

Pools use preservation capacity during reconstruction, whereas volume groups use hot spare drives for the same purpose. The preservation capacity method is an improvement over hot spare drives because it allows reconstruction to happen faster. Preservation capacity is spread over a number of drives in the pool instead of on one drive in the case of a hot spare drive, so you are not limited by the speed or availability of one drive.

What RAID level is best for my application?

To maximize the performance of a volume group, you must select the appropriate RAID level.

You can determine the appropriate RAID level by knowing the read and write percentages for the applications that are accessing the volume group. Use the Performance page to obtain these percentages.

RAID levels and application performance

RAID relies on a series of configurations, called levels, to determine how user and redundancy data is written and retrieved from the drives. Each RAID level provides different performance features. Applications with a high read percentage will perform well using RAID 5 volumes or RAID 6 volumes because of the outstanding read performance of the RAID 5 and RAID 6 configurations.

Applications with a low read percentage (write-intensive) do not perform as well on RAID 5 volumes or RAID 6 volumes. The degraded performance is the result of the way that a controller writes data and redundancy data to the drives in a RAID 5 volume group or a RAID 6 volume group.

Select a RAID level based on the following information.

RAID 0

Description:

- Non-redundant, striping mode.
- RAID 0 stripes data across all of the drives in the volume group.

Data protection features:

- RAID 0 is not recommended for high availability needs. RAID 0 is better for non-critical data.
- If a single drive fails in the volume group, all of the associated volumes fail, and all data is lost.

Drive number requirements:

- A minimum of one drive is required for RAID Level 0.
- RAID 0 volume groups can have more than 30 drives.
- You can create a volume group that includes all of the drives in the storage array.

RAID 1 or RAID 10

Description:

- Striping/mirror mode.

How it works:

- RAID 1 uses disk mirroring to write data to two duplicate disks simultaneously.
- RAID 10 uses drive striping to stripe data across a set of mirrored drive pairs.

Data protection features:

- RAID 1 and RAID 10 offer high performance and the best data availability.
- RAID 1 and RAID 10 use drive mirroring to make an exact copy from one drive to another drive.
- If one of the drives in a drive pair fails, the storage array can instantly switch to the other drive without any loss of data or service.
- A single drive failure causes associated volumes to become degraded. The mirror drive allows access to the data.
- A drive-pair failure in a volume group causes all of the associated volumes to fail, and data loss could

occur.

Drive number requirements:

- A minimum of two drives is required for RAID 1: one drive for the user data, and one drive for the mirrored data.
- If you select four or more drives, RAID 10 is automatically configured across the volume group: two drives for user data, and two drives for the mirrored data.
- You must have an even number of drives in the volume group. If you do not have an even number of drives and you have some remaining unassigned drives, go to **Pools & Volume Groups** to add additional drives to the volume group, and retry the operation.
- RAID 1 and RAID 10 volume groups can have more than 30 drives. A volume group can be created that includes all of the drives in the storage array.

RAID 5

Description:

- High I/O mode.

How it works:

- User data and redundant information (parity) are striped across the drives.
- The equivalent capacity of one drive is used for redundant information.

Data protection features

- If a single drive fails in a RAID 5 volume group, all of the associated volumes become degraded. The redundant information allows the data to still be accessed.
- If two or more drives fail in a RAID 5 volume group, all of the associated volumes fail, and all data is lost.

Drive number requirements:

- You must have a minimum of three drives in the volume group.
- Typically, you are limited to a maximum of 30 drives in the volume group.

RAID 6

Description:

- High I/O mode.

How it works:

- User data and redundant information (dual parity) are striped across the drives.
- The equivalent capacity of two drives is used for redundant information.

Data protection features:

- If one or two drives fail in a RAID 6 volume group, all of the associated volumes become degraded, but the redundant information allows the data to still be accessed.

- If three or more drives fail in a RAID 6 volume group, all of the associated volumes fail, and all data is lost.

Drive number requirements:

- You must have a minimum of five drives in the volume group.
- Typically, you are limited to a maximum of 30 drives in the volume group.



You cannot change the RAID level of a pool. The user interface automatically configures pools as RAID 6.

RAID levels and data protection

RAID 1, RAID 5, and RAID 6 write redundancy data to the drive media for fault tolerance. The redundancy data might be a copy of the data (mirrored) or an error-correcting code derived from the data. You can use the redundancy data to quickly reconstruct information on a replacement drive if a drive fails.

You configure a single RAID level across a single volume group. All redundancy data for that volume group is stored within the volume group. The capacity of the volume group is the aggregate capacity of the member drives minus the capacity reserved for redundancy data. The amount of capacity needed for redundancy depends on the RAID level used.

Why are some drives not showing up?

In the Add Capacity dialog, not all drives are available for adding capacity to an existing pool or volume group.

Drives are not eligible for any of the following reasons:

- A drive must be unassigned and not secure-enabled. Drives already part of another pool, another volume group, or configured as a hot spare are not eligible. If a drive is unassigned but is secure-enabled, you must manually erase that drive for it to become eligible.
- A drive that is in a non-optimal state is not eligible.
- If the capacity of a drive is too small, it is not eligible.
- The drive media type must match within a pool or volume group. You cannot mix the following:
 - Hard Disk Drives (HDDs) with Solid State Disks (SSDs)
 - NVMe with SAS drives
 - Drives with 512-byte and 4KiB volume block sizes
- If a pool or volume group contains all secure-capable drives, non-secure-capable drives are not listed.
- If a pool or volume group contains all Federal Information Processing Standards (FIPS) drives, non-FIPS drives are not listed.
- If a pool or volume group contains all Data Assurance (DA)-capable drives and there is at least one DA-enabled volume in the pool or volume group, a drive that is not DA capable is not eligible, so it cannot be added to that pool or volume group. However, if there is no DA-enabled volume in the pool or volume group, a drive that is not DA capable can be added to that pool or volume group. If you decide to mix these drives, keep in mind that you cannot create any DA-enabled volumes.



Capacity can be increased in your storage array by adding new drives or by deleting pools or volume groups.

Why can I not increase my preservation capacity?

If you have created volumes on all available usable capacity, you might not be able to increase preservation capacity.

Preservation capacity is the amount of capacity (number of drives) that is reserved on a pool to support potential drive failures. When a pool is created, the system automatically reserves a default amount of preservation capacity depending on the number of drives in the pool. If you have created volumes on all available usable capacity, you cannot increase preservation capacity without adding capacity to the pool by either adding drives or deleting volumes.

You can change the preservation capacity from Pools & Volume Groups. Select the pool that you want to edit. Click **View/Edit Settings**, and then select the **Settings** tab.



Preservation capacity is specified as a number of drives, even though the actual preservation capacity is distributed across the drives in the pool.

What is Data Assurance?

Data Assurance (DA) implements the T10 Protection Information (PI) standard, which increases data integrity by checking for and correcting errors that might occur as data is transferred along the I/O path.

The typical use of the Data Assurance feature will check the portion of the I/O path between the controllers and drives. DA capabilities are presented at the pool and volume group level.

When this feature is enabled, the storage array appends error-checking codes (also known as cyclic redundancy checks or CRCs) to each block of data in the volume. After a data block is moved, the storage array uses these CRC codes to determine if any errors occurred during transmission. Potentially corrupted data is neither written to disk nor returned to the host. If you want to use the DA feature, select a pool or volume group that is DA capable when you create a new volume (look for **Yes** next to **DA** in the pool and volume group candidates table).

Make sure you assign these DA-enabled volumes to a host using an I/O interface that is capable of DA. I/O interfaces that are capable of DA include Fibre Channel, SAS, iSCSI over TCP/IP, NVMe/FC, NVMe/IB, NVMe/RoCE and iSER over InfiniBand (iSCSI Extensions for RDMA/ IB). DA is not supported by SRP over InfiniBand.

What is FDE/FIPS security?

FDE/FIPS security refers to secure-capable drives that encrypt data during writes and decrypt data during reads using a unique encryption key.

These secure-capable drives prevent unauthorized access to the data on a drive that is physically removed from the storage array. Secure-capable drives can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives. FIPS drives have undergone certification testing.



For volumes that require FIPS support, use only FIPS drives. Mixing FIPS and FDE drives in a volume group or pool will result in all drives being treated as FDE drives. Also, an FDE drive cannot be added to or used as a spare in an all-FIPS volume group or pool.

What is secure-capable (Drive Security)?

Drive Security is a feature that prevents unauthorized access to data on secure-enabled drives when removed from the storage array.

These drives can be either Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives.

How do I view and interpret all SSD Cache statistics?

You can view nominal statistics and detailed statistics for SSD Cache.

Nominal statistics are a subset of the detailed statistics. The detailed statistics can be viewed only when you export all SSD statistics to a .csv file. As you review and interpret the statistics, keep in mind that some interpretations are derived by looking at a combination of statistics.

Nominal statistics

To view SSD Cache statistics, go to the **Manage** page. Select **Provisioning > Configure Pools & Volume Groups**. Select the SSD Cache that you want to view statistics for, and then select **More > View Statistics**. The nominal statistics are shown on the View SSD Cache Statistics dialog.



This feature is not available on the EF600 or EF300 storage system.

The list includes nominal statistics, which are a subset of the detailed statistics.

Detailed statistics

The detailed statistics consist of the nominal statistics, plus additional statistics. These additional statistics are saved along with the nominal statistics, but unlike the nominal statistics, they do not display in the View SSD Cache Statistics dialog. You can view the detailed statistics only after exporting the statistics to a .csv file.

The detailed statistics are listed after the nominal statistics.

What is shelf loss protection and drawer loss protection?

Shelf loss protection and drawer loss protection are attributes of pools and volume groups that allow you to maintain data access in the event of a single shelf or drawer failure.

Shelf loss protection

A shelf is the enclosure that contains either the drives or the drives and the controller. Shelf loss protection guarantees accessibility to the data on the volumes in a pool or volume group if a total loss of communication occurs with a single drive shelf. An example of total loss of communication might be loss of power to the drive shelf or failure of both I/O modules (IOMs).



Shelf loss protection is not guaranteed if a drive has already failed in the pool or volume group. In this situation, losing access to a drive shelf and consequently another drive in the pool or volume group causes loss of data.

The criteria for shelf loss protection depends on the protection method, as described in the following table.

Level	Criteria for shelf loss protection	Minimum number of shelves required
Pool	The pool must include drives from at least five shelves and there must be an equal number of drives in each shelf. Shelf loss protection is not applicable to high-capacity shelves; if your system contains high-capacity shelves, refer to Drawer Loss Protection.	5
RAID 6	The volume group contains no more than two drives in a single drawer.	3
RAID 3 or RAID 5	Each drive in the volume group is located in a separate shelf.	3
RAID 1	Each drive in a RAID 1 pair must be located in a separate shelf.	2
RAID 0	Cannot achieve Shelf Loss Protection.	Not applicable

Drawer loss protection

A drawer is one of the compartments of a shelf that you pull out to access the drives. Only the high-capacity shelves have drawers. Drawer loss protection guarantees accessibility to the data on the volumes in a pool or volume group if a total loss of communication occurs with a single drawer. An example of total loss of communication might be loss of power to the drawer or failure of an internal component within the drawer.



Drawer loss protection is not guaranteed if a drive has already failed in the pool or volume group. In this situation, losing access to a drawer (and consequently another drive in the pool or volume group) causes loss of data.

The criteria for drawer loss protection depends on the protection method, as described in the following table:

Level	Criteria for drawer loss protection	Minimum number of drawers required
Pool	<p>Pool candidates must include drives from all drawers, and there must be an equal number of drives in each drawer.</p> <p>The pool must include drives from at least five drawers and there must be an equal number of drives in each drawer.</p> <p>A 60-drive shelf can achieve Drawer Loss Protection when the pool contains 15, 20, 25, 30, 35, 40, 45, 50, 55, or 60 drives. Increments in multiples of 5 can be added to the pool after initial creation.</p>	5

Level	Criteria for drawer loss protection	Minimum number of drawers required
RAID 6	The volume group contains no more than two drives in a single drawer.	3
RAID 3 or 5	Each drive in the volume group is located in a separate drawer	3
RAID 1	Each drive in a mirrored pair must be located in a separate drawer.	2
RAID 0	Cannot achieve Drawer Loss Protection.	Not applicable

How do I maintain shelf and drawer loss protection?

To maintain shelf and drawer loss protection for a pool or volume group, use the criteria specified in the following table.

Level	Criteria for shelf/drawer loss protection	Minimum number of shelves/drawers required
Pool	For shelves, the pool must contain no more than two drives in a single shelf. For drawers, the pool must include an equal number of drives from each drawer.	6 for shelves 5 for drawers
RAID 6	The volume group contains no more than two drives in a single shelf or drawer.	3
RAID 3 or RAID 5	Each drive in the volume group is located in a separate shelf or drawer.	3
RAID 1	Each drive in a mirrored pair must be located in a separate shelf or drawer.	2
RAID 0	Cannot achieve shelf/drawer loss protection.	Not applicable



Shelf/drawer loss protection is not maintained if a drive has already failed in the pool or volume group. In this situation, losing access to a drive shelf or drawer, and consequently another drive in the pool or volume group, causes loss of data.

What is optimization capacity for pools?

SSD drives will have longer life and better maximum write performance when a portion of their capacity is unallocated.

For drives associated with a pool, unallocated capacity is comprised of a pool's preservation capacity, the free

capacity (capacity not used by volumes), and a portion of the usable capacity set aside as additional optimization capacity. The additional optimization capacity ensures a minimum level of optimization capacity by reducing the usable capacity, and as such, is not available for volume creation.

When a pool is created, a recommended optimization capacity is generated that provides a balance of performance, drive wear life, and available capacity. The Additional Optimization Capacity slider located in the Pool Settings dialog allows adjustments to the pool's optimization capacity. Adjusting the slider provides for better performance and drive wear life at the expense of available capacity, or additional available capacity at the expense of performance and drive wear life.



The Additional Optimization Capacity slider is only available for EF600 and EF300 storage systems.

What is optimization capacity for volume groups?

SSD drives will have longer life and better maximum write performance when a portion of their capacity is unallocated.

For drives associated with a volume group, unallocated capacity is comprised of a volume group's free capacity (capacity not used by volumes), and a portion of the usable capacity set aside as optimization capacity. The additional optimization capacity ensures a minimum level of optimization capacity by reducing the usable capacity, and as such, is not available for volume creation.

When a volume group is created, a recommended optimization capacity is generated that provides a balance of performance, drive wear life, and available capacity. The Additional Optimization Capacity slider in the Volume Group Settings dialog allows adjustments to a volume group's optimization capacity. Adjusting the slider provides for better performance and drive wear life at the expense of available capacity, or additional available capacity at the expense of performance and drive wear life.



Additional Optimization Capacity slider is only available for EF600 and EF300 storage systems.

What is resource provisioning capable?

Resource Provisioning is a feature available in the EF300 and EF600 storage arrays, which allows volumes to be put in use immediately with no background initialization process.

A resource-provisioned volume is a thick volume in an SSD volume group or pool, where drive capacity is allocated (assigned to the volume) when the volume is created, but the drive blocks are deallocated (unmapped). By comparison, in a traditional thick volume, all drive blocks are mapped or allocated during a background volume initialization operation in order to initialize the Data Assurance protection information fields and to make data and RAID parity consistent in each RAID stripe. With a resource provisioned volume, there is no time-bound background initialization. Instead, each RAID stripe is initialized upon the first write to a volume block in the stripe.

Resource-provisioned volumes are supported only on SSD volume groups and pools, where all drives in the group or pool support the NVMe Deallocated or Unwritten Logical Block Error Enable (DULBE) error recovery capability. When a resource-provisioned volume is created, all drive blocks assigned to the volume are deallocated (unmapped). In addition, hosts can deallocate logical blocks in the volume using the NVMe Dataset Management command. Deallocating blocks can improve SSD wear life and increase maximum write performance. The improvement varies with each drive model and capacity.

What do I need to know about the resource-provisioned volumes feature?

Resource Provisioning is a feature available in the EF300 and EF600 storage arrays, which allows volumes to be put in use immediately with no background initialization process.



The Resource Provisioning capability is not available at this time. In some views, components might be reported as resource-provisioning capable, but the ability to create resource-provisioned volumes has been disabled until it can be re-enabled in a future update.

Resource-provisioned volumes

A resource-provisioned volume is a thick volume in an SSD volume group or pool, where drive capacity is allocated (assigned to the volume) when the volume is created, but the drive blocks are deallocated (unmapped). By comparison, in a traditional thick volume, all drive blocks are mapped or allocated during a background volume initialization operation in order to initialize the Data Assurance protection information fields and to make data and RAID parity consistent in each RAID stripe. With a resource provisioned volume, there is no time-bound background initialization. Instead, each RAID stripe is initialized upon the first write to a volume block in the stripe.

Resource-provisioned volumes are supported only on SSD volume groups and pools, where all drives in the group or pool support the NVMe Deallocated or Unwritten Logical Block Error Enable (DULBE) error recovery capability. When a resource-provisioned volume is created, all drive blocks assigned to the volume are deallocated (unmapped). In addition, hosts can deallocate logical blocks in the volume using the NVMe Dataset Management command. Deallocating blocks can improve SSD wear life and increase maximum write performance. The improvement varies with each drive model and capacity.

Enabling and disabling the feature

Resource provisioning is enabled by default on systems where the drives support DULBE. You can disable that default setting from Pools & Volume Groups. Disabling resource provisioning is a permanent action for existing volumes and cannot be reversed (i.e., you cannot re-enable resource provisioning for these volume groups and pools).

However, if you want to re-enable resource provisioning again for any new volumes you create, you can do so from **Settings > System**. Be aware that when you re-enable resource provisioning, only newly created volume groups and pools are affected. Any existing volume groups and pools will remain unchanged. If desired, you can also disable resource provisioning again from **Settings > System**.

What is the difference between internal security key and external security key management?

When you implement the Drive Security feature, you can use an internal security key or an external security key to lock down data when a secure-enabled drive is removed from the storage array.

A security key is a string of characters, which is shared between the secure-enabled drives and controllers in a storage array. Internal keys are maintained on the controller's persistent memory. External keys are maintained on a separate key management server, using a Key Management Interoperability Protocol (KMIP).

What do I need to know before creating a security key?

A security key is shared by controllers and secure-enabled drives within a storage array.

If a secure-enabled drive is removed from the storage array, the security key protects the data from unauthorized access.

You can create and manage security keys using one of the following methods:

- Internal key management on the controller's persistent memory.
- External key management on an external key management server.

Internal key management

Internal keys are maintained and "hidden" in a non-accessible location on the controller's persistent memory. Before creating an internal security key, you must do the following:

1. Install secure-capable drives in the storage array. These drives can be Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives.
2. Make sure the Drive Security feature is enabled. If necessary, contact your storage vendor for instructions on enabling the Drive Security feature.

You can then create an internal security key, which involves defining an identifier and a pass phrase. The identifier is a string that is associated with the security key, and is stored on the controller and on all drives associated with the key. The pass phrase is used to encrypt the security key for backup purposes. When you are finished, the security key is stored on the controller in a non-accessible location. You can then create secure-enabled volume groups or pools, or you can enable security on existing volume groups and pools.

External key management

External keys are maintained on a separate key management server, using a Key Management Interoperability Protocol (KMIP). Before creating an external security key, you must do the following:

1. Install secure-capable drives in the storage array. These drives can be Full Disk Encryption (FDE) drives or Federal Information Processing Standard (FIPS) drives.
2. Make sure the Drive Security feature is enabled. If necessary, contact your storage vendor for instructions on enabling the Drive Security feature
3. Obtain a signed, client certificate file. A client certificate validates the storage array's controllers, so the key management server can trust their KMIP requests.
 - a. First, you complete and download a client Certificate Signing Request (CSR). Go to **Settings > Certificates > Key Management > Complete CSR**.
 - b. Next, you request a signed client certificate from a CA that is trusted by the key management server. (You can also create and download a client certificate from the key management server using the downloaded CSR file.)
 - c. Once you have a client certificate file, copy that file to the host where you are accessing System Manager.
4. Retrieve a certificate file from the key management server, and then copy that file to the host where you are accessing System Manager. A key management server certificate validates the key management server, so the storage array can trust its IP address. You can use a root, intermediate, or server certificate for the key management server.

You can then create an external key, which involves defining the IP address of the key management server and the port number used for KMIP communications. During this process, you also load certificate files. When you are finished, the system connects to the key management server with the credentials you entered. You can then create secure-enabled volume groups or pools, or you can enable security on existing volume groups and

pools.

Why do I need to define a pass phrase?

The pass phrase is used to encrypt and decrypt the security key file stored on the local management client. Without the pass phrase, the security key cannot be decrypted and used to unlock data from a secure-enabled drive if it is re-installed in another storage array.

Legacy solutions

Cloud connector

Overview of the SANtricity® Cloud Connector

The SANtricity Cloud Connector is a host-based Linux application that enables you to perform full block-based and file-based backup and recovery of E-Series volumes to S3 complaint accounts (for example, Amazon Simple Storage Service and NetApp StorageGRID) and NetApp AltaVault appliance.

Available for installation on RedHat and SUSE Linux platforms, the SANtricity Cloud Connector is a packaged solution (.bin file). After you install SANtricity Cloud Connector, you can configure the application to perform backup and restore jobs for E-Series volumes to an AltaVault appliance or to your existing Amazon S3 or StorageGRID accounts. All jobs performed through the SANtricity Cloud Connector use REST-based APIs.



The SANtricity Cloud Connector tool is deprecated and no longer available for download.

Considerations

When using these procedures, be aware that:

- Configuration and backup/restore jobs described in these procedures apply to the graphical user interface version of the SANtricity Cloud Connector.
- REST API workflows for the SANtricity Cloud Connector application are not described in these procedures. For experienced developers, endpoints are available for each SANtricity Cloud Connector operation under the API documentation. The API documentation is accessible by navigating to <http://<hostname.domain>:<port>/docs> through a browser.

Types of backups

The SANtricity Cloud Connector provides two types of backups: image-based and file-based backups.

- **Image-based backup**

An image-based backup reads the raw data blocks from a snapshot volume and backs them up to a file known as an image. All of the data blocks on the snapshot volume are backed up, including empty blocks, blocks occupied by deleted files, blocks associated with partitioning, and filesystem metadata. Image backups have the advantage of storing all information with the snapshot volume regardless of the partitioning scheme or filesystems on it.

The image is not stored on the Backup Target as a single file, but is instead broken up into a series of data

chunks, which are 64MB in size. The data chunks allow SANtricity Cloud Connector to use multiple connections to the backup target, thereby improving the performance of the backup process.

For backups to StorageGRID and Amazon Web Services (S3), each data chunk uses a separate encryption key to encrypt the chunk. The key is a SHA256 hash consisting of the combination of a user supplied passphrase and the SHA256 hash of the user data. For backups to AltaVault, SANtricity Cloud Connector does not encrypt the data chunks as AltaVault performs this operation.

- **File-based backup**

A file-based backup reads the files contained with a filesystem partition and backs them up into a series of data chunks that are 64MB in size. A file-based backup does not back up deleted files or partitioning and filesystem metadata. As with image-based backups, the data chunks allow SANtricity Cloud Connector to use multiple connections to the backup target, thereby improving performance of the backup process.

For backups to StorageGRID and Amazon Web Services, each data chunk uses a separate encryption key to encrypt the chunk. The key is a SHA256 hash consisting of the combination of user-supplied passphrase and the SHA256 hash of the user data. For backups to AltaVault, the data chunks are not encrypted by SANtricity Cloud Connector because AltaVault performs this operation.

System requirements for Cloud Connector

Your system must meet compatibility requirements for the SANtricity Cloud Connector.

Host hardware requirements

Your hardware must meet the following minimum requirements:

- At least 5 GB of memory; 4 GB for the maximum configured heap size
- At least 5 GB of free disk space is required from the software installation

You must install the SANtricity Web Services Proxy to use the SANtricity Cloud Connector. You can install the Web Services Proxy locally or you can run the application remotely on a different sever. For information on installing the SANtricity Web Services Proxy, see the [Web Services Proxy topics](#).

Supported browsers

The following browsers are supported with the SANtricity Cloud Connector application (minimum versions noted):

- Firefox v31
- Google Chrome v47
- Microsoft Internet Explorer v11
- Microsoft Edge, EdgeHTML 12
- Safari v9



API documentation for the SANtricity Cloud Connector application will not load when using the Compatibility View setting within the Microsoft Internet Explorer v11 browser. To ensure the API documentation displays properly under the Microsoft Internet Explorer v11 browser, it is recommended that the Compatibility View setting is disabled.

Compatible storage arrays and controller firmware

You should verify the compatibility of your storage arrays and firmware before using the SANtricity Cloud Connector application.

For a complete and up-to-date listing of all compatible storage arrays and firmware for the SANtricity Cloud Connector, see the [NetApp Interoperability Matrix Tool](#).

Compatible operating systems

The SANtricity Cloud Connector 4.0 application is compatible with and supported on the following operating systems:

Operating System	Version	Architecture
Red Hat Enterprise Linux (RHEL)	7.x	64 bit
SUSE Linux Enterprise Server (SLES)	12.x	64 bit

Supported file systems

You must use supported file systems to perform backups and restores through the SANtricity Cloud Connector application.

The following file systems are supported for backup and restore operations under the SANtricity Cloud Connector application:

- ext2
- ext3
- ext4

Install SANtricity Cloud Connector

The SANtricity Cloud Connector packaged solution (.bin file) is available for RedHat and SUSE Linux platforms only.

You can install the SANtricity Cloud Connector application through graphical mode or console mode on a compatible Linux operating system. During the installation process, you must specify the non-SSL and SSL port numbers for the SANtricity Cloud Connector. When installed, the SANtricity Cloud Connector runs as a daemon process.



The SANtricity Cloud Connector tool is deprecated and no longer available for download.

Before you begin

Review the following notes:

- If SANtricity Web Services Proxy is already installed on the same server as the SANtricity Cloud Connector, conflicts will occur between non-SSL port numbers and SSL port numbers conflicts. In this case, choose appropriate numbers for the non-SSL port and the SSL port during the SANtricity Cloud Connector installation.

- If any hardware changes are performed on your host, re-install the SANtricity Cloud Connector application to ensure encryption consistency.
- Backups created through version 3.1 of the SANtricity Cloud Connector application are not compatible with version 4.0 of the SANtricity Cloud Connector application. If you intend to maintain these backups, you must continue to use your previous version of the SANtricity Cloud Connector. To ensure successful installation of separate 3.1 and 4.0 releases of the SANtricity Cloud Connector, unique port numbers must be assigned for each version of the application.

Install Device Mapper Multipath (DM-MP)

Any host running the SANtricity Cloud Connector also must run Linux Device Mapper Multipath (DM-MP) and have the multipath-tools package installed.

The SANtricity Cloud Connector discovery process relies on the multipath tools package for discovery and recognition of the volumes and files to backup or restore. For more information on how to set up and configure the Device Mapper, see the *SANtricity Storage Manager Multipath Drivers Guide* for the release of SANtricity you are using under the [E-Series and SANtricity Document Resources](#).

Install Cloud Connector

You can install SANtricity Cloud Connector on Linux operating systems in either graphical mode or console mode.

Graphical mode

You can use graphical mode to install the SANtricity Cloud Connector on a Linux operating system.

Before you begin

Designate a host location for the SANtricity Cloud Connector installation.

Steps

1. Download the SANtricity Cloud Connector installation file to the desired host location.
2. Open a terminal window.
3. Navigate to the directory file containing the SANtricity Cloud Connector installation file.
4. Start the SANtricity Cloud Connector installation process:

```
./cloudconnector-xxxx.bin -i gui
```

In this command, `xxxx` designates the version number of the application.

The Installer window is displayed.

5. Review the Introduction statement, and then click **Next**.

The License Agreement for NetApp, Inc. Software is displayed within the installer window.

6. Accept the terms of the License Agreement, and then click **Next**.

The Backups created with previous releases of SANtricity Cloud Connector page is displayed.

7. To acknowledge the Backups created with previous releases of SANtricity Cloud Connector message, click

Next.



To install version 4.0 of the SANtricity Cloud Connector while maintaining a previous version, unique port numbers must be assigned for each version of the application.

The Choose Install page is displayed within the Installer window. The Where Would You Like to Install field displays the following default install folder: `opt/netapp/santricity_cloud_connector4/`

8. Choose one of the following options:

- To accept the default location, click **Next**.
- To change the default location, enter a new folder location.
An Enter the Non SSL Jetty Port Number page is displayed. A default value of 8080 is assigned to the non-SSL port.

9. Choose one of the following options:

- To accept the default SSL port number, click **Next**.
- To change the default SSL port number, enter the new desired port number value.

10. Choose one of the following options:

- To accept the default Non SSL port number, click **Next**.
- To change the default Non SSL port number, enter the new desired port number value.
The Pre-Installation Summary page is displayed.

11. Review the displayed Pre-Installation Summary, and then click **Install**.

The installation of the SANtricity Cloud Connector begins and a Webserver Daemon Setup prompt is displayed.

12. Click **OK** to acknowledge the Webserver Daemon Setup prompt.

The Installation Complete message is displayed.

13. Click **Done** to exit the SANtricity Cloud Connector installer.

Console mode

You can use the console mode to install the SANtricity Cloud Connector on a Linux operating system.

Before you begin

Designate a host location for the SANtricity Cloud Connector installation.

Steps

1. Download the SANtricity Cloud Connector installation file to the desired IO host location.
2. Open a terminal window.
3. Navigate to the directory file containing the SANtricity Cloud Connector installation file.
4. Start the SANtricity Cloud Connector installation process:

```
./cloudconnector-xxxx.bin -i console
```

In this command, `xxxx` indicates the version number of the application.

The installation process for the SANtricity Cloud Connector is initialized.

5. Press **Enter** to proceed with the installation process.

The End User License Agreement for NetApp, Inc. Software is displayed within the installer window.



To cancel the installation process at any time, type `quit` under the installer window.

6. Press **Enter** to proceed through each portion of the End User License Agreement.

The License Agreement acceptance statement is displayed under the installer window.

7. To accept the terms of the End User License Agreement and proceed with the installation of the SANtricity Cloud Connector, enter `y` and press **Enter** under the installer window.

The Backups created with previous releases of SANtricity Cloud Connector page is displayed.



If you do not accept the terms of the End User Agreement, type `n` and press **Enter** to terminate the installation process for the SANtricity Cloud Connector.

8. To acknowledge the Backups created with previous releases of SANtricity Cloud Connector message, press **Enter**.



To install version 4.0 of the SANtricity Cloud Connector while maintaining a previous version, unique port numbers must be assigned for each version of the application.

A Choose Install Folder message with the following default install folder for the SANtricity Cloud Connector is displayed: `/opt/netapp/santricity_cloud_connector4/`.

9. Choose one of the following options:

- To accept the default install location, press **Enter**.
- To change the default install location, enter the new folder location.
An Enter the Non SSL Jetty Port Number message is displayed. A default value of 8080 is assigned to the Non SSL port.

10. Choose one of the following options:

- To accept the default SSL port number, press **Next**.
- To change the default SSL port number, enter the new desired port number value.

11. Choose one of the following options:

- To accept the default Non SSL port number, press **Enter**.
- To change the default Non SSL port number, enter the new port number value.
The Pre-Installation Summary for the SANtricity Cloud Connector is displayed.

12. Review the displayed Pre-Installation Summary, and press **Enter**.

13. Press **Enter** to acknowledge the Webserver Daemon Setup prompt.

The Installation Complete message is displayed.

14. Press **Enter** to exit the SANtricity Cloud Connector installer.

Add server certificate and CA certificate into a keystore

To use a secure https connection from the browser to the SANtricity Cloud Connector host, you can accept the self-signed certificate from the SANtricity Cloud Connector host or add a certificate and a trust chain recognized by both the browser and the SANtricity Cloud Connector application.

Before you begin

The SANtricity Cloud Connector application must be installed on a host.

Steps

1. Stop the service using the `systemctl` command.
2. From the default install location, access the working directory.



The default install location for the SANtricity Cloud Connector is `/opt/netapp/santricity_cloud_connector4`.

3. Using the `keytool` command, create your server certificate, and certificate signing request (CSR).

EXAMPLE

```
keytool -genkey -dname "CN=host.example.com, OU=Engineering, O=Company, L=<CITY>, S=<STATE>, C=<COUNTRY>" -alias cloudconnect -keyalg "RSA" -sigalg SHA256withRSA -keysize 2048 -validity 365 -keystore keystore_cloudconnect.jks -storepass changeit
keytool -certreq -alias cloudconnect -keystore keystore_cloudconnect.jks -storepass changeit -file cloudconnect.csr
```

4. Send the generated CSR to the certificate authority (CA) of your choosing.

The certificate authority signs the certificate request and returns a signed certificate. In addition, you receive a certificate from the CA itself. This CA certificate must be imported into your keystore.

5. Import the certificate and the CA certificate chain into the application keystore: `<install Path>/working/keystore`

EXAMPLE

```
keytool -import -alias ca-root -file root-ca.cer -keystore keystore_cloudconnect.jks -storepass <password> -noprompt
keytool -import -alias ca-issuing-1 -file issuing-ca-1.cer -keystore keystore_cloudconnect.jks -storepass <password> -noprompt
keytool -import -trustcacerts -alias cloudconnect -file certnew.cer -keystore keystore_cloudconnect.jks -storepass <password>
```

6. Restart the service.

Add StorageGRID certificate into a keystore

If you are configuring StorageGRID as the target type for the SANtricity Cloud Connector application, you must first add a StorageGRID certificate into the SANtricity Cloud Connector keystore.

Before you begin

- You have a signed StorageGRID certificate.
- You have the SANtricity Cloud Connector application installed on a host.

Steps

1. Stop the service using the `systemctl` command.
2. From the default install location, access the working directory.



The default install location for the SANtricity Cloud Connector is `/opt/netapp/santricity_cloud_connector4`.

3. Import the StorageGRID certificate into the application keystore: `<install Path>/working/keystore`

EXAMPLE

```
opt/netapp/santricity_cloud_connector4/jre/bin/keytool -import
-trustcacerts -storepass changeit -noprompt -alias StorageGrid_SSL -file
/home/ictlabs01.cer -keystore
/opt/netapp/santricity_cloud_connector/jre/lib/security/cacerts
```

4. Restart the service.

Configure the SANtricity Cloud Connector for the first time

Upon successful installation, you can set up the SANtricity Cloud Connector application through the configuration wizard. The configuration wizard is displayed after you initially log in to the SANtricity Cloud Connector.

Log in to the SANtricity Cloud Connector for the first time

When initializing the SANtricity Cloud Connector for the first time, you must enter a default password to access the application.

Before you begin

Make sure you have access to an internet-connected browser.

Steps

1. Open a supported browser.
2. Connect to the configured SANtricity Cloud Connector server (e.g., `http://localhost:8080/`).

The initial login page for the SANtricity Cloud Connector application is displayed.

3. In the Administrator Password field, enter the default password of `password`.

4. Click **Log In**.

The SANtricity Cloud Connector Configuration Wizard is displayed.

Using the Configuration Wizard

The Configuration Wizard is displayed upon successful initial login to the SANtricity Cloud Connector.

Through the Configuration Wizard, you set up the administrator password, Web Services Proxy login management credentials, desired backup target type, and encryption pass phrase for the SANtricity Cloud Connector.

Step 1: Set administrator password

You can customize the password used for subsequent logins to the SANtricity Cloud Connector through the Set Administrator Password page.

Establishing a password through the Set Administrator Password page effectively replaces the default password used during the initial login for the SANtricity Cloud Connector application.

Steps

1. On the Set Administrator Password page, enter the desired login password for the SANtricity Cloud Connector in the **Enter the new administrator password** field.
2. In the **Re-enter the new administrator password** field, re-enter the password from first field.
3. Click **Next**.

The password setup for the SANtricity Cloud Connector is accepted and the Set Pass Phrase page is displayed under the Configuration Wizard.



The user defined administrator password is not set until you complete the configuration wizard.

Step 2: Set pass phrase

Under the Enter the Encryption Pass Phrase page, you can specify an alphanumeric pass phrase between 8 and 32 characters.

A user-specified pass phrase is required as part of the data encryption key used by the SANtricity Cloud Connector application.

Steps

1. In the **Define a pass phrase** field, enter the desired pass phrase.
2. In the **Re-enter your pass phrase** field, re-enter the pass phrase from the first field.
3. Click **Next**.

The entered pass phrase for the SANtricity Cloud Connector application is accepted and the Select Target Type page for the configuration wizard is displayed.

Step 3: Select target type

Backup and restore capabilities are available for Amazon S3, AltaVault, and StorageGRID target types through the SANtricity Cloud Connector. You can specify the desired storage target type for the SANtricity Cloud Connector application under the Select the Target Type page.

Before you begin

Make sure you have one of the following: AltaVault mount point, Amazon AWS account, or StorageGRID account.

Steps

1. In the dropdown menu, select one of the following options:

- Amazon AWS
- AltaVault
- StorageGRID

A Target Type page for the selected option is displayed in the Configuration Wizard.

2. Refer to the appropriate configuration instructions for AltaVault, Amazon AWS, or StorageGRID.

Configure AltaVault appliance

After selecting the AltaVault appliance option under the Select the Target Type page, configuration options for the AltaVault target type are displayed.

Before you begin

- You have the NFS mount path for an AltaVault appliance.
- You specified AltaVault appliance as the target type.

Steps

1. In the **NFS Mount Path** field, enter the mount point for the AltaVault target type.



Values in the **NFS Mount Path** field must follow the Linux path format.

2. Select the **Save a backup of the configuration database on this target** check box to create a backup of the configuration database on the selected target type.



If an existing database configuration is detected on the specified target type when testing the connection, you have the option of replacing the existing database configuration information on the SANtricity Cloud Connector host with the new backup information entered under the configuration wizard.

3. Click **Test Connection** to test the connection for the specified AltaVault settings.
4. Click **Next**.

The specified target type for the SANtricity Cloud Connector is accepted and the Web Services Proxy page is displayed in the Configuration Wizard.

5. Proceed to "Step 4: Connect to Web Services Proxy."

Configure Amazon AWS account

After selecting the Amazon AWS option under the Select the Target Type page, configuration options for the Amazon AWS target type are displayed.

Before you begin

- You have an established Amazon AWS account.
- You specified Amazon AWS as the target type.

Steps

1. In the **Access Key ID** field, enter the access ID for the Amazon AWS target.
2. In the **Secret Access Key** field, enter the secret access key for the target.
3. In the **Bucket Name** field, enter the bucket name for the target.
4. Select the **Save a backup of the configuration database on this target** checkbox to create a backup of the configuration database on the selected target type.



It is recommended you enable this setting to ensure that data from the backup target can be restored if the database is lost.



If an existing database configuration is detected on the specified target type when testing the connection, you have the option of replacing the existing database configuration information on the SANtricity Cloud Connector host with the new backup information entered under the configuration wizard.

5. Click **Test Connection** to verify the entered Amazon AWS credentials.
6. Click **Next**.

The specified target type for the SANtricity Cloud Connector is accepted, and the Web Services Proxy page is displayed under the Configuration Wizard.

7. Proceed to "Step 4: Connect to Web Services Proxy."

Configure StorageGRID account

After selecting the StorageGRID option under the Select the Target Type page, configuration options for the StorageGRID target type are displayed.

Before you begin

- You have an established StorageGRID account.
- You have a signed StorageGRID certificate in the SANtricity Cloud Connector keystore.
- You specified StorageGRID as the target type.

Steps

1. In the **URL** field, enter the URL for the Amazon S3 cloud service
2. In the **Access Key ID** field, enter the access ID for the S3 target.
3. In the **Secret Access Key** field, enter the secret access key for the S3 target.
4. In the **Bucket Name** field, enter the bucket name for the S3 target.
5. To use path style access, select the **Use path-style access** checkbox.



If unchecked, virtual host-style access is used.

6. Select the **Save a backup of the configuration database on this target** checkbox to create a backup of the configuration database on the selected target type.



It is recommended you enable this setting to ensure that data from the backup target can be restored if the database is lost.



If an existing database configuration is detected on the specified target type when testing the connection, you have the option of replacing the existing database configuration information on the SANtricity Cloud Connector host with the new backup information entered in the configuration wizard.

7. Click **Test Connection** to verify the entered S3 credentials.



Some S3-compliant accounts may require secured HTTP connections. For information on placing a StorageGRID certificate in the keystore, see [Add StorageGRID certificate into a keystore](#).

8. Click **Next**.

The specified target type for the SANtricity Cloud Connector is accepted and the Web Services Proxy page is displayed under the Configuration Wizard.

9. Proceed to "Step 4: Connect to Web Services Proxy."

Step 4: Connect to Web Services Proxy

Login and connection information for the Web Services Proxy used in conjunction with the SANtricity Cloud Connector is entered through the Enter Web Services Proxy URL and Credentials page.

Before you begin

Make sure you have an established connection to the SANtricity Web Services Proxy.

Steps

1. In the **URL** field, enter the URL for the Web Services Proxy used for the SANtricity Cloud Connector.
2. In the **User Name** field, enter the user name for the Web Services Proxy connection.
3. In the **Password** field, enter the password for the Web Services Proxy connection.
4. Click **Test Connection** to verify the connection for the entered Web Services Proxy credentials.
5. After verifying the entered Web Services Proxy credentials through the test connection.
6. Click **Next**

The Web Services Proxy credentials for the SANtricity Cloud Connector is accepted and the Select Storage Arrays page is displayed in the Configuration Wizard.

Step 5: Select storage arrays

Based on the SANtricity Web Services Proxy credentials entered through the Configuration Wizard, a list of available storage arrays is displayed under the Select Storage Arrays page. Through this page, you can select

which storage arrays the SANtricity Cloud Connector uses for backup and restore jobs.

Before you begin

Make sure you have storage arrays configured to your SANtricity Web Services Proxy application.



Unreachable storage arrays observed by the SANtricity Cloud Connector application will result in API exceptions in the log file. This is the intended behavior of the SANtricity Cloud Connector application whenever a volume list is pulled from an unreachable array. To avoid these API exceptions in the log file, you can resolve the root issue directly with the storage array or remove the affected storage array from the SANtricity Web Services Proxy application.

Steps

1. Select each checkbox next to the storage array that you want to assign to the SANtricity Cloud Connector application for backup and restore operations.
2. Click **Next**.

The selected storage arrays are accepted, and the Select Hosts page is displayed in the Configuration Wizard.



You must configure a valid password for any storage array selected under the Select Storage Arrays page. You can configure storage array passwords through the SANtricity Web Services Proxy API Documentation.

Step 6: Select hosts

Based on the Web Services Proxy-hosted storage arrays selected through the Configuration Wizard, you can select an available host to map backup and restore candidate volumes to the SANtricity Cloud Connector application through the Select Hosts page.

Before you begin

Make sure you have a host available through the SANtricity Web Services Proxy.

Steps

1. In the drop-down menu for the listed storage array, select the desired host.
2. Repeat step 1 for any additional storage arrays listed under the Select Host page.
3. Click **Next**.

The selected host for the SANtricity Cloud Connector is accepted and the Review page is displayed in the Configuration Wizard.

Step 7: Review the initial configuration

The final page of the SANtricity Cloud Connector configuration wizard provides a summary of the entered results for your review.

Review the results of the validated configuration data.

- If all configuration data is successfully validated and established, click **Finish** to complete the configuration process.
- If any section of the configuration data cannot be validated, click **Back** to navigate to the applicable page of

the configuration wizard to revise the submitted data.

Log into the SANtricity Cloud Connector

You can access the graphical user interface for the SANtricity Cloud Connector application through the configured server in a supported browser. Make sure you have an established SANtricity Cloud Connector account.

Steps

1. In a supported browser, connect to the configured SANtricity Cloud Connector server (for example, `http://localhost:8080/`).

The login page for the SANtricity Cloud Connector application is displayed.

2. Enter your configured administrator password.
3. Click **Login**.

The landing page for the SANtricity Cloud Connector application is displayed.

Backups

You can access the Backups option in the left navigation panel of the SANtricity Cloud Connector application. The Backups option displays the Backups page, which allows you to create new image-based or file-based backup jobs.

Use the **Backups** page of the SANtricity Cloud Connector application to create and process backups of E-Series volumes. You can create image-based or file-based backups and then perform those operations immediately or at a later time. In addition, you can choose to perform full backups or incremental backups based on the last performed full backup. A maximum of six incremental backups can be performed based on the last full backup performed through the SANtricity Cloud Connector application.



All timestamps for backup and restore jobs listed under the SANtricity Cloud Connector application use local time.

Create a new image-based backup

You can create new image-based backups through the Create function on the Backups page of the SANtricity Cloud Connector application.

Before you begin

Make sure you have storage arrays from the Web Services Proxy registered to the SANtricity Cloud Connector.

Steps

1. In the Backups page, click **Create**.

The Create Backup window is displayed.

2. Select **Create an image-based backup**.
3. Click **Next**.

A list of available E-Series volumes is displayed in the Create Backup window.

4. Select the desired E-Series volume and click **Next**.

The **Name the backup and provide a description** page of Create Backup confirmation window is displayed.

5. To modify the auto-generated backup name, enter the desired name in the **Job Name** field.
6. If needed, add a description for the backup in the **Job Description** field.



You should enter a job description that allows you to easily identify the contents of the backup.

7. Click **Next**.

A summary of the selected image-based backup is displayed under the **Review backup information** page of the Create Backup window.

8. Review the selected backup and click **Finish**.

The confirmation page of the Create Backup window is displayed.

9. Select one of the following options:

- **YES** — Initiates a full backup for the selected backup.
- **NO** — A full backup for the selected image-based backup is not performed.



A full backup for the selected image-based backup can be performed at a later time through the Run function on the Backups page.

10. Click **OK**.

The backup for the selected E-Series volume is initiated, and the status for the task is displayed under the result list section of the Backups page.

Create a new folder/file-based backup

You can create new folder/file-based backups through the Create function on the Backups page of the SANtricity Cloud Connector application.

Before you begin

Make sure you have storage arrays from the Web Services Proxy registered to the SANtricity Cloud Connector.

A file-based backup unconditionally backs up all files on the filesystem you specify. However, you can perform a selective restore of files and folders.

Steps

1. In the Backups page, click **Create**.

The Create Backup window is displayed.

2. Select **Create a folder/file-based backup**.
3. Click **Next**.

A list of volumes containing file systems available for backup is displayed in the Create Backup window.

4. Select the desired volume and click **Next**.

A list of available filesystems on the selected volume is displayed in the Create Backup window.



If your filesystem does not appear, verify your filesystem type is supported by the SANtricity Cloud Connector application. For more information, refer to [Supported file systems](#).

5. Select the desired filesystem containing the folder or files to backup, and click **Next**.

The **Name the backup and provide a description** page of Create Backup confirmation window is displayed.

6. To modify the auto-generated backup name, enter the desired name in the **Job Name** field.
7. If needed, add a description for the backup in the **Job Description** field.



You should enter a job description that allows you to easily identify the contents of the backup.

8. Click **Next**.

A summary of the selected folder/file-based backup is displayed under the **Review backup information** page of the Create Backup window.

9. Review the selected folder/file-based backup and click **Finish**.

The confirmation page of the Create Backup window is displayed.

10. Select one of the following options:

- **YES** — Initiates a full backup for the selected backup.
- **NO** — A full backup for the selected backup is not performed.



A full backup for the selected file-based backup can also be performed at a later time through the Run function on the Backups page.

11. Click **Close**.

The backup for the selected E-Series volume is initiated, and the status for the task is displayed under the result list section of the Backup page.

Run Full and Incremental Backups

You can perform full and incremental backups through the Run function on the Backups page. Incremental backups are only available for file-based backups.

Before you begin

Make sure you have created a backup job through the SANtricity Cloud Connector.

Steps

1. In the Backups tab, select the desired backup job and click **Run**.



A full backup is performed automatically whenever an image-based backup job or a backup job without a previously performed initial backup is selected.

The Run Backup window is displayed.

2. Select one of the following options:

- **Full** — Backs up all data for the selected file-based backup.
- **Incremental** — Backs up changes made only since the last performed backup.



A maximum number of six incremental backups can be performed based on the last full backup performed through the SANtricity Cloud Connector application.

3. Click **Run**.

The backup request is initiated.

Delete a backup job

The Delete function deletes backed up data at the specified target location for the selected backup along with backup set.

Before you begin

Make sure there is a backup with a status of Completed, Failed, or Canceled.

Steps

1. In the Backups page, select the desired backup and click **Delete**.



If a full base backup is selected for deletion, all associated incremental backups are also deleted.

The Confirm Delete window is displayed.

2. In the **Type delete** field, type `DELETE` to confirm the delete action.

3. Click **Delete**.

The selected backup is deleted.

Restores

You can access the Restore option in the left navigation panel of the SANtricity Cloud Connector application. The Restore option displays the Restore page, which allows you to create new image-based or file-based restore jobs.

The SANtricity Cloud Connector uses the concept of jobs to perform the actual restore of an E-Series volume. Before performing a restore, you must identify which E-Series volume will be used for the operation. After you add an E-Series volume for restore to the SANtricity Cloud Connector host, you can use the `Restore` page of the SANtricity Cloud Connector application to create and process restores.



All timestamps for backup and restore jobs listed under the SANtricity Cloud Connector application use local time.

Create a new image-based restore

You can create new image-based restores through the Create function on the Restore page of the SANtricity Cloud Connector application.

Before you begin

Make sure you have an image-based backup available through the SANtricity Cloud Connector.

Steps

1. In the Restore page of the SANtricity Cloud Connector application, click **Create**.

The Restore window is displayed.

2. Select the desired backup.
3. Click **Next**.

The Select Backup Point page is displayed in the Restore window.

4. Select the desired completed backup.
5. Click **Next**.

The Select Restore Target page is displayed in the Restore window.

6. Select the restore volume and click **Next**.

The Review page is displayed in the Restore window.

7. Review the selected restore operation and click **Finish**.

The restore for the selected target host volume is initiated, and the status for the task is displayed in the result list section of the Restore page.

Create a new file-based restore

You can create new file-based restores through the Create function in the Restore page of the SANtricity Cloud Connector application.

Before you begin

Make sure you have a file-based backup available through the SANtricity Cloud Connector.

Steps

1. In the Restore page of the SANtricity Cloud Connector application, click **Create**.

The Restore window is displayed.

2. In the Restore window, select the desired file-based backup.
3. Click **Next**.

The Select Backup Point page is displayed in the Create Restore Job window.

4. In the Select Backup Point page, select the desired completed backup.
5. Click **Next**.

A list of available filesystems or folders/files page is displayed in the Restore window.

6. Select the desired folders or files to restore and click **Next**.

The Select Restore Target page is displayed in the Restore window.

7. Select the restore volume and click **Next**.

The Review page is displayed in the Restore window.

8. Review the selected restore operation and click **Finish**.

The restore for the selected target host volume is initiated, and the status for the task is displayed in the result list section of the Restore page.

Delete a restore

You can use the Delete function to delete a selected restore item from the result list section of the Restore page.

Before you begin

Make sure there is a restore job with a status of Completed, Failed or Canceled.

Steps

1. In the Restore page, click **Delete**.

The Confirm Delete window is displayed.

2. In the **Type delete** field, type `delete` to confirm the delete action.
3. Click **Delete**.



You cannot delete a suspended restore.

The selected restore is deleted.

Modify the SANtricity Cloud Connector settings

The Settings option allows you to modify the application's current configurations for the S3 account, managed storage arrays and hosts, and Web Services Proxy credentials. You can also change the password for the SANtricity Cloud Connector application through the Settings option.

Modify S3 Account settings

You can modify existing S3 settings for the SANtricity Cloud Connector application in the S3 Account Settings window.

Before you begin

When modifying the URL or S3 Bucket Label settings, be aware that access to any existing backups configured through the SANtricity Cloud Connector will be affected.

Steps

1. In the left toolbar, click **Settings > Configuration**.

The Settings - Configuration page is displayed.

2. Click **View/Edit Settings** for S3 Account Settings.

The S3 Account Settings page is displayed.

3. In the URL field, enter the URL for the S3 cloud service.
4. In the **Access Key ID** field, enter the access ID for the S3 target.
5. In the **Secret Access Key** field, enter the access key for the S3 target.
6. In the **S3 Bucket Name** field, enter the bucket name for the S3 target.
7. Select the **Use Path Style Access** check box if needed.
8. Click **Test Connection** to verify the connection for the entered S3 credentials.
9. Click **Save** to apply the modifications.

The modified S3 account settings are applied.

Manage storage arrays

You can add or remove storage arrays from the Web Services Proxy registered to the SANtricity Cloud Connector host in the Manage Storage Arrays page.

The Manage Storage Arrays page displays a list of storage arrays from the Web Services Proxy available for registration with the SANtricity Cloud Connector host.

Steps

1. In the left toolbar, click **Settings > Storage Arrays**.

The Settings - Storage Arrays screen is displayed.

2. To add storage arrays to the SANtricity Cloud Connector, click **Add**.
 - a. In the Add Storage Arrays window, select each checkbox next to the desired storage arrays from the result list.
 - b. Click **Add**.

The selected storage array is added to the SANtricity Cloud Connector and displays in the result list section of the Settings - Storage Arrays screen.

3. To modify the host for an added storage array, click **Edit** for the line item in the result list section of the Settings - Storage Arrays screen.
 - a. In the Associated Host drop-down menu, select the desired host for the storage array.
 - b. Click **Save**.

The selected host is assigned to the storage array.

4. To remove an existing storage array from the SANtricity Cloud Connector host, select the desired storage arrays from the bottom result list, and click **Remove**.
 - a. In the Confirm Remove Storage Array field, type `REMOVE`.
 - b. Click **Remove**.

The selected storage array is removed from the SANtricity Cloud Connector host.

Modify Web Services Proxy settings

You can modify existing Web Services Proxy settings for the SANtricity Cloud Connector application in the Web Services Proxy Settings window.

Before you begin

The Web Services Proxy used with the SANtricity Cloud Connector needs to have the appropriate arrays added and the corresponding password set.

Steps

1. In the left toolbar, click **Settings > Configuration**.

The Settings - Configuration screen is displayed.

2. Click **View/Edit Settings** for Web Services Proxy.

The Web Services Proxy settings screen is displayed.

3. In the URL field, enter the URL for the Web Services proxy used for the SANtricity Cloud Connector.
4. In the User Name field, enter the user name for the Web Services Proxy connection.
5. In the Password field, enter the password for the Web Services Proxy connection.
6. Click **Test Connection** to verify the connection for the entered Web Services Proxy credentials.
7. Click **Save** to apply the modifications.

Change SANtricity Cloud Connector password

You can change the password for the SANtricity Cloud Connector application in the Change Password screen.

Steps

1. In the left toolbar, click **Settings > Configuration**.

The Settings - Configuration screen is displayed.

2. Click **Change Password** for SANtricity Cloud Connector.

The Change Password screen is displayed.

3. In the Current password field, enter your current password for the SANtricity Cloud Connector application.
4. In the New Password field, enter your new password for the SANtricity Cloud Connector application.
5. In the Confirm new password field, re-enter the new password.
6. Click **Change** to apply the new password.

The modified password is applied to the SANtricity Cloud Connector application.

Uninstall the SANtricity Cloud Connector

You can uninstall the SANtricity Cloud Connector through the graphical uninstaller or console mode.

Uninstall using graphical mode

You can use the graphical mode to uninstall the SANtricity Cloud Connector on a Linux operating system.

Steps

1. From a terminal window, navigate to the directory containing the SANtricity Cloud Connector uninstall file.

The uninstall file for the SANtricity Cloud Connector is available at the following default directory location:

```
/opt/netapp/santricity_cloud_connector4/uninstall_cloud_connector4
```

2. From the directory containing the SANtricity Cloud Connector uninstall file, run the following command:

```
./uninstall_cloud_connector4 -i gui
```

The uninstall process for the SANtricity Cloud Connector is initialized.

3. In the uninstall window, click **Uninstall** to proceed with uninstalling the SANtricity Cloud Connector.

The uninstall process is completed, and the SANtricity Cloud Connector application is uninstalled in the Linux operating system.

Uninstall using console mode

You can use the console mode to uninstall the SANtricity Cloud Connector on a Linux operating system.

Steps

1. From a terminal window, navigate to the directory containing the SANtricity Cloud Connector uninstall file.

The uninstall file for the SANtricity Cloud Connector is available at the following default directory location:

```
/opt/netapp/santricity_cloud_connector4/uninstall_cloud_connector4
```

2. From the directory containing the SANtricity Cloud Connector uninstall file, run the following command:

```
./uninstall_cloud_connector4 -i console
```

The uninstall process for the SANtricity Cloud Connector is initialized.

3. In the uninstall window, press **Enter** to proceed with uninstalling the SANtricity Cloud Connector.

The uninstall process is completed, and the SANtricity Cloud Connector application is uninstalled in the

Linux operating system.

Copyright information

Copyright © 2022 NetApp, Inc. All Rights Reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

LIMITED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (b)(3) of the Rights in Technical Data -Noncommercial Items at DFARS 252.227-7013 (FEB 2014) and FAR 52.227-19 (DEC 2007).

Data contained herein pertains to a commercial product and/or commercial service (as defined in FAR 2.101) and is proprietary to NetApp, Inc. All NetApp technical data and computer software provided under this Agreement is commercial in nature and developed solely at private expense. The U.S. Government has a non-exclusive, non-transferrable, nonsublicensable, worldwide, limited irrevocable license to use the Data only in connection with and in support of the U.S. Government contract under which the Data was delivered. Except as provided herein, the Data may not be used, disclosed, reproduced, modified, performed, or displayed without the prior written approval of NetApp, Inc. United States Government license rights for the Department of Defense are limited to those rights identified in DFARS clause 252.227-7015(b) (FEB 2014).

Trademark information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.